


CPUs

- Input and output.
- Supervisor mode, exceptions, traps.
- Co-processors.

I/O devices

- Usually includes some non-digital component.
- Typical digital interface to CPU:




Application: 8251 UART

- Universal asynchronous receiver transmitter (UART) : provides serial communication.
- 8251 functions are integrated into standard PC interface chip.
- Allows many communication parameters to be programmed.

Serial communication


- Characters are transmitted separately:



Serial communication parameters

- Baud (bit) rate.
- Number of bits per character.
- Parity/no parity.
- Even/odd parity.
- Length of stop bit (1, 1.5, 2 bits).

8251 CPU interface



Serial port mode bits

- mode[1:0]: baud rate
 - 00: synchronous
 - 01: async, no clock prescaler
 - 10: async, 16x prescaler
 - 11: async, 64x prescaler
- mode[3:2]: bits/char
 - 00: 5 bits
 - 01: 6 bits
 - 10: 7 bits
 - 11: 8 bits.
- mode[5:4]: parity
 - 00, 10: no parity
 - 01: odd parity
 - 11: even parity
- mode[7:6]: stop bit
 - 00: invalid
 - 01: 1 stop bit
 - 10: 1.5 stop bits
 - 11: 2 stop bits

Serial port command register

- mode[0]: transmit enable
- mode[1]: set nDTR output
- mode[2]: enable receiver
- mode[3]: send break
- mode[4]: reset error flags
- mode[5]: set nRTS
- mode[6]: internal reset
- mode[7]: hunt mode

Serial port status register

- status[0]: transmitter ready
- status[1]: receiver ready
- status[2]: transmission complete
- status[3]: parity
- status[4]: overrun
- status[5]: frame error
- status[6]: sync char deleted
- status[7]: nDSR value

Programming I/O

- Two types of instructions can support I/O:
 - special-purpose I/O instructions;
 - memory-mapped load/store instructions.
- Intel x86 provides `in`, `out` instructions. Most other CPUs use memory-mapped I/O.
- I/O instructions do not preclude memory-mapped I/O.

ARM memory-mapped I/O

- Define location for device:

```
DEV1 EQU 0x1000
```

- Read/write code:

```
LDR r1,#DEV1 ; set up device adrs
```

```
LDR r0,[r1] ; read DEV1
```

```
LDR r0,#8 ; set up value to write
```

```
STR r0,[r1] ; write value to device
```

SHARC memory mapped I/O

- Device must be in external memory space (above 0x400000).
- Use DM to control access:

`I0 = 0x400000;`

`M0 = 0;`

`R1 = DM(I0,M0);`

Peek and poke

- Traditional HLL interfaces:

```
int peek(char *location) {  
    return *location; }
```

```
void poke(char *location, char newval) {  
    (*location) = newval; }
```

Busy/wait output

- Simplest way to program device.
 - Use instructions to test when device is ready.

```
current_char = mystring;
while (*current_char != '\0') {
    poke(OUT_CHAR, *current_char);
    while (peek(OUT_STATUS) != 0);
    current_char++;
}
```


Simultaneous busy/wait input and output

```
while (TRUE) {  
    /* read */  
    while (peek(IN_STATUS) == 0);  
    achar = (char)peek(IN_DATA);  
    /* write */  
    poke(OUT_DATA, achar);  
    poke(OUT_STATUS, 1);  
    while (peek(OUT_STATUS) != 0);  
}
```

Interrupt I/O

- Busy/wait is very inefficient.
 - CPU can't do other work while testing device.
 - Hard to do simultaneous I/O.
- Interrupts allow a device to change the flow of control in the CPU.
 - Causes subroutine call to handle device.

Interrupt interface



Interrupt behavior

- Based on subroutine call mechanism.
- Interrupt forces next instruction to be a subroutine call to a predetermined location.
 - Return address is saved to resume executing foreground program.

Interrupt physical interface

- CPU and device are connected by CPU bus.
- CPU and device handshake:
 - device asserts interrupt request;
 - CPU asserts interrupt acknowledge when it can handle the interrupt.

Example: character I/O handlers


```
void input_handler() {  
    achar = peek(IN_DATA);  
    gotchar = TRUE;  
    poke(IN_STATUS, 0);  
}  
void output_handler() {  
}
```

Example: interrupt-driven main program

```
main() {  
    while (TRUE) {  
        if (gotchar) {  
            poke(OUT_DATA, achar);  
            poke(OUT_STATUS, 1);  
            gotchar = FALSE;  
        }  
    }  
}
```

Example: interrupt I/O with buffers


- Queue for characters:



Buffer-based input handler

```
void input_handler() {  
    char achar;  
    if (full_buffer()) error = 1;  
    else { achar = peek(IN_DATA); add_char(achar); }  
    poke(IN_STATUS,0);  
    if (nchars() >= 1)  
        { poke(OUT_DATA,remove_char()); poke(OUT_STATUS,1); }  
}
```

I/O sequence diagram




Debugging interrupt code

- What if you forget to change registers?
 - Foreground program can exhibit mysterious bugs.
 - Bugs will be hard to repeat---depend on interrupt timing.

Priorities and vectors

- Two mechanisms allow us to make interrupts more specific:
 - **Priorities** determine what interrupt gets CPU first.
 - **Vectors** determine what code is called for each type of interrupt.
- Mechanisms are orthogonal: most CPUs provide both.


Prioritized interrupts



Interrupt prioritization


- **Masking**: interrupt with priority lower than current priority is not recognized until pending interrupt is complete.
- **Non-maskable interrupt (NMI)**: highest-priority, never masked.
 - Often used for power-down.

Example: Prioritized I/O




Interrupt vectors


- Allow different devices to be handled by different code.
- Interrupt vector table:



Interrupt vector acquisition



Generic interrupt mechanism



Interrupt sequence

- CPU acknowledges request.
- Device sends vector.
- CPU calls handler.
- Software processes request.
- CPU restores state to foreground program.

Sources of interrupt overhead

- Handler execution time.
- Interrupt mechanism overhead.
- Register save/restore.
- Pipeline-related penalties.
- Cache-related penalties.

ARM interrupts

- ARM7 supports two types of interrupts:
 - Fast interrupt requests (FIQs).
 - Interrupt requests (IRQs).
 - Priorities: FIQs > IRQs
- Interrupt table starts at location 0.

ARM interrupt procedure

- CPU actions:
 - Save PC. Copy CPSR to SPSR.
 - Force bits in CPSR to record interrupt.
 - Force PC to vector.
- Handler responsibilities:
 - Restore proper PC.
 - Restore CPSR from SPSR.
 - Clear interrupt disable flags.

ARM interrupt latency

- Worst-case latency to respond to interrupt is 27 cycles:
 - Two cycles to synchronize external request.
 - Up to 20 cycles to complete current instruction.
 - Three cycles for data abort.
 - Two cycles to enter interrupt handling state.
- Best-case can be 4 cycles

C55x interrupts

- Latency is between 7 and 13 cycles.
- Maskable interrupt sequence:
 - Interrupt flag register is set.
 - Interrupt enable register is checked.
 - Interrupt mask register is checked.
 - Interrupt flag register is cleared.
 - Appropriate registers are saved.
 - INTM set to 1, DBGGM set to 1, EALLOW set to 0.
 - Branch to ISR.
- Two styles of return: fast and slow.

PIC16F interrupts

- Synchronous interrupts come from inside CPU.
- Asynchronous interrupts come from outside CPU.
- INTCON register holds interrupt control bits.
 - GIE allows unmasked interrupts.
 - PIE allows peripheral interrupts.
 - TMR0 enables timer 0.
 - INT enables external interrupts.
 - PIR1/PIR2 control peripheral interrupts.
- RETFIE used to return from interrupts.
- Latency:
 - Synchronous: $3T_{cy}$
 - Asynchronous: 3 to $3.75 T_{cy}$

Supervisor mode

- May want to provide protective barriers between programs.
 - Avoid memory corruption.
- Need **supervisor mode** to manage the various programs.
- SHARC does not have a supervisor mode.

ARM supervisor mode

- Use SWI instruction to enter supervisor mode, similar to subroutine:
SWI CODE_1
- Sets PC to 0x08.
- Argument to SWI is passed to supervisor mode code.
- Saves CPSR in SPSR.

Exception

- **Exception**: internally detected error.
- Exceptions are synchronous with instructions but unpredictable.
- Build exception mechanism on top of interrupt mechanism.
- Exceptions are usually prioritized and vectorized.

Trap

- **Trap (software interrupt)**: an exception generated by an instruction.
 - Call supervisor mode.
- ARM uses SWI instruction for traps.
- SHARC offers three levels of software interrupts.
 - Called by setting bits in IRPTL register.

Co-processor


- **Co-processor**: added function unit that is called by instruction.
 - Floating-point units are often structured as co-processors.
- ARM allows up to 16 designer-selected co-processors.
 - Floating-point co-processor uses units 1, 2.
- C55x uses co-processors as well.

Fixed-point

- $0100.0110 = 4 + 3/8$
- $0001.0000 = 1$
- $0000.1000 = 1/2$
- $0000.0101 = 5/16$
- $0000.0100 = 1/4$

Floating-point

$$12.345 = \underbrace{12345}_{\text{significand}} \times \underbrace{10^{-3}}_{\text{base}}^{\text{exponent}}$$




- 0.15625 (Dec) = 0.00101 (Bin)

C55x image/video hardware extensions

- Available in 5509 and 5510.
 - Equivalent C-callable functions for other devices.
- Available extensions:
 - DCT/IDCT.
 - Pixel interpolation
 - Motion estimation.

DCT/IDCT


- 2-D DCT/IDCT is computed from two 1-D DCT/IDCT.
 - Put data in different banks to maximize throughput.



C55 DCT/IDCT coprocessor extensions

- Load, compute, transfer to accumulators:
 - $ACy = \text{copr}(k8, ACx, Xmem, Ymem)$
- Compute, transfer, mem write:
 - $ACy = \text{copr}(k8, ACx, ACy), Lmem = ACz$
- Special:
 - $ACy = \text{copr}(k8, ACx, ACy)$

Software pipelined load/compute/store for DCT




C55 motion estimation

- Search strategy:
 - Full vs. non-full.
- Accuracy:
 - Full-pixel vs. half-pixel.
- Number of returned motion vectors:
 - 1 (one 16x16) vs. 4 (four 8x8).
- Algorithms:
 - 3-step algorithm (distance 4,2,1).
 - 4-step algorithm (distance 8,4,2,1).
 - 4-step with half-pixel refinement.

Four-step motion estimation breakdown

```
d = {8,4,2,1};  
for (i=0; i<4; i++) {  
    compute 3 upper differences  
    for d[i];  
    compute 3 middle differences  
    for d[i];  
    compute 3 lower differences for  
    d[i];  
    compute minimum value;  
    move to next d;  
}
```




C55 motion estimation accelerator

- Includes 3 16-bit pixel data paths, 3 16-bit absolute differences (ADs).
- Basic operation:
 - $[ACx, ACy] = \text{copr}(k8, ACx, ACy, Xmem, Ymem, \text{Coeff})$
 - K8 = control bits (enable AD units, etc.)
 - ACx, ACy = accumulated absolute differences
 - Xmem, Ymem = pointers to odd, even lines of the search window
 - Pointer to two adjacent pixels from reference window

C55 pixel interpolation

- Given four pixels A, B, C, D, interpolate three half-pixels:




Pixel interpolation coprocessor operations

- Load pixels and compute:
 - $ACy = \text{copr}(k8, AC, Lmem)$
- Load pixels, compute, and store:
 - $ACy = \text{copr}(k8, AACx, Lmem) \parallel Lmem = ACz$

CPUs

- Caches.
- Memory management.

Caches and CPUs



Cache operation

- Many main memory locations are mapped onto one cache entry.
- May have caches for:
 - instructions;
 - data;
 - data + instructions (**unified**).
- Memory access time is no longer deterministic.

Terms

- **Cache hit**: required location is in cache.
- **Cache miss**: required location is not in cache.
- **Working set**: set of locations used by program in a time interval.


Types of misses

- **Compulsory** (**cold**): location has never been accessed.
- **Capacity**: working set is too large.
- **Conflict**: multiple locations in working set map to same cache entry.

Memory system performance

- h = cache hit rate.
- t_{cache} = cache access time, t_{main} = main memory access time.
- Average memory access time:
 - $t_{\text{av}} = ht_{\text{cache}} + (1-h)t_{\text{main}}$ (Wrong)
 - $= ht_{\text{cache}} + (1-h)(t_{\text{cache}} + t_{\text{main}})$

Multiple levels of cache



Multi-level cache access time

- h_1 = cache hit rate.
- h_2 = rate for miss on L1, hit on L2.
- Average memory access time:
 - $t_{av} = h_1 t_{L1} + (h_2 - h_1) t_{L2} + (1 - h_2 - h_1) t_{main}$ (Wrong)
 - $t_{av} = h_1 t_{L1} + (1 - h_1)(h_2(t_{L1} + t_{L2}) + (1 - h_2)(t_{main} + t_{L1} + t_{L2}))$

Replacement policies

- **Replacement policy**: strategy for choosing which cache entry to throw out to make room for a new memory location.
- Two popular strategies:
 - Random.
 - Least-recently used (LRU).


Cache organizations

- **Fully-associative**: any memory location can be stored anywhere in the cache (almost never implemented).
- **Direct-mapped**: each memory location maps onto exactly one cache entry.
- **N-way set-associative**: each memory location can go into one of n sets.

Cache performance benefits

- Keep frequently-accessed locations in fast cache.
- Cache retrieves more than one word at a time.
 - Sequential accesses are faster after first access.

Direct-mapped cache



Write operations


- **Write-through**: immediately copy write to main memory.
- **Write-back**: write to main memory only when location is removed from cache.
 - eliminate the write when a location is written several times before it is removed from the cache

Direct-mapped cache locations

- Many locations map onto the same cache block.
- Conflict misses are easy to generate:
 - Array `a[]` uses locations 0, 1, 2, ...
 - Array `b[]` uses locations 1024, 1025, 1026, ...
 - Operation `a[i] + b[i]` generates conflict misses.

Set-associative cache

- A set of direct-mapped caches:
 - extra overhead, higher hit rate



Example: direct-mapped vs. set-associative

address

000

001

010

011

100

101

110

111

data

0101

1111

0000

0110

1000

0001

1010

0100

Direct-mapped cache behavior

- After 001 access:

block	tag	data
00	-	-
01	0	1111
10	-	-
11	-	-

- After 010 access:

block	tag	data
00	-	-
01	0	1111
10	0	0000
11	-	-

Direct-mapped cache behavior, cont'd.

- After 011 access:

block	tag	data
00	-	-
01	0	1111
10	0	0000
11	0	0110

- After 100 access:

block	tag	data
00	1	1000
01	0	1111
10	0	0000
11	0	0110

Direct-mapped cache behavior, cont'd.

- After 101 access:

block	tag	data
00	1	1000
01	1	0001
10	0	0000
11	0	0110

- After 111 access:

block	tag	data
00	1	1000
01	1	0001
10	0	0000
11	1	0100

2-way set-associative cache behavior

- Final state of cache (twice as big as direct-mapped):

blk	set 0 tag	set 0 data	set 1 tag	set 1 data
00	1	1000	-	-
01	0	1111	1	0001
10	0	0000	-	-
11	0	0110	1	0100

2-way set-associative cache behavior

- Final state of cache (same size as direct-mapped):


blk set 0 tag		set 0 data	set 1 tag	set 1 data
0	01	0000	10	1000
1	10	0001	11	0100

Example caches

- StrongARM:
 - 16 Kbyte, 32-way, 32-byte block instruction cache.
 - 16 Kbyte, 32-way, 32-byte block data cache (write-back).
- SHARC:
 - 32-instruction, 2-way instruction cache.

Memory management units

- Memory management unit (MMU) translates addresses:




Memory management tasks

- Allows programs to move in physical memory during execution.
- Allows **virtual memory**:
 - memory images kept in secondary storage;
 - images returned to main memory on demand during execution.
- **Page fault**: request for location not resident in memory.


Address translation

- Requires some sort of register/table to allow arbitrary mappings of logical to physical addresses.
- Two basic schemes:
 - segmented;
 - paged.
- Segmentation and paging can be combined (x86).


Segments and pages




Segment address translation



Page address translation



Page table organizations




Caching address translations

- Large translation tables require main memory access.
- **TLB**: Translation Lookaside Buffer
 - cache for address translation.
 - Typically small.

ARM memory management

- Memory region types:
 - section: 1 Mbyte block;
 - large page: 64 kbytes;
 - small page: 4 kbytes.
- An address is marked as section-mapped or page-mapped.
- Two-level translation scheme.

ARM address translation



CPUs

- CPU performance.
- CPU power consumption and power management.
- Safety and security.

Elements of CPU performance

- Cycle time.
- CPU pipeline.
- Memory system.



Pipelining

- Several instructions are executed simultaneously at different stages of completion.
- Various conditions can cause **pipeline bubbles** that reduce utilization:
 - branches;
 - memory system delays;
 - etc.

Performance measures

- **Latency**: time it takes for an instruction to get through the pipeline.
- **Throughput**: number of instructions executed per time period.
- Pipelining increases throughput without reducing latency.

Single-cycle, nonpipelining & Pipelining



Pipelining is an implementation technique in which multiple instructions are overlapped in execution.


Feature

- Optimal Pipeline
 - One phase of instruction execution per cycle
 - At each clock cycle, one instruction is completed
 - On average, the execution is much faster
- What makes this work?
 - The similarity between instructions makes it possible for all instructions to use the same stage (in general)
 - Each stage takes about the same amount of time: less time wasted


Problems for Pipelining CPUs

- Limits to pipelining: [Hazards](#) prevent next instruction from executing during its designated clock cycle
 - [Structural hazards](#): HW cannot support some combination of instructions that we want to execute in the same clock cycle.
 - [Data hazards](#): Instruction depends on result of prior instruction still in the pipeline.
 - [Control hazards](#): Pipelining of branches causes later instruction fetches to wait for the result of the branch.
- So we need [pipeline stalls](#) or “[bubbles](#)”.

Structural Hazard #1: Single Memory



Structural Hazard #2: Registers




Can we read and write to registers simultaneously?


Resolving Structural Hazards

- Structural hazards occurs when two instruction need same hardware resource at same time
 - Can resolve in hardware by stalling newer instruction till older instruction finished with resource
- A structural hazard can always be avoided by adding more hardware to design
 - E.g., if two instructions both need a port to memory at same time, could avoid hazard by adding second port to memory

Datapath with Control



Pipelined Dependencies




Forwarding

	Time (in clock cycles) →								
	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register \$2 :	10	10	10	10	10/-20	-20	-20	-20	-20
Value of EX/MEM :	X	X	X	-20	X	X	X	X	X
Value of MEM/WB :	X	X	X	X	-20	X	X	X	X

MEM data hazard detection:


Program execution order (in instructions)




If (MEM/WB.RegWrite &&
(MEM/WB.RegisterRd ≠ 0)
&& (MEM/WB.RegisterRd
= ID/EX.RegisterRs))

If (MEM/WB.RegWrite &&
(MEM/WB.RegisterRd ≠ 0)
&& (MEM/WB.RegisterRd
= ID/EX.RegisterRt))


Forwarding




Stalling



Hazard Detection Unit



Control Hazard: Branching




Control Hazard: Branching

- Stall until decision is made
 - Insert the "no-op" instruction (which does nothing but consume time) or hold up the fetch of the next instruction (for 2 cycles).
 - Cons: Each branch statement takes 3 clock cycles (assuming the comparison is made in the ALU phase)

Control Hazard: Branching

- Optimization #1:
 - insert **special branch comparator** in Stage 2
 - as soon as instruction is decoded (Opcode identifies it as a branch), immediately make a decision and set the new value of the PC
 - Benefit: since branch is complete in Stage 2, only one unnecessary instruction is fetched, so only one no-op is needed
 - Side Note: This means that branches are idle in Stages 3, 4 and 5.

Control Hazard: Branching (6/8)



Control Hazard: Branching

- Optimization #2: Redefining branches
 - Old definition: if we take the branch, none of the instructions after the branch get executed by accident
 - New definition: whether or not we take the branch, the single instruction immediately following the branch gets executed (called the **branch-delay slot**)
- The term “**Delayed Branch**” means **we always execute inst after branch**

Control Hazard: Branching

- Notes on **Branch-Delay Slot**
 - Worst case: can always put a no-op in the branch-delay slot
 - Better case: Find an instruction before the branch, which can be put into the branch-delay slot without affecting the flow of the program
 - re-ordering instructions is a common method of speeding up programs
 - compiler must be very smart in order to find instructions to do this
 - usually can find such an instruction at least 50% of the time
 - Jumps also have a delay slot...

Example: Nondelayed vs. Delayed Branch

Nondelayed Branch

or \$8, \$9, \$10

add \$1, \$2, \$3

sub \$4, \$5, \$6

beq \$1, \$4, Exit

xor \$10, \$1, \$11

Exit:

Delayed Branch

add \$1, \$2, \$3


sub \$4, \$5, \$6

beq \$1, \$4, Exit

or \$8, \$9, \$10

xor \$10, \$1, \$11


Exit:



ARM7 pipeline

- ARM 7 has 3-stage pipe:
 - **fetch** instruction from memory;
 - **decode** opcode and operands;
 - **execute**.


ARM pipeline execution



Pipeline stalls

- If every step cannot be completed in the same amount of time, pipeline stalls.
- Bubbles introduced by stall increase latency, reduce throughput.


ARM multi-cycle LDmia instruction



Control stalls

- Branches often introduce stalls (branch penalty).
 - Stall time may depend on whether branch is taken.
- May have to squash instructions that already started executing.
- Don't know what to fetch until condition is evaluated.

ARM pipelined branch



Delayed branch

- To increase pipeline efficiency, delayed branch mechanism requires n instructions after branch always executed whether branch is executed or not.
- SHARC supports delayed and non-delayed branches.
 - Specified by bit in branch instruction.
 - 2 instruction branch delay slot.

Example: ARM execution time

- Determine execution time of FIR filter:
 for (i=0, f=0; i<N; i++)
 f = f + c[i]*x[i];
- Only branch in loop test may take more than one cycle.
 - BLT loop takes 1 cycle best case, 3 worst case.

FIR filter ARM code

; loop initiation code

MOV r0,#0 ; use r0 for i, set to 0

MOV r8,#0 ; use a separate index for arrays

ADR r2,N ; get address for N

LDR r1,[r2] ; get value of N

MOV r2,#0 ; use r2 for f, set to 0

ADR r3,c ; load r3 with address of base of c

ADR r5,x ; load r5 with address of base of x

; test for exit

loop CMP r0,r1

BGE loopend ; if I >= N, exit loop

; loop body

LDR r4,[r3,r8] ; get value of c[i]

LDR r6,[r5,r8] ; get value of x[i]

MUL r4,r4,r6 ; compute c[i]*x[i]

ADD r2,r2,r4 ; add into running sum

; update


ADD r8,r8,#4 ; add one to array index

ADD r0,r0,#1 ; add 1 to i

B loop

loopend ...

FIR filter block diagram



FIR filter performance by block

Block	Variable	# instructions	# cycles
Initiation	t_{init}	7	7
Test	t_{test}	2	[2,4]
Body	t_{body}	4	4
Update	t_{update}	3	3

$$t_{\text{loop}} = t_{\text{init}} + N(t_{\text{body}} + t_{\text{update}} + t_{\text{test,best}}) + t_{\text{test,worst}}$$

Loop exit fails is best case

Loop exit succeeds is worst case

PIC16F execution time

- Instruction divided into four Q cycles:
 - Q1 decodes.
 - Q2 reads operands.
 - Q3 performs data operations.
 - Q4 writes data.
- One Q cycle per clock cycle.
- Execution time for one instruction is T_{cy} .

PIC16F instruction timing

- Most instructions execute in one cycle, except:
 - CALL, GOTO, RETFIE, RETLW, RETURN require 2 cycles.
 - Skip-if instructions require two cycles if skip is taken.

PIC16F loop execution timing


- This program sets a bit on an I/O device for a data-dependent amount of time.
- Computed goto determines number of bsfs executed.
 - len=1 -> go to len3, etc.

```
movf len, w ; get ready for computed goto
addwf pcl, f ; computed goto (PCL is low byte of PC)
len3: bsf x, l ; set the bit at t-3
len2: bsf x, l ; set the bit at t-2
len1: bsf x, l ; set the bit at t-1
bcf x, l ; clear the bit at t
```

C55x pipeline

- C55x has 7-stage pipe:
 - **fetch**;
 - **decode**;
 - **address**: computes data/branch addresses;
 - **access 1**: reads data;
 - **access 2**: finishes data read;
 - **Read stage**: puts operands on internal busses;
 - **execute**.

C55x organization



C55x pipeline hazards

- Processor structure:
 - Three computation units.
 - 14 operators.
- Can perform two operations per instruction.
- Some combinations of operators are not legal.

C55x hazards

- A-unit ALU/A-unit ALU.
- A-unit swap/A-unit swap.
- D-unit ALU,shifter,MAC/D-unit ALU,shifter,MAC
- D-unit shifter/D-unit shift, store
- D-unit shift, store/D-unit shift, store
- D-unit swap/D-unit swap
- P-unit control/P-unit control

Memory system performance

- Caches introduce indeterminacy in execution time.
 - Depends on order of execution.
- **Cache miss penalty**: added time due to a cache miss.

Types of cache misses

- Compulsory miss: location has not been referenced before.
- Conflict miss: two locations are fighting for the same block.
- Capacity miss: working set is too large.

CPU power consumption

- Most modern CPUs are designed with power consumption in mind to some degree.
- Power vs. energy:
 - heat depends on power consumption;
 - battery life depends on energy consumption.

CMOS power consumption

- **Voltage drops**: power consumption proportional to V^2 .
- **Toggling**: more activity means more power.
- **Leakage**: basic circuit characteristics; can be eliminated by disconnecting power.

Power and Energy

- Power is drawn from a voltage source attached to the V_{DD} pin(s) of a chip.

- Instantaneous Power:

$$P(t) = I(t)V(t)$$

- Energy:


$$E = \int_0^T P(t)dt$$

- Average Power:

$$P_{\text{avg}} = \frac{E}{T} = \frac{1}{T} \int_0^T P(t)dt$$

Switching Power and Activity Factor

$$\begin{aligned}P_{\text{switching}} &= \frac{1}{T} \int_0^T i_{DD}(t) V_{DD} dt \\&= \frac{V_{DD}}{T} \int_0^T i_{DD}(t) dt \\&= \frac{V_{DD}}{T} [T f_{\text{sw}} C V_{DD}] \\&= C V_{DD}^2 f_{\text{sw}}\end{aligned}$$




Suppose the system clock frequency = f, Let $f_{\text{sw}} = \alpha f$, where α = activity factor

If the signal is a clock, $\alpha = 1$

If the signal switches once per cycle, $\alpha = \frac{1}{2}$

So, Dynamic power: $P_{\text{switching}} = \alpha C V_{DD}^2 f$

Leakage



Sub-threshold current one of most compelling issues in low-energy circuit design!

CPU power-saving strategies

- Reduce power supply voltage.
- Run at lower clock frequency.
- Disable function units with control signals when not in use.
- Disconnect parts from power supply when not in use.

C55x low power features

- Parallel execution units---longer idle shutdown times.
- Multiple data widths:
 - 16-bit ALU vs. 40-bit ALU.
- Instruction caches minimizes main memory accesses.
- Power management:
 - Function unit idle detection.
 - Memory idle detection.
 - User-configurable IDLE domains allow programmer control of what hardware is shut down.

Power management styles

- **Static power management:** does not depend on CPU activity.
 - Example: user-activated power-down mode.
- **Dynamic power management:** based on CPU activity.
 - Example: disabling off function units.

Power-down costs

- Going into a power-down mode costs:
 - time;
 - energy.
- Must determine if going into mode is worthwhile.
- Can model CPU power states with power state machine.


ARMv8-A standby mode

- Standby keeps power on CPU but stops or gates most clocks.
- WFI instruction waits for interrupt or debug request.
- WFE instruction waits for certain specified events.

Application: NXP LPC 1311

- Four power modes:
 - Run: normal operation.
 - Sleep: stops CPU clock, with CPU and peripheral logic still powered. Peripherals may initiate interrupts that wake up the CPU.
 - Deep sleep: stops CPU clock, powers down most analog blocks, maintains CPU and peripheral registers and SRAM values.
 - Deep power-down: power and clocks shut down to entire chip, registers and SRAM not maintained.

LPC1311 power state machine



Power management

- Different physical mechanisms require different strategies:
 - Dynamic voltage and frequency scaling (DVFS) for low leakage.
 - Race-to-dark for high leakage.
- Total energy consumption comes from dynamic and static power consumption mechanisms:
 - $E_{tot} = \int_0^T P(t)dt = \int_0^T [P_{dyn}(t) + P_{static}(t)]dt$

Two power management strategies



- DVFS used for cases when power consumption is dominated by dynamic power.
- Characteristics as a function of power supply voltage:
 - Clock frequency proportional to power supply.
 - Power consumption proportional to square of power supply.
- Reducing power supply voltage reduces performance but reduces power consumption more.
- Adjust power supply voltage to level just fast enough for current tasks.
- Race-to-dark used when power consumption is dominated by static power.
- Run as fast as possible, then remove power supply from logic to eliminate static power consumption.

What is DVFS

- Dynamic voltage and frequency scaling
 - Circuits can work at a range of V_{dd} values
 - A given V_{dd} can support a range of clock frequencies with a
 - $F \propto (V_{dd} - V_{th})^\chi / V_{dd}$ $\chi \in (1.0, 2.0)$
- Why DVFS saves power and energy?
 - Reduce V_{dd} to γV_{dd}
 - f_{max} reduces to roughly γf_{max}
 - Dynamic power reduces by roughly γ^3
 - Energy reduces by roughly γ^2


Voltage / Frequency

- **Dynamic Voltage Scaling**
 - Adjust V_{DD} and f according to workload




Dynamic Voltage Scaling (DVS)


- Stretch the execution by lowering the supply voltage
 - Quadratic power saving
 - Working only fast enough to meet deadlines!
- DVS algorithms can be implemented as HW or SW



Dynamic Voltage and Frequency Scaling (DVFS)




If you don't need to hustle,
relax and save power




- Through SW-HW cooperation
- SW automatically predicts future workloads (by OS and application software interaction)

How DVFS Works

- Suppose that a data sample comes every 1 ms
- Requires processing time of 250 μ s at 600MHz
- DVS: reduce voltage such that clock slows down to 150MHz



DVFS Roadmap




Safety and security

- Programs that perform important operations should run in trusted mode.
- Supervisor mode, memory management provide some level of trust but does not ensure that the code comes from a trusted source.
- Testing a digital signature for software before executing ensures that the software comes from a trusted source.

Trusted programs and root of trust

- Checking a digital signature for code requires a public key.
 - Public key must come from a trusted source.
- A trusted program may rely on another trusted program, etc.
 - New program's level of trustworthiness depends on the previous program's trustworthiness.
- Trust must eventually be traced back to a root of trust that provides known trust.
- Root of trust can be provided with a hardware key that cannot be altered.

Chain-of-trust



ARM Trusted Execution Environment

- Four-compartment model:
 - Normal world includes user and system modes.
 - Hypervisor mode supports virtual machines.
 - Trusted World partitions into secure and non-secure components.
 - SecurCore provides physically separate chips for additional protection.
- Trusted Execution Environment (TEE) provides secure execution mode.

Secure operation according to ARM

- Processor security mechanisms should limit system access.
- DMA should provide secure access for memory and I/O.
- Access to CPU from other parts of system should be limited.
- Trusted software should be able to identify and control processes that consume resources in the non-trusted world.
- JTAG etc. should not compromise security.
- I/O devices should support secure/non-secure operation.
- Platform should have a unique, non-modifiable key.
- Platform should provide secure trusted private key storage.
- Process and update should be secure and verifiable.
- Secure execution should ensure that control flow cannot be subverted.
- Platform should provide security primitives.

Smart cards

- Smart cards are used for financial, medical, other types of secure information.
 - Must operate at very low energy levels and very low cost.
- Runs from external power supply.
- Some code held in ROM.
- Permanent storage in non-volatile EEPROM.
- Cryptography engine provides low power operation, trusted functions.

