


Program design and analysis

- Software components.
- Representations of programs.
- Assembly and linking.

Software state machine

- State machine keeps internal state as a variable, changes state based on inputs.
- Uses:
 - control-dominated code;
 - reactive systems.

State machine example



C implementation

```
#define IDLE 0
#define SEATED 1
#define BELTED 2
#define BUZZER 3
switch (state) {
    case IDLE: if (seat) { state = SEATED; timer_on = TRUE; }
               break;
    case SEATED: if (belt) state = BELTED;
                else if (timer) state = BUZZER;
               break;
    ...
}
```


```

switch(state) { /*check the current state */
    case IDLE:
        if (seat){ state = SEATED; timer_on = TRUE; }
        /*default case is self-loop */
        break;
    case SEATED:
        if (belt) state = BELTED; /*won't hear the buzzer */
        else if (timer) state = BUZZER; /*didn't put on
belt in time */
        /*default case is self-loop */
        break;
    case BELTED:
        if (!seat) state = IDLE; /* person left */
        else if (!belt) state = SEATED; /* person still
in seat */
        break;
    case BUZZER:
        if (belt) state = BELTED; /*belt is on---turn off
buzzer */
        else if (!seat) state = IDLE; /* no one in seat--
turn off buzzer */
        break;
}

```


Signal processing and circular buffer

- Commonly used in signal processing:
 - new data constantly arrives;
 - each datum has a limited lifetime.




- Use a circular buffer to hold the data stream.

Circular buffer





Data stream



Circular buffer

Circular buffers

- Indexes locate currently used data, current input data:



Circular buffer in C

```
#define CMAX 6 /* filter order */
int circ[CMAX]; /* circular buffer */
int pos; /* position of current sample */


void circ_update(int xnew) {
    /* compute the new head value with wraparound; the pos pointer moves from 0 to CMAX-1 */
    pos = ((pos == CMAX-1) ? 0 : (pos+1));
    /* insert the new value at the new head */
    circ[pos] = xnew;
}
```

Circular buffer in C, cont'd.

```
void circ_init() {
    int i;
    for (i=0; i<CMAX; i++) /* set values to 0 */
        circ[i] = 0; pos=CMAX-1; /* start at tail so first element will be at 0 */
}

int circ_get(int i) {
    int ii;
    /* compute the buffer position */
    ii = (pos - i) % CMAX;
    return circ[ii]; /* return the value */
}
```

FIR filter




FIR filter update function

```
void circ_update(int xnew) {  
    /* add the new sample and push off the oldest one */  
    /* compute the new head value with wraparound; the pos pointer moves from  
    CMAX-1 down to 0 */  
    pos = ((pos == 0) ? CMAX-1 : (pos-1));  
    /* insert the new value at the new head */  
    circ[pos] = xnew;  
}
```

FIR filter using circular buffer

```
int fir(int xnew) {  
    /* given a new sample value, update the queue and      compute the filter  
    output */  
    int i;  
    int result; /* holds the filter output */  
    circ_update(xnew); /* put the new value in */  
    for (i=0, result=0; i<CMAX; i++)  
        result += b[i] * circ_get(i);  
    return result;  
}
```

IIR direct form type II filter



IIR filter in C

```
int iir2(int xnew) {  
    int i, aside, bside, result;  
    for (i=0, aside=0; i<ZMAX; i++)  
        aside += -a[i+1] * circ_get(i);  
    for (i=0, bside=0; i<ZMAX; i++)  
        bside += b[i+1] * circ_get(i);  
    result = b[0] * (xnew + aside) + bside;  
    circ_update(xnew); /* put the new value in */  
    return result;  
}
```

Array-based queue in C

```
#define Q_SIZE 5 /* your queue size may vary */
#define Q_MAX (Q_SIZE-1) /* this is the maximum index value into the array */
int q[Q_SIZE]; /* the array for our queue */
int head, tail; /* indexes for the current queue head and tail */
void queue_init() {
    /* initialize the queue data structure */
    head = 0;
    tail = 0;
}
```



Array based queue, cont'd.

```
void enqueue(int val) {  
    if (((tail+1) % Q_SIZE) == head) error("enqueue onto full queue",tail);  
    q[tail] = val;  
    /* update the tail */  
    if (tail == Q_MAX)  
        tail = 0;  
    else  
        tail++;  
}
```

Array based queue, cont'd.

```
int dequeue() {  
    int returnval;  
    if (head == tail) error("dequeue from empty queue",head);  
    returnval = q[head];  
    if (head == Q_MAX)  
        head = 0;  
    else  
        head++;  
    return returnval;  
}
```

Producer-consumer system



- Queues allow varying input and output rates.

Models of programs

- Source code is not a good representation for programs:
 - clumsy;
 - leaves much information implicit.
- Compilers derive intermediate representations to manipulate and optimize the program.

Data flow graph

- **DFG**: data flow graph.
- Does not represent control.
- Models basic block: code with no entry or exit.
- Describes the minimal ordering requirements on operations.

Single assignment form

$x = a + b;$

$y = c - d;$

$z = x * y;$

$y = b + d;$

original basic block

$x = a + b;$

$y = c - d;$

$z = x * y;$

$y1 = b + d;$

single assignment form

Data flow graph


$x = a + b;$

$y = c - d;$


$z = x * y;$

$y1 = b + d;$

single assignment form



DFGs and partial orders



Partial order:

- $a+b, c-d; b+d \times y$


Can do pairs of operations in any order.

Control-data flow graph

- **CDFG**: represents control and data.
- Uses data flow graphs as components.
- Two types of nodes:
 - decision;
 - data flow.

Data flow node


Encapsulates a data flow graph:



```
x = a + b;  
y = c + d
```

Write operations in basic block form for simplicity.


Control



Equivalent forms

CDFG example

```
if (cond1) bb1();  
else bb2();  
bb3();  
switch (test1) {  
  case c1: bb4(); break;  
  case c2: bb5(); break;  
  case c3: bb6(); break;  
}
```




for loop

```
for (i=0; i<N; i++)  
    loop_body();
```


for loop

```
i=0;  
while (i<N) {  
    loop_body(); i++; }  
  
equivalent
```



Assembly and linking

- Last steps in compilation:



Multiple-module programs

- Programs may be composed from several files.
- Addresses become more specific during processing:
 - **relative addresses** are measured relative to the start of a module;
 - **absolute addresses** are measured relative to the start of the CPU address space.

Assemblers

- Major tasks:
 - generate binary for symbolic instructions;
 - translate labels into addresses;
 - handle pseudo-ops (data, etc.).
- Generally one-to-one translation.
- Assembly labels:
 ORG 100
label1 ADR r4,c

Symbol table

	ADD r0,r1,r2
xx	ADD r3,r4,r5
	CMP r0,r3
yy	SUB r5,r6,r7

assembly code


xx	0x8
yy	0x10

symbol table

Symbol table generation

- Use program location counter (**PLC**) to determine address of each location.
- Scan program, keeping count of PLC.
- Addresses are generated at assembly time, not execution time.

Symbol table example



Two-pass assembly

- Pass 1:
 - generate symbol table
- Pass 2:
 - generate binary instructions

Relative address generation

- Some label values may not be known at assembly time.
- Labels within the module may be kept in relative form.
- Must keep track of external labels---can't generate full binary for instructions that use external labels.

Pseudo-operations

- Pseudo-ops do not generate instructions:
 - **ORG** sets program location.
 - **EQU** generates symbol table entry without advancing PLC.
 - **Data statements** define data blocks.

```
                ORG 100
label1  ADR r4,c
                LDR r0,[r4]
label2  ADR r4,d
                LDR r1,[r4]
label3  SUB r0,r0,r1
```

PLC = 116 $\longrightarrow$ label3

	ORG 100
label1	ADR r4,c
	LDR r0,[r4]
label2	ADR r4,d
	LDR r1,[r4]
label3	SUB r0,r0,r1

Code

label1	100
label2	108
label3	116


Symbol table


```
        ADD  r0,r1,r2
FOO     EQU  5
BAZ     SUB  r3,r4,#FOO
```

Linking


- Combines several object modules into a single executable module.
- Jobs:
 - put modules in order;
 - resolve labels across modules.

Externals and entry points



Module ordering

- Code modules must be placed in absolute positions in the memory space.
- **Load map** or linker flags control the order of modules.



Dynamic linking

- Some operating systems link modules dynamically at run time:
 - shares one copy of library among all executing programs;
 - allows programs to be updated with new versions of libraries.

Reentrancy

- Interrupting program with another call to the function does not change results.
 - Changing global variables compromises reentrancy.
- Recursive code:

```
int foo = 1;  
int task1() {  
    foo = foo + 1;  
    return foo;  
}
```


Program design and analysis

- Compilation flow.
- Basic statement translation.
- Basic optimizations.
- Interpreters and just-in-time compilers.

Compilation

- Compilation strategy (Wirth):
 - compilation = translation + optimization
- Compiler determines quality of code:
 - use of CPU resources;
 - memory access scheduling;
 - code size.

Basic compilation phases




Statement translation and optimization

- Source code is translated into intermediate form such as CDFG.
- CDFG is transformed/optimized.
- CDFG is translated into instructions with optimization decisions.
- Instructions are further optimized.

Arithmetic expressions


$a*b + 5*(c-d)$

expression



DFG

Arithmetic expressions, cont'd.



DFG

```
ADR r4,a
MOV r1,[r4]
ADR r4,b
MOV r2,[r4]
MUL r3,r1,r2

ADR r4,c
MOV r1,[r4]
ADR r4,d
MOV r5,[r4]
SUB r6,r4,r5
MUL r7,r6,#5
ADD r8,r7,r3
```

code

Compiled code for arithmetic expressions

ldr r2, [fp, #-16]

ldr r3, [fp, #-20]

mul r1, r3, r2 ; multiply

ldr r2, [fp, #-24]

ldr r3, [fp, #-28]

rsb r2, r3, r2 ; subtract

mov r3, r2

mov r3, r3, asl #2

add r3, r3, r2 ; add

add r3, r1, r3 ; add

str r3, [fp, #-32] ; assign


Control code generation

if ($a+b > 0$)


$x = 5$;

else

$x = 7$;



Control code generation, cont'd.



```
ADR r5,a
LDR r1,[r5]
ADR r5,b
LDR r2,[r5]
ADD r3,r1,r2
BLE label3
  LDR r3,#5
  ADR r5,x
  STR r3,[r5]
  B stmtent
label3 LDR r3,#7
  ADR r5,x
  STR r3,[r5]
stmtent ...
```

Compiled code for control

ldr r2, [fp, #-16]

ldr r3, [fp, #-20]

add r3, r2, r3

cmp r3, #0 ; test the branch
condition

ble .L3 ; branch to false block if <=

mov r3, #5 ; true block

str r3, [fp, #-32]

b .L4 ; go to end of if statement

.L3: ; false block

mov r3, #7


str r3, [fp, #-32]

.L4:

Procedure linkage

- Need code to:
 - call and return;
 - pass parameters and results.
- Parameters and returns are passed on stack.
 - Procedures with few parameters may use registers.
- Local variables are stored in the stack.

Procedure stacks



ARM procedure linkage

- APCS (ARM Procedure Call Standard):
 - r0-r3 pass parameters into procedure. Extra parameters are put on stack frame.
 - r0 holds return value.
 - r4-r7 hold register values.
 - r11 is frame pointer, r13 is stack pointer.
 - r10 holds limiting address on stack size to check for stack overflows.

```
int p1(int a, int b, int c, int d, int e) {  
    return a + e;  
}
```

```
mov    ip, sp          ; procedure entry  
stmfd  sp!, {fp, ip, lr, pc}  
sub    fp, ip, #4  
sub    sp, sp, #16  
str    r0, [fp, #-16]   ; put first four args on stack  
str    r1, [fp, #-20]  
str    r2, [fp, #-24]  
str    r3, [fp, #-28]  
ldr    r2, [fp, #-16]   ; load a  
ldr    r3, [fp, #4]     ; load e  
add    r3, r2, r3       ; compute a + e  
mov    r0, r3           ; put the result into r0 for return  
ldmea  fp, {fp, sp, pc} ; return
```

Compiled procedure call code

```
ldr r3, [fp, #-32] ; get e
str r3, [sp, #0] ; put into p1()'s stack frame
ldr r0, [fp, #-16] ; put a into r0
ldr r1, [fp, #-20] ; put b into r1
ldr r2, [fp, #-24] ; put c into r2
ldr r3, [fp, #-28] ; put d into r3
bl p1 ; call p1()
mov r3, r0 ; move return value into r3
str r3, [fp, #-36] ; store into y in stack frame
```


`y = p1(a, b, c, d, x);`

Data structures

- Different types of data structures use different data layouts.
- Some offsets into data structure can be computed at compile time, others must be computed at run time.


One-dimensional arrays

- C array name points to 0th element:



Two-dimensional arrays

- Row-major layout:




Structures

- Fields within structures are static offsets:

```
struct {  
    int field1;  
    char field2;  
} mystruct;
```

```
struct mystruct a, *aptr = &a;
```



Expression simplification

- Constant folding:
 - $8+1 = 9$
- Algebraic:
 - $a*b + a*c = a*(b+c)$
- Strength reduction:
 - $a*2 = a<<1$


Dead code elimination

- Dead code:

```
#define DEBUG 0
```

```
if (DEBUG) dbg(p1);
```

- Can be eliminated by analysis of control flow, constant folding.



Procedure inlining

- Eliminates procedure linkage overhead:

```
int foo(a,b,c) { return a + b - c;}
```

```
z = foo(w,x,y);
```



```
z = w + x - y;
```

Loop transformations

- Goals:
 - reduce loop overhead;
 - increase opportunities for pipelining;
 - improve memory system performance.

Loop unrolling

- Reduces loop overhead, enables some other optimizations.

```
for (i=0; i<4; i++)
```

```
    a[i] = b[i] * c[i];
```



```
for (i=0; i<2; i++) {
```

```
    a[i*2] = b[i*2] * c[i*2];
```

```
    a[i*2+1] = b[i*2+1] * c[i*2+1];
```

```
}
```

Loop fusion and distribution

- Fusion combines two loops into 1:

```
for (i=0; i<N; i++) a[i] = b[i] * 5;
```

```
for (j=0; j<N; j++) w[j] = c[j] * d[j];
```

```
⇒ for (i=0; i<N; i++) {  
    a[i] = b[i] * 5; w[i] = c[i] * d[i];  
}
```

- Distribution breaks one loop into two.
- Changes optimizations within loop body.

Register allocation

- Goals:
 - choose register to hold each variable;
 - determine lifespan of variable in the register.
- Basic case: within basic block.

Register lifetime graph

$w = a + b;$


$x = c + w;$

$y = c + d;$


$t=1$

$t=2$


$t=3$



```
w = a + b; /* statement 1 */  
x = c + d; /* statement 2 */  
y = x + e; /* statement 3 */  
z = a - b; /* statement 4 */
```



```
w = a + b; /*statement 1 */  
z = a - b; /* statement 29 */  
x = c + d; /*statement 39 */  
y = x + e; /*statement 49 */
```



Instruction scheduling

- Non-pipelined machines do not need instruction scheduling: any order of instructions that satisfies data dependencies runs equally fast.
- In pipelined machines, execution time of one instruction depends on the nearby instructions: **opcode**, **operands**.

Reservation table

- A reservation table relates instructions/time to CPU resources.


Time/instr	A	B
instr1	X	
instr2	X	X
instr3	X	
instr4		X

Software pipelining


- Schedules instructions across loop iterations.
- Reduces instruction latency in iteration i by inserting instructions from iteration $i+1$.

Instruction selection


- May be several ways to implement an operation or sequence of operations.
- Represent operations as graphs, match possible instruction sequences onto graph.




expression



MUL



ADD



MADD

templates

Using your compiler

- Understand various optimization levels (-O1, -O2, etc.)
- Look at mixed compiler/assembler output.
- Modifying compiler output requires care:
 - correctness;
 - loss of hand-tweaked code.

Interpreters and JIT compilers


- **Interpreter**: translates and executes program statements on-the-fly.
- **JIT compiler**: compiles small sections of code into instructions during program execution.
 - Eliminates some translation overhead.
 - Often requires more memory.

Program design and analysis

- Program-level performance analysis.
- Optimizing for:
 - Execution time.
 - Energy/power.
 - Program size.
- Program validation and testing.
- Safety and security.

Program-level performance analysis

- Need to understand performance in detail:
 - Real-time behavior, not just typical.
 - On complex platforms.
- Program performance $\neq$ CPU performance:
 - Pipeline, cache are windows into program.
 - We must analyze the entire program.



Complexities of program performance

- Varies with input data:
 - Different-length paths.
- Cache effects.
- Instruction-level performance variations:
 - Pipeline interlocks.
 - Fetch times.

How to measure program performance

- Simulate execution of the CPU.
 - Makes CPU state visible.
- Measure on real CPU using timer.
 - Requires modifying the program to control the timer.
- Measure on real CPU using logic analyzer.
 - Requires events visible on the pins.

Program performance metrics

- Average-case execution time.
 - Typically used in application programming.
- Worst-case execution time.
 - A component in deadline satisfaction.
- Best-case execution time.
 - Task-level interactions can cause best-case program behavior to result in worst-case system behavior.

Elements of program performance

- Basic program execution time formula:
 - execution time = program path + instruction timing
- Solving these problems independently helps simplify analysis.
 - Easier to separate on simpler CPUs.
- Accurate performance analysis requires:
 - Assembly/binary code.
 - Execution platform.


Data-dependent paths in an if statement

```
if (a || b) { /* T1 */  
    if ( c ) /* T2 */  
        x = r*s+t; /* A1 */  
    else y=r+s; /* A2 */  
    z = r+s+u; /* A3 */  
}  
else {  
    if ( c ) /* T3 */  
        y = r-t; /* A4 */  
}
```

a	b	c	path
0	0	0	T1=F, T3=F: no assignments
0	0	1	T1=F, T3=T: A4
0	1	0	T1=T, T2=F: A2, A3
0	1	1	T1=T, T2=T: A1, A3
1	0	0	T1=T, T2=F: A2, A3
1	0	1	T1=T, T2=T: A1, A3
1	1	0	T1=T, T2=F: A2, A3
1	1	1	T1=T, T2=T: A1, A3


Paths in a loop

```
for (i=0, f=0; i<N; i++)  
    f = f + c[i] * x[i];
```



Paths in a loop with conditions

```
for (i=0, f=0; i<N; i++) {  
    if (x[i] > 0)  
        f = f + c[i] * x[i];  
}
```



executed 1 time

executed $N+1$ times

executed N times

executed $[0, N]$ times

executed N times

Instruction timing

- Not all instructions take the same amount of time.
 - Multi-cycle instructions.
 - Fetches.
- Execution times of instructions are not independent.
 - Pipeline interlocks.
 - Cache effects.
- Execution times may vary with operand value.
 - Floating-point operations.
 - Some multi-cycle integer operations.

Caching effects

```
for (i = 0, f = 0; i < N; i++)  
    f = f + c[i] * x[i];
```

$$t_{loop} = 2N + \frac{N}{L}t_{miss} + N\left(1 - \frac{1}{L}\right)t_{hit}$$

line 0	Word 0	Word 1	Word 2	Word 3
line 1	Word 4	Word 5	Word 6	Word 7

- First access to a cache line causes a miss and prefetch.
- Later accesses to that line result in cache hits.

Mesaurement-driven performance analysis

- Not so easy as it sounds:
 - Must actually have access to the CPU.
 - Must know data inputs that give worst/best case performance.
 - Must make state visible.
- Still an important method for performance analysis.

Feeding the program

- Need to know the desired input values.
- May need to write software scaffolding to generate the input values.
- Software scaffolding may also need to examine outputs to generate feedback-driven inputs.

Trace-driven measurement

- Trace-driven:
 - Instrument the program.
 - Save information about the path.
- Requires modifying the program.
- Trace files are large.
- Widely used for cache analysis.

Physical measurement

- In-circuit emulator allows tracing.
 - Affects execution timing.
- Logic analyzer can measure behavior at pins.
 - Address bus can be analyzed to look for events.
 - Code can be modified to make events visible.
- Particularly important for real-world input streams.

CPU simulation

- Some simulators are less accurate.
- Cycle-accurate simulator provides accurate clock-cycle timing.
 - Simulator models CPU internals.
 - Simulator writer must know how CPU works.

SimpleScalar FIR filter simulation

```
int x[N] = {8, 17, ... };  
int c[N] = {1, 2, ... };  
main() {  
    int i, k, f;  
    for (k=0; k<COUNT; k++)  
        for (i=0; i<N; i++)  
            f += c[i]*x[i];  
}
```

COUNT	total sim cycles	sim cycles per filter execution
100	25854	259
1,000	155759	156
1,0000	1451840	145

Performance optimization motivation

- Embedded systems must often meet deadlines.
 - Faster may not be fast enough.
- Need to be able to analyze execution time.
 - Worst-case, not typical.
- Need techniques for reliably improving execution time.

Programs and performance analysis


- Best results come from analyzing optimized instructions, not high-level language code:
 - non-obvious translations of HLL statements into instructions;
 - code may move;
 - cache effects are hard to predict.

Loop optimizations

- Loops are good targets for optimization.
- Basic loop optimizations:
 - code motion;
 - induction-variable elimination;
 - strength reduction ($x*2 \rightarrow x \ll 1$).

Code motion

```
for (i=0; i<N*M; i++)  
    z[i] = a[i] + b[i];
```



Induction variable elimination

- **Induction variable**: loop index.
- Consider loop:
 for (i=0; i<N; i++)
 for (j=0; j<M; j++)
 z[i,j] = b[i,j];
- Rather than recompute $i*M+j$ for each array in each iteration, share induction variable between arrays, increment at end of loop body.


```
for (i = 0; i < N; i++)  
    for (j = 0; j < M; j++) {  
        zbinduct = i*M + j;  
        *(zptr + zbinduct) = *(bptr + zbinduct);  
  
zbinduct = 0;  
for (i = 0; i < N; i++) {  
    for (j = 0; j < M; j++) {  
        *(zptr + zbinduct) = *(bptr + zbinduct);  
        zbinduct++;  
    }  
}
```


Cache analysis

- **Loop nest**: set of loops, one inside other.
- **Perfect loop nest**: no conditionals in nest.
- Because loops use large quantities of data, cache conflicts are common.

Array conflicts in cache

```
for (j = 0; j < M; j++)  
  for (i = 0; i < N; i++)  
    a[j][i] = b[j][i] * c;
```



Array conflicts, cont'd.

- Array elements conflict because they are in the same line, even if not mapped to same location.
- Solutions:
 - move one array;
 - pad array to change alignment.

Array padding

- Add array elements to change mapping into cache:

	a[1,1]	a[1,2]

before

a[0,0]	a[0,1]	a[0,2]	a[0,2]
a[1,0]	a[1,1]	a[1,2]	a[1,2]


after

Loop tiling


- Breaks one loop into a nest of loops.
- Changes order of accesses within array.
 - Changes cache behavior.

Loop tiling example

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        z[i][j] = x[i] * y[j];  
    }  
}
```



```
for (j=0; j < N; j += TILE) {  
    for (i=0; i<N; i++) {  
        for (jj=0; jj<TILE; jj++) {  
            z[i][j + jj] = x[i] * y[j + jj];  
        }  
    }  
}
```



Performance optimization hints


- Use registers efficiently.
 - Group accesses to a value together so that the value can be brought into a register and kept there.
- Use page mode memory accesses.
 - Page mode reads and writes eliminate one step in the memory access. You can increase the use of page mode by rearranging your variables so that more can be referenced contiguously.
- Analyze cache behavior:
 - instruction conflicts can be handled by rewriting code, rescheduling;
 - conflicting scalar data can easily be moved;
 - conflicting array data can be moved, padded.

Energy/power optimization

- **Energy**: ability to do work.
 - Most important in battery-powered systems.
- **Power**: energy per unit time.
 - Important even in wall-plug systems---power becomes heat.

Measuring energy consumption

- Execute a small loop, measure current:





Sources of energy consumption

- Relative energy per operation (Catthoor et al):
 - memory transfer: 33
 - external I/O: 10
 - SRAM write: 9
 - SRAM read: 4.4
 - multiply: 3.6
 - add: 1

Cache behavior is important

- Energy consumption has a sweet spot as cache size changes:
 - **cache too small**: program thrashes, burning energy on external memory accesses;
 - **cache too large**: cache itself burns too much power.

Cache sweet spot



[Li98] © 1998 IEEE

Optimizing for energy

- First-order optimization:
 - high performance = low energy.
- Not many instructions trade speed for energy.

Optimizing for energy, cont'd.

- Use registers efficiently.
- Identify and eliminate cache conflicts.
- Moderate loop unrolling eliminates some loop overhead instructions.
 - When the loop is unrolled too much, the power increases owing to the lower hit rates of straight-line code.
- Eliminate pipeline stalls.
 - Software pipelining
- Inlining procedures may help: reduces linkage, but may increase cache thrashing.

Efficient loops

- General rules:
 - Don't use function calls.
 - Keep loop body small to enable local repeat (only forward branches).
 - Use unsigned integer for loop counter.
 - Use $\leq$ to test loop counter.
 - Make use of compiler---global optimization, software pipelining.

Single-instruction repeat loop example

STM #4000h,AR2

; load pointer to source

STM #100h,AR3

; load pointer to destination

RPT #(1024-1)

MVDD *AR2+,*AR3+

; move

Optimizing for program size

- Goal:
 - reduce hardware cost of memory;
 - reduce power consumption of memory units.
- Two opportunities:
 - data;
 - instructions.

Data size minimization

- Reuse constants, variables, data buffers in different parts of code.
 - Requires careful verification of correctness.
- Generate data using instructions.

Reducing code size

- Avoid function inlining.
- Choose CPU with compact instructions.
- Use specialized instructions where possible.

Program validation and testing

- But does it work?
- Concentrate here on functional verification.
- Major testing strategies:
 - Black box doesn't look at the source code.
 - Clear box (white box) does look at the source code.

Clear-box testing

- Examine the source code to determine whether it works:
 - Can you actually exercise a path?
 - Do you get the value you expect along a path?
- Testing procedure:
 - **Controllability**: provide program with inputs.
 - Execute.
 - **Observability**: examine outputs.

Controlling and observing programs

```
firout = 0.0;
for (j=curr, k=0; j<N; j++, k++)
    firout += buff[j] * c[k];
for (j=0; j<curr; j++, k++)
    firout += buff[j] * c[k];
if (firout > 100.0) firout = 100.0;
if (firout < -100.0) firout = -100.0;
```


- **Controllability:**
 - Must fill circular buffer with desired N values.
 - Other code governs how we access the buffer.
- **Observability:**
 - Want to examine firout before limit testing.

Execution paths and testing

- Paths are important in functional testing as well as performance analysis.
- In general, an exponential number of paths through the program.
 - Show that some paths dominate others.
 - Heuristically limit paths.


Choosing the paths to test

- Possible criteria:
 - Execute every statement at least once.
 - Execute every branch direction at least once.
- Equivalent for structured programs.
- Not true for gotos.



Basis paths

- Approximate CDFG with undirected graph.
- Undirected graphs have basis paths:
 - All paths are linear combinations of basis paths.



Graph

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>
<i>a</i>	0	0	1	0	0
<i>b</i>	0	0	1	0	1
<i>c</i>	1	1	0	1	0
<i>d</i>	0	0	1	0	1
<i>e</i>	0	1	0	1	0

Incidence matrix


<i>a</i>	1	0	0	0	0
<i>b</i>	0	1	0	0	0
<i>c</i>	0	0	1	0	0
<i>d</i>	0	0	0	1	0
<i>e</i>	0	0	0	0	1

Basis set

Cyclomatic complexity

- Cyclomatic complexity is a bound on the size of basis sets:

- $e = \# \text{ edges}$
- $n = \# \text{ nodes}$
- $p = \text{number of graph components}$
- $M = e - n + 2p.$



Branch testing


- Heuristic for testing branches.
 - Exercise true and false branches of conditional.
 - Exercise every simple condition at least once.

Branch testing example



- Correct:
 - `if (a || (b >= c)) { printf("OK\n"); }`
- Incorrect:
 - `if (a && (b >= c)) { printf("OK\n"); }`
- Test:
 - `a = F`
 - `(b >= c) = T`
- Example:
 - Correct: `[0 || (3 >= 2)] = T`
 - Incorrect: `[0 && (3 >= 2)] = F`

Domain testing

- Heuristic test for linear inequalities.
- Test on each side + boundary of inequality.



Correct test



Incorrect tests

Def-use pairs

- Variable def-use:
 - Def when value is assigned (defined).
 - Use when used on right-hand side.
- Exercise each def-use pair.
 - Requires testing correct path.

```
a = mypointer;  
if (c > 5){  
    while (a->field1 != val1)  
        a = a->next;  
}  
if (a->field2 == val2)  
    someproc(a,b);
```

The diagram illustrates a control flow graph with the following nodes and edges:

- Node 1:** `a = mypointer;` (Definition of `a`)
- Node 2:** `if (c > 5){` (Control flow from Node 1)
- Node 3:** `while (a->field1 != val1)` (Control flow from Node 2)
- Node 4:** `a = a->next;` (Definition of `a`, part of a loop from Node 3 back to Node 3)
- Node 5:** `}` (Control flow from Node 3 to Node 5)
- Node 6:** `if (a->field2 == val2)` (Control flow from Node 5)
- Node 7:** `someproc(a,b);` (Control flow from Node 6)

Arrows indicate the flow of execution between these statements, highlighting the paths where the variable `a` is defined and then used.

Loop testing

- Loops need specialized tests to be tested efficiently.
- Heuristic testing strategy:
 - Skip loop entirely.
 - One loop iteration.
 - Two loop iterations.
 - # iterations much below max.
 - $n-1$, n , $n+1$ iterations where n is max.

Black-box testing

- Complements clear-box testing.
 - May require a large number of tests.
- Tests software in different ways.

Black-box test vectors

- Random tests.
 - May weight distribution based on software specification.
- Regression tests.
 - Tests of previous versions, bugs, etc.
 - May be clear-box tests of previous versions.

How much testing is enough?

- Exhaustive testing is impractical.
- One important measure of test quality---bugs escaping into field.
- Good organizations can test software to give very low field bug report rates.
- Error injection measures test quality:
 - Add known bugs.
 - Run your tests.
 - Determine % injected bugs that are caught.

Security-related bugs

- Buffer overflows provide attackers with attack vectors.
- Failure to initialize buffers can leak information to attackers.


Code certificates

- Provide assurance that code comes from a trusted source:
 - Certification authority provides certificate service.
- Certificate:
 - Certification authority is given code/data.
 - Generates certificate with source for data and associated encryption date.
 - May include expiration date.
- Recipient checks source identifier before decrypting and using code/data.


Code signing

- Combines signing certificate with digital signature.
 - Hash code to provide a digest.
 - Signature includes original program, encrypted digest, public key, signing certificate.

Code signing process



Checking signed code



Passwords

- Passwords should be stored in encrypted form.