

# Program Design and Analysis

# 5

---

## CHAPTER POINTS

- Some useful components for embedded software.
- Models of programs, such as data flow and control flow graphs.
- An introduction to compilation methods.
- Analyzing and optimizing programs for performance, size, and power consumption.
- How to test programs to verify their accuracy.
- Design examples: software modem and digital still camera (DSC).

---

## 5.1 Introduction

In this chapter, we study in detail the process of creating programs for embedded processors. The creation of embedded programs is at the heart of embedded system design. If you are reading this book, you almost certainly know how to write code, but designing and implementing embedded programs is different and more challenging than writing typical workstation or PC programs. Embedded code must not only provide rich functionality, but also often run at a required rate to meet system deadlines, fit into the allowed amount of memory, and meet power consumption requirements. Designing code that simultaneously meets multiple design constraints is a considerable challenge, but luckily there are techniques and tools that we can use to help us through the design process. Ensuring that the program works is also a challenge, but once again methods and tools come to our aid.

Throughout the discussion we concentrate on high-level programming languages, specifically C. High-level languages were once shunned as too inefficient for embedded microcontrollers, but better compilers, more compiler friendly architectures, and faster processors and memory have made high-level language programs common. Some sections of a program may still need to be written in assembly language if the compiler doesn't provide sufficiently good results, but even when coding in assembly language, it is often helpful to think about the program's functionality in high-level form. Many of the analysis and optimization techniques that we study in this chapter are equally applicable to programs that are written in assembly language.

The next section talks about some software components that are commonly used in embedded software. Section 5.3 introduces the control/data flow graph as a model for high-level language programs (which can also be applied to programs originally written in assembly language). Section 5.4 reviews the assembly and linking process, and Section 5.5 introduces some compilation techniques. Section 5.6 introduces methods for analyzing the performance of programs. We talk about optimization techniques that are specific to embedded computing in the next three sections: performance in Section 5.7, energy consumption in Section 5.8, and size in Section 5.9. In Section 5.10, we discuss techniques for ensuring that the programs you write are correct. In Section 5.11, we consider the related problem of program design for safety and security. We close with two design examples: a software modem in Section 5.12 and a DSC in Section 5.13.

---

## 5.2 Components for embedded programs

In this section, we consider code for three structures or components that are commonly used in embedded software: the state machine, circular buffer, and queue. State machines are well suited to **reactive systems** such as user interfaces; circular buffers and queues are useful in digital signal processing.

### 5.2.1 State machines

#### State machine style

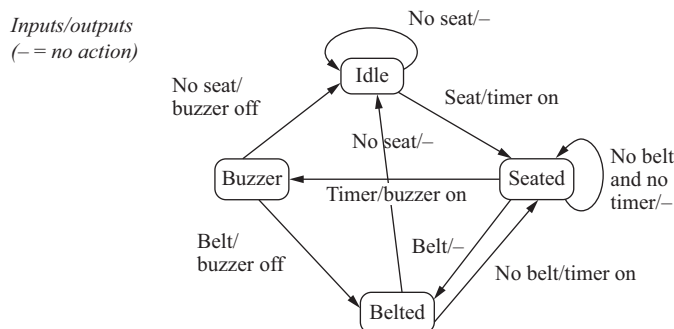
When inputs appear intermittently rather than as periodic samples, it is often convenient to think of the system as reacting to those inputs. The reaction of most systems can be characterized in terms of the input received and the current state of the system. This naturally leads to a **finite-state machine** style of describing the reactive system's behavior. Moreover, if the behavior is specified in that way, it is natural to write the program implementing that behavior in a state machine style. The state machine style of programming is also an efficient implementation of such computations. Finite-state machines are usually first encountered in the context of hardware design.

Programming Example 5.1 shows how to write a finite-state machine in a high-level programming language.

---

#### Programming Example 5.1: A State Machine in C

The behavior we want to implement is a simple seat belt controller [Chi94]. The controller's job is to turn on a buzzer if a person sits in a seat and does not fasten the seat belt within a fixed amount of time. This system has three inputs and one output. The inputs are a sensor for the seat to know when a person has sat down, a seat belt sensor that tells when the belt is fastened, and a timer that goes off when the required time interval has elapsed. The output is the buzzer. Here is a state diagram that describes the seat belt controller's behavior:



The idle state is in force when there is no person in the seat. When the person sits down, the machine goes into the seated state and turns on the timer. If the timer goes off before the seat belt is fastened, the machine goes into the buzzer state. If the seat belt goes on first, it enters the belted state. When the person leaves the seat, the machine goes back to idle.

To write this behavior in C, we will assume that we have loaded the current values of all three inputs (seat, belt, timer) into variables and will similarly hold the outputs in variables temporarily (timer\_on, buzzer\_on). We will use a variable named state to hold the current state of the machine and a switch statement to determine which action to take in each state. Here is the code:

```

#define IDLE 0
#define SEATED 1
#define BELTED 2
#define BUZZER 3

switch(state) { /*check the current state */
    case IDLE:
        if (seat){ state = SEATED; timer_on = TRUE; }
        /*default case is self-loop */
        break;
    case SEATED:
        if (belt) state = BELTED; /*won't hear the buzzer */
        else if (timer) state = BUZZER; /*didn't put on
belt in time */
        /*default case is self-loop */
        break;
    case BELTED:
        if (!seat) state = IDLE; /* person left */
        else if (!belt) state = SEATED; /* person still
in seat */
        break;
    case BUZZER:
        if (belt) state = BELTED; /*belt is on---turn off
buzzer */
        else if (!seat) state = IDLE; /* no one in seat--
turn off buzzer */
        break;
}

```

This code takes advantage of the fact that the state will remain the same unless explicitly changed; this makes self-loops back to the same state easy to implement. This state machine may be executed forever in a `while (TRUE)` loop or periodically called by some other code. In either case, the code must be executed regularly so that it can check on the current value of the inputs, and if necessary, go into a new state.

---

5.2.2 Circular buffers and stream-oriented programming

Data stream style

The data stream style makes sense for data that come in regularly and must be processed on the fly. The finite impulse response (FIR) filter of Application Example 2.1 is a classic example of stream-oriented processing. For each sample, the filter must emit one output that depends on the values of the last  $n$  inputs. In a typical workstation application, we would process the samples over a given interval by reading them all in from a file, and then, computing the results all at once in a batch process. In an embedded system, we must not only emit outputs in real time, but we must also do so using a minimum amount of memory.

Circular buffer

The **circular buffer** is a data structure that lets us handle streaming data in an efficient way. Fig. 5.1 illustrates how a circular buffer stores a subset of the data stream. At each point in time, the algorithm needs a subset of the data stream that forms a window into the stream. The window slides with time as we throw out old values that are no longer needed and add new values. Because the size of the window does not




FIGURE 5.1

A circular buffer.

change, we can use a fixed-sized buffer to hold the current data. To avoid constantly copying data within the buffer, we will move the head of the buffer in time. The buffer points to the location at which the next sample will be placed; every time we add a sample, we automatically overwrite the oldest sample, which is the one that needs to be thrown out. When the pointer reaches the end of the buffer, it wraps around to the top.

#### Instruction set support

Many digital signal processors provide addressing modes to support circular buffers. For example, the C55x [Tex04] provides five circular buffer start address registers (their names start with BSA). These registers allow circular buffers to be placed without alignment constraints.

#### High-level language implementation

In the absence of specialized instructions, we can write our own C code for a circular buffer. This code also helps us to understand the operation of the buffer. Programming Example 5.2 provides an efficient implementation of a circular buffer.

---

### Programming Example 5.2: A Circular Buffer in C

Once we build a circular buffer, we can use it in a variety of ways. We will use an array as the buffer:

```
#define CMAX 6 /*filter order */

int circ[CMAX]; /*circular buffer */
int pos; /*position of current sample */
```

The variable `pos` holds the position of the current sample. As we add new values to the buffer, this variable moves.

Here is the function that adds a new value to the buffer:

```
void circ_update(int xnew) {
    /*add the new sample and push off the oldest one */

    /*compute the new head value with wraparound; the pos
    pointer moves from 0 to CMAX-1 */
    pos = ((pos == CMAX-1) ? 0 : (pos+1));
    /*insert the new value at the new head */
    circ[pos] = xnew;
}
```

The assignment to `pos` takes care of wraparound; when `pos` hits the end of the array, it returns to zero. We then put the new value into the buffer at the new position. This overwrites the old value that was there. Note that as we go to higher index values in the array, we march through the older values.

We can now write an initialization function, which sets the buffer values to zero. More importantly, it sets `pos` to the initial value. For ease of debugging, we want the first data

element to go into `circ[0]`. To do this, we set `pos` to the end of the array so that it is set to zero before the first element is added:

```
void circ_init() {
    int i;

    for (i=0; i<CMAX; i++) /* set values to 0 */
        circ[i] = 0;
    pos=CMAX-1; /*start at tail so first element will be at 0 */
}
```


We can also make use of a function to get the  $i^{\text{th}}$  value of the buffer. This function has to translate the index in temporal order, with zero being the newest value, to its position in the array:

```
int circ_get(int i) {
    /*get the ith value from the circular buffer */
    int ii;
    /*compute the buffer position */
    ii = (pos - i) % CMAX;
    /*return the value */
    return circ[ii];
}
```

### Signal flow graph

We can now write C code for a digital filter. To help us to understand the filter algorithm, we can introduce a widely used representation for filter functions.

The FIR filter is only one type of digital filter. We can represent many different filtering structures using a **signal flow graph**, as shown in Fig. 5.2. The filter operates at a sample rate, with inputs arriving and outputs generated at the sample rate. The inputs  $x[n]$  and  $y[n]$  are sequences indexed by  $n$ , which corresponds to the sequence of samples. Nodes in the graph can represent either arithmetic operators or delay operators. The  $+$  node adds its two inputs and produces the output  $y[n]$ . The box labeled  $z^{-1}$  is a delay operator. The  $z$  notation comes from the  $z$  transform that is used in digital signal processing; the  $-1$  superscript means that the operation performs a time delay of one sample period. The edge from the delay operator to the



**FIGURE 5.2**

A signal flow graph.

## Filters and buffering

addition operator is labeled with  $b_1$ , meaning that the output of the delay operator is multiplied by  $b_1$ .

The code to produce one FIR filter output looks like this:

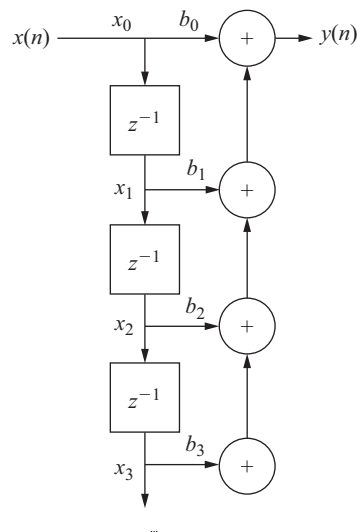
```
for (i=0, y=0.0; i<N; i++)
    y += x[i] * b[i];
```

However, the filter takes in a new sample in every sample period. The new input becomes  $x_1$ , the old  $x_1$  becomes  $x_2$ , and so on.  $x_0$  is stored directly in the circular buffer, but must be multiplied by  $b_0$  before being added to the output sum. Early digital filters were built-in hardware, where we could build a shift register to perform this operation. If we used an analogous operation in software, we would move every value in the filter in every sample period. We can avoid this with a circular buffer, moving the head without moving the data elements.

The next example uses our circular buffer class to build a FIR filter.

### Programming Example 5.3: An FIR Filter in C

Here is a signal flow graph for an FIR filter:



The delay elements running vertically hold the input samples, with the most recent sample at the top and the oldest one at the bottom. Unfortunately, the signal flow graph doesn't explicitly label all of the values that we use as inputs to operations, so the figure also shows the values ( $x_i$ ) that we need to operate on in our FIR loop.

When we compute the filter function, we want to match the  $b_i$ 's and  $x_i$ 's. We will use our circular buffer for the  $x$ 's, which change over time. We will use a standard array for the  $b$ 's that don't change. In order for the filter function to be able to use the same  $i$  value for both sets of data, we need to put the  $x$  data in the proper order. We can put the  $b$  data in a standard

array, with  $b_0$  being the first element. When we add a new  $x$  value, it becomes  $x_0$  and replaces the oldest data value in the buffer. This means that the buffer head moves from higher to lower values, and not from lower to higher as we might expect.

Here is the modified version of `circ_update()` that puts a new sample into the buffer in the desired order:

```
void circ_update(int xnew) {
    /*add the new sample and push off the oldest one */

    /*compute the new head value with wraparound; the pos
    pointer moves from CMAX-1 down to 0 */
    pos = ((pos == 0) ? CMAX-1 : (pos-1));
    /*insert the new value at the new head */
    circ[pos] = xnew;
}
```

We also need to change `circ_init()` to set `pos = 0` initially. We don't need to change `circ_get()`.

Given these functions, the filter itself is simple. Here is our code for the FIR filter function:

```
int fir(int xnew) {
    /*given a new sample value, update the queue and compute the
    filter output */
    int i;

    int result; /*holds the filter output */

    circ_update(xnew); /*put the new value in */
    for (i=0, result=0; i<CMAX; i++) /* compute the filter
    function */
        result += b[i] * circ_get(i);
    return result;
}
```

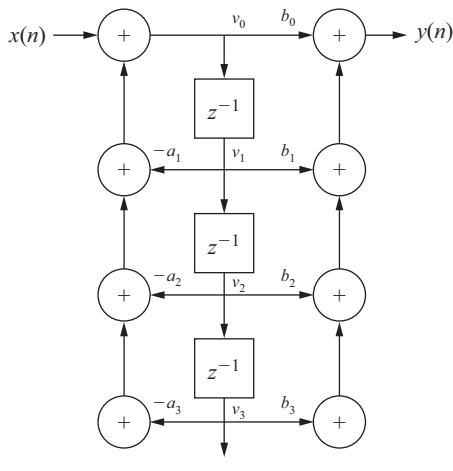
---

There is only one major structure for FIR filters, but several possible structures for infinite impulse response (IIR) filters, depending on the application requirements. One of the important reasons for so many different IIR forms is numerical properties; depending on the filter structure and coefficients, one structure may give significantly less numerical noise than another. However, numerical noise is beyond the scope of our discussion, so let's concentrate on one form of the IIR filter that highlights buffering issues. The next example looks at one form of the IIR filter.



### Programming Example 5.4: A Direct Form II IIR Filter Class in C

Here is what is known as the direct form II of an IIR filter:



This structure is designed to minimize the amount of buffer space required. Other forms of the IIR filter have other advantages, but require more storage. We will store the  $v_i$  values as in the FIR filter. In this case,  $v_0$  does not represent the input, but rather the left-hand sum. However,  $v_0$  is stored before multiplication by  $b_0$  so that we can move  $v_0$  to  $v_1$  in the following sample period.

We can use the same `circ_update()` and `circ_get()` functions that we used for the FIR filter. We need two coefficient arrays: one for  $a$ s and one for  $b$ s; as with the FIR filter, we can use standard C arrays for the coefficients because they don't change over time. Here is the IIR filter function:

```
int iir2(int xnew) {
    /*given a new sample value, update the queue and compute
    the filter output */
    int i, aside, bside, result;

    for (i=0, aside=0; i<ZMAX; i++)
        aside += -a[i+1] *circ_get(i);
    for (i=0, bside=0; i<ZMAX; i++)
        bside += b[i+1] *circ_get(i);
    result = b[0] *(xnew + aside) + bside;
    circ_update(xnew); /*put the new value in */
    return result;
}
```

### 5.2.3 Queues and producer/consumer systems

Queues are also used in signal processing and event processing. Queues are used whenever data may arrive and depart at somewhat unpredictable times, or when variable amounts of data may arrive. A queue is often referred to as an **elastic buffer**. We saw how to use elastic buffers for I/O in [Chapter 3](#).

One way to build a queue is with a linked list. This approach allows the queue to grow to an arbitrary size. However, in many applications, we are unwilling to pay the price of dynamically allocating memory. Another way to design the queue is to use an array to hold all of the data. Although some writers use both circular buffer and queue to mean the same thing, we use the term *circular buffer* to refer to a buffer that always has a fixed number of data elements, while a *queue* may have varying numbers of elements.

Programming Example 5.5 gives the C code for a queue that is built from an array.

---

#### Programming Example 5.5: An Array-Based Queue

The first step in designing the queue is to declare the array that we will use for the buffer:

```
#define Q_SIZE 5 /*your queue size may vary */
#define Q_MAX (Q_SIZE-1) /*this is the maximum index value into the
array */

int q[Q_SIZE]; /*the array for our queue */
int head, tail; /*indexes for the current queue head and tail */
```

The variables `head` and `tail` keep track of the two ends of the queue.

Here is the initialization code for the queue:

```
void queue_init() {

    /*initialize the queue data structure */

    head = 0;
    tail = 0;
}
```

We initialize the head and tail to the same position. As we add a value to the tail of the queue, we increment `tail`. Similarly, when we remove a value from the head, we increment `head`. The value of `head` is always equal to the location of the first element of the queue (except when the queue is empty). The value of `tail`, in contrast, points to the location in which the next queue entry will go. When we reach the end of the array, we must wrap around these values; for example, when we add a value into the last element of `q`, the new value of `tail` becomes the zeroth entry of the array.

We need to check for two error conditions: removing from an empty queue and adding to a full queue. In the first case, we know the queue is empty if `head == tail`. In the second case, we know the queue is full if incrementing `tail` will cause it to equal `head`. Testing for fullness, however, is a little harder because we must worry about wraparound.

Here is the code for adding an element to the tail of the queue, which is known as **enqueueing**:

```
void enqueue(int val) {
    /*check for a full queue */
    if (((tail+1) % Q_SIZE) == head) error("enqueue onto full
queue",tail);
    /*add val to the tail of the queue */
    q[tail] = val;
    /*update the tail */
    if (tail == Q_MAX)
        tail = 0;
    else
        tail++;
}
```


And here is the code for removing an element from the head of the queue, known as **dequeueing**:

```
int dequeue() {
    int returnval; /*use this to remember the value that you will
return */
    /*check for an empty queue */
    if (head == tail) error("dequeue from empty queue",head);
    /*remove from the head of the queue */
    returnval = q[head];
    /*update head */
    if (head == Q_MAX)
        head = 0;
    else
        head++;
    /*return the value */
    return returnval;
}
```

Digital filters always take in the same amount of data in each time period. Many systems, even signal processing systems, don't fit that mold. Rather, they may take in varying amounts of data over time and produce varying amounts. When several of these systems operate in a chain, the variable-rate output of one stage becomes the variable-rate input of another stage.

#### Producer/consumer

Fig. 5.3 shows a block diagram of a simple **producer/consumer system**. p1 and p2 are the two blocks that perform algorithmic processing. The data are fed to them by queues that act as elastic buffers. The queues modify the flow of control in the system and store data. If, for example, p2 runs ahead of p1, it will eventually run out of data in its q12 input queue. At that point, the queue will return an empty signal to p2. At this point, p2 should stop working until more data are available. This sort of complex control is easier to implement in a multitasking environment, as we will see in [Chapter 6](#),

**FIGURE 5.3**

A producer/consumer system.

#### Data structures in queues

but it is also possible to make effective use of queues in programs that are structured as nested procedures.

The queues in a producer/consumer may hold either uniform-sized data elements or variable-sized data elements. In some cases, the consumer needs to know how many of a given type of data element are associated together. The queue can be structured to hold a complex data type. Alternatively, the data structure can be stored as bytes or integers in the queue with, for example, the first integer holding the number of successive data elements.

## 5.3 Models of programs

In this section, we develop models for programs that are more general than source code. Why not use the source code directly? First, there are many different types of source code, such as assembly languages and C code, but we can use a single model to describe all of them. Once we have such a model, we can perform many useful analyses on the model more easily than we could on the source code.

Our fundamental model for programs is the control/data flow graph (CDFG). We can also model hardware behavior with the CDFG. As the name implies, the CDFG has constructs that model both data operations (arithmetic and other computations) and control operations (conditionals). Part of the power of the CDFG comes from its combination of control and data constructs. To understand the CDFG, we start with pure data descriptions, and then, extend the model to control.

### 5.3.1 Data flow graphs

A **data flow graph** is a model of a program with no conditionals. In a high-level programming language, a code segment with no conditionals—more precisely, with only one entry and exit point—is known as a basic block. Fig. 5.4 shows a simple basic block. As the C code is executed, we would enter this basic block at the beginning and execute all of the statements.

Before we are able to draw the data flow graph for this code, we need to modify it slightly. There are two assignments to the variable  $x$ ; it appears twice on the left side of an assignment. We need to rewrite the code in **single-assignment form**, in which a variable appears only once on the left side. Because our specification is C code, we

```

w = a + b;
x = a - c;
y = x + d;
x = a + c;
z = y + e;

```

**FIGURE 5.4**

A basic block in C.

```

w = a + b;
x1 = a - c;
y = x1 + d;
x2 = a + c;
z = y + e;

```

**FIGURE 5.5**

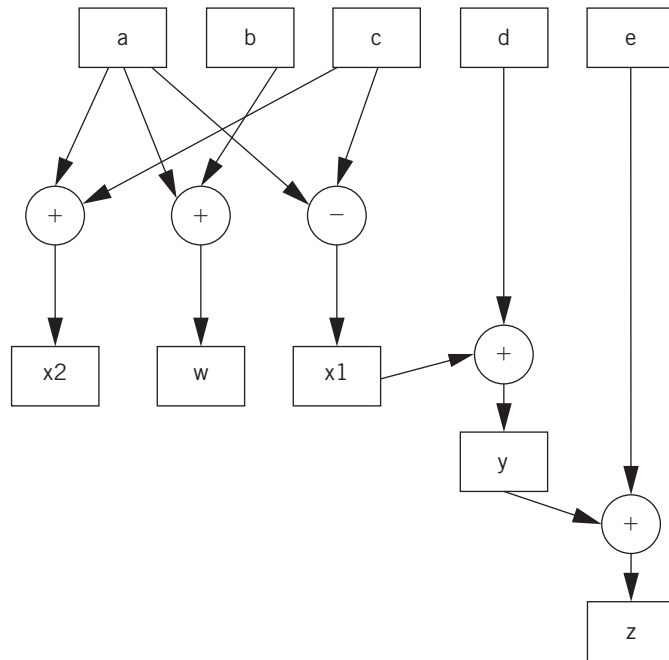
The basic block in single-assignment form.

assume that the statements are executed sequentially, so that any use of a variable refers to its latest assigned value. In this case,  $x$  is not reused in this block (it is presumably used elsewhere), so we must eliminate the multiple assignments to  $x$ . The result is shown in [Fig. 5.5](#), where we have used the names  $x1$  and  $x2$  to distinguish the separate uses of  $x$ .

The single-assignment form is important because it allows us to identify a unique location in the code where each named location is computed. As an introduction to the data flow graph, we use two types of nodes in the graph: round nodes denote operators and square nodes represent values. The value nodes may be either inputs to the basic block, such as  $a$  and  $b$ , or variables that are assigned to within the block, such as  $w$  and  $x1$ . The data flow graph for our single-assignment code is shown in [Fig. 5.6](#). The single-assignment form means that the data flow graph is acyclic; if we assigned to  $x$  multiple times, the second assignment would form a cycle in the graph including  $x$  and the operators used to compute  $x$ . Keeping the data flow graph acyclic is important in many types of analyses that we want to do on the graph. Of course, it is important to know whether the source code assigns to a variable multiple times, because some of those assignments may be mistakes. We consider the analysis of source code for proper use of assignments in [Section 5.5](#).


The data flow graph is generally drawn in the form that is shown in [Fig. 5.7](#). Here, the variables are not explicitly represented by nodes. Instead, the edges are labeled with the variables that they represent. As a result, a variable can be represented by more than one edge. However, the edges are directed and all of the edges for a variable must come from a single source. We use this form for its simplicity and compactness.

The data flow graph for the code makes the order in which the operations are performed in the C code much less obvious. This is one of the advantages of the data flow graph. We can use it to determine feasible reorderings of the operations, which may help us to reduce pipeline or cache conflicts. We can also use it when the exact order of operations simply doesn't matter. The data flow graph defines a partial ordering of the operations in the basic block. We must ensure that a value is computed before it is used, but generally, there are several possible orderings of evaluating expressions that satisfy this requirement.



**FIGURE 5.6**

An extended data flow graph for our sample basic block.



**FIGURE 5.7**

Standard data flow graph for our sample basic block.

### 5.3.2 Control/data flow graphs

A CDFG uses a data flow graph as an element, adding constructs to describe control. In a basic CDFG, we have two types of nodes: **decision nodes** and **data flow nodes**. A data flow node encapsulates a complete data flow graph to represent a basic block. We can use one type of decision node to describe all types of control in a sequential program. The jump/branch is, after all, the way that we implement all of those high-level control constructs.

Fig. 5.8 shows a portion of C code with control constructs and the CDFG constructed from it. The rectangular nodes in the graph represent the basic blocks. The basic blocks in the C code have been represented by function calls for simplicity. The diamond-shaped nodes represent the conditionals. The node's condition is given by the label, and the edges are labeled with the possible outcomes of evaluating the condition.

Building a CDFG for a `while` loop is straightforward, as shown in Fig. 5.9. The `while` loop consists of both a test and a loop body, each of which we know how to represent in a CDFG. We can represent `for` loops by remembering that, in C, a `for` loop is defined in terms of a `while` loop [Ker88]. This `for` loop

```
for (i = 0; i < N; i++) {
    loop_body();
}
```

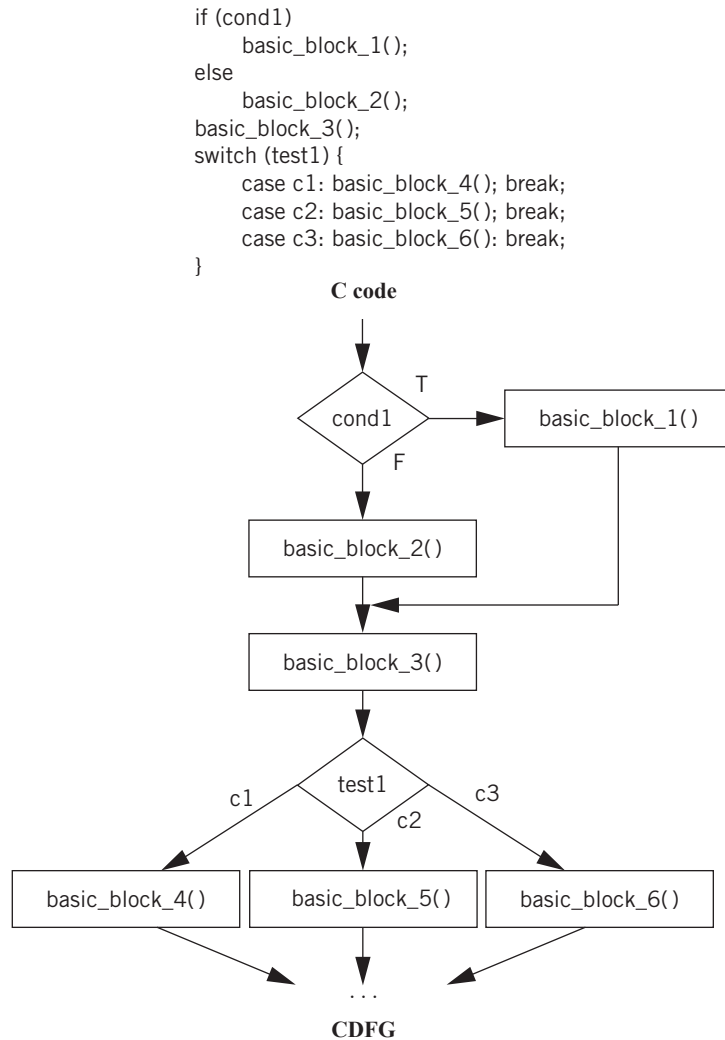
is equivalent to

```
i = 0;
while (i < N) {
    loop_body();
    i++;
}
```

Hierarchical  
representation

For a complete CDFG model, we can use a data flow graph to model each data flow node. Thus, the CDFG is a hierarchical representation; a data flow CDFG can be expanded to reveal a complete data flow graph.

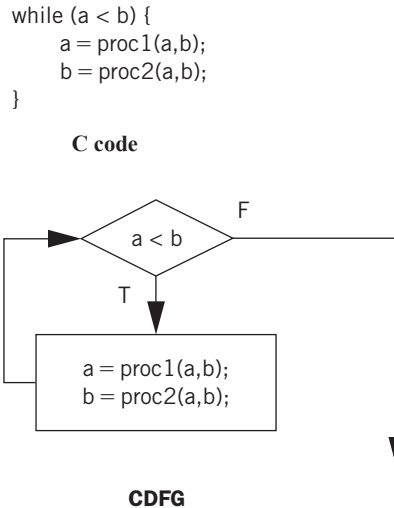
An execution model for a CDFG is very much like the execution of the program that it represents. The CDFG does not require the explicit declaration of variables and we assume that the implementation has sufficient memory for all variables. We can define a state variable that represents a program counter in a CPU. (When studying a drawing of a CDFG, a finger works well for keeping track of the program counter state.) As we execute the program, we either execute the data flow node or compute the decision in the decision node and follow the appropriate edge, depending on the type of node that the program counter points to. Even though the data flow nodes may specify only a partial ordering on the data flow computations, the CDFG is a sequential representation of the program. There is only one program counter in our execution model of the CDFG, and operations are not executed in parallel.

**FIGURE 5.8**

C code and its control/data flow graph.

The CDFG is not necessarily tied to high-level language control structures. We can also build a CDFG for an assembly language program. A jump instruction corresponds to a nonlocal edge in the CDFG. Some architectures, such as ARM and many VLIW processors, support the predicated execution of instructions, which may be represented by special constructs in the CDFG.




**FIGURE 5.9**

A while loop and its CDFG.

## 5.4 Assembly, linking, and loading

Assembly and linking are the last steps in the compilation process. They change a list of instructions into an image of the program's bits in memory. Loading puts the program into memory so that it can be executed. In this section, we survey the basic techniques required for assembly linking to help us to understand the complete compilation and loading process.

Fig. 5.10 highlights the role of assemblers and linkers in the compilation process. This process is often hidden from us by compilation commands that do everything that is required to generate an executable program. As the figure shows, most compilers do not directly generate machine code, but instead create the instruction-level program in the form of human-readable assembly language. Generating assembly language rather than binary instructions frees the compiler writer from details that are extraneous to the compilation process, including the instruction format as well as the exact addresses of instructions and data. The assembler's job is to translate symbolic assembly language statements into bit-level representations of instructions that are known as **object code**. The assembler takes care of instruction formats and does part of the job of translating labels into addresses. However, because the program may be built from many files, the final steps in determining the addresses of instructions and data are performed by the linker, which produces an **executable binary** file. However, this file may not necessarily be in the CPU's memory, unless the linker happens to

**FIGURE 5.10**

Program generation from compilation through loading.

#### Absolute and relative addresses


create the executable directly in RAM. The program that brings the program into memory for execution is called a **loader**.

The simplest form of the assembler assumes that the starting address of the assembly language program has been specified by the programmer. The addresses in such a program are known as **absolute addresses**. However, in many cases, particularly when we are creating an executable out of several component files, we do not want to specify the starting addresses for all of the modules before assembly. If we did, we would have to determine not only the length of each program in memory, but also the order in which they would be linked into the program before assembly. Therefore, most assemblers allow us to use **relative addresses** by specifying at the start of the file that the origin of the assembly language module is to be computed later. Addresses within the module are then computed relative to the start of the module, and the linker is responsible for translating relative addresses into addresses.

### 5.4.1 Assemblers

When translating assembly code into object code, the assembler must translate opcodes, format the bits in each instruction, and translate labels into addresses. In this section, we review the translation of assembly language into binary.

Labels make the assembly process more complex, but they are the most important abstraction provided by the assembler. Labels let the programmer (a human programmer or a compiler generating assembly code) avoid worrying about the locations of instructions and data. Label processing requires making two passes through the assembly source code:



**FIGURE 5.11**  
Symbol table processing during assembly.

1. The first pass scans the code to determine the address of each label.
2. The second pass assembles the instructions using the label values computed in the first pass.

Symbol table

As shown in Fig. 5.11, the name of each symbol and its address are stored in a **symbol table** that is built during the first pass. The symbol table is built by scanning from the first instruction to the last. For the moment, we assume that we know the address of the first instruction in the program. During scanning, the current location in memory is kept in a **program location counter (PLC)**. Despite the similarity in name to the program counter, the PLC is not used to execute the program, but only to assign memory locations to labels. For example, the PLC always makes exactly one pass through the program, whereas the program counter makes many passes over code in a loop. Thus, at the start of the first pass, the PLC is set to the program’s starting address and the assembler looks at the first line. After examining the line, the assembler updates the PLC to the next location (because ARM instructions are four bytes long, the PLC would be incremented by four) and looks at the next instruction. If the instruction begins with a label, a new entry is made in the symbol table, which includes the label name and its value. The value of the label is equal to the current value of the PLC. At the end of the first pass, the assembler rewinds to the beginning of the assembly language file to make the second pass. During the second pass, when a label name is found, the label is looked up in the symbol table and its value substituted into the appropriate place in the instruction.

However, how do we know the starting value of the PLC? The simplest case is addressing. In this case, one of the first statements in the assembly language program is a pseudo-op that specifies the **origin** of the program; that is, the location of the first address in the program. A common name for this pseudo-op (e.g., the one used for the ARM) is the `ORG` statement

```
ORG 2000
```

which puts the start of the program at location 2000. This pseudo-op accomplishes this by setting the PLC’s value to its argument’s value, which is 2000 in this case. Assemblers generally allow a program to have many `ORG` statements in case the instructions or data must be spread around various spots in memory.


Example 5.1 illustrates the use of the PLC in generating the symbol table.

**Example 5.1: Generating a Symbol Table**


Let's use the following simple example of ARM assembly code:

```
ORG 100
label1 ADR r4,c
      LDR r0,[r4]
label2 ADR r4,d
      LDR r1,[r4]
label3 SUB r0,r0,r1
```

The initial ORG statement tells us the starting address of the program. To begin, let's initialize the symbol table to an empty state and put the PLC at the initial ORG statement:




The PLC value shown is at the beginning of this step before we have processed the ORG statement. The ORG tells us to set the PLC value to 100.




To process the next statement, we move the PLC to point to the next statement. However, because the last statement was a pseudo-op that generates no memory values, the PLC value remains at 100.


Because there is a label in this statement, we add it to the symbol table, taking its value from the current PLC value.



To process the next statement, we advance the PLC to point to the next line of the program and increment its value by the length in memory of the last line, namely 4.



We continue this process as we scan the program until we reach the end, at which the state of the PLC and symbol table are as shown below.



Assemblers allow labels to be added to the symbol table without occupying space in the program memory. A typical name for this pseudo-op is `EQU` for equate. For example, in the code,

```
ADD r0,r1,r2
F00 EQU 5
BAZ SUB r3,r4,#F00
```

the `EQU` pseudo-op adds a label named `F00` with value 5 to the symbol table. The value of the `BAZ` label is the same as if the `EQU` pseudo-op were not present, because `EQU` does not advance the PLC. The new label is used in the subsequent `SUB` instruction as the name for a constant. `EQU`s can be used to define symbolic values to help to make the assembly code more structured.

ARM ADR pseudo-op

The ARM assembler supports one pseudo-op that is particular to the ARM instruction set. In other architectures, an address would be loaded into a register (e.g., for an indirect access) by reading it from a memory location. ARM does not have an instruction that can load an effective address, so the assembler supplies the `ADR` pseudo-op to create the address in the register. It does so by using `ADD` or `SUB` instructions to generate the address. The address to be loaded can be register relative, program relative, or

**Object code formats**

numeric, but it must assemble to a single instruction. More complicated address calculations must be programmed explicitly.

The assembler produces an object file that describes the instructions and data in binary format. A commonly used object file format, which was originally developed for Unix but is now used in other environments as well, is known as the Common Object File Format (COFF). The object file must describe the instructions, data, and any addressing information, and also usually carries along the symbol table for later use in debugging.

Generating relative code rather than code introduces some new challenges into the assembly language process. Rather than using an `ORG` statement to provide the starting address, the assembly code uses a pseudo-op to indicate that the code is in fact relocatable. Relative code is the default for the ARM assembler. Similarly, we must mark the output object file as being relative code. We can initialize the PLC to 0 to ensure that addresses are relative to the start of the file. However, we must be careful when we generate code that makes use of those labels, because we do not yet know the actual value that must be put into the bits. Instead, we must generate relocatable code. We use extra bits in the object file format to mark the relevant fields as relocatable, and then, insert the label's relative value into the field. The linker must therefore modify the generated code; when it finds a field that is marked as relative, it uses the addresses that it has generated to replace the relative value with a correct value for the address. To understand the details of changing relocatable code into executable code, we must understand the linking process that is described in the next section.

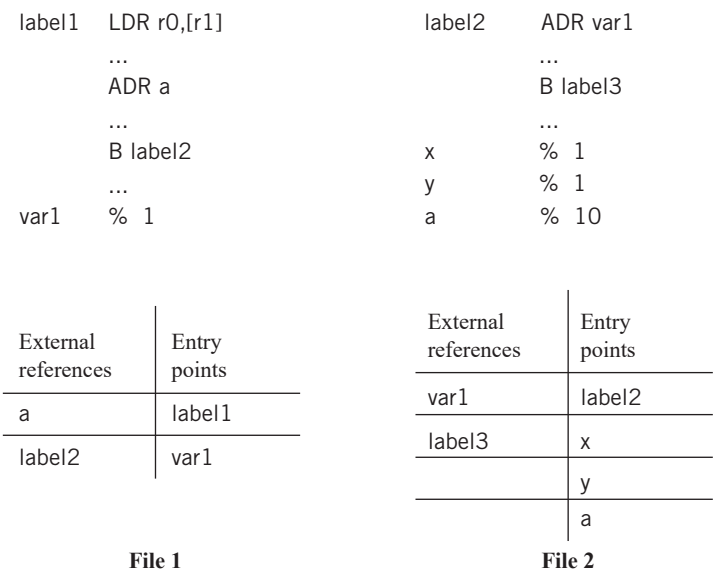
## 5.4.2 Linking

Many assembly language programs are written as several smaller pieces rather than as a single large file. Breaking a large program into smaller files helps to delineate program modularity. If the program uses library routines, those will already be preassembled, and assembly language source code for the libraries may not be available for purchase. A **linker** allows a program to be stitched together out of several smaller pieces. The linker operates on the object files created by the assembler and modifies the assembled code to make the necessary links between files.

Some labels will be both defined and used in the same file. Other labels will be defined in a single file but used elsewhere, as illustrated in [Fig. 5.12](#). The place in the file where a label is defined is known as an **entry point**. The place in the file where the label is used is called an **external reference**. The main job of the loader is to *resolve* external references based on available entry points. As a result of the need to know how definitions and references connect, the assembler passes not only the object file, but also the symbol table, to the linker. Even if the entire symbol table is not kept for later debugging purposes, it must at least pass the entry points. External references are identified in the object code by their relative symbol identifiers.

**Linking process**

The linker proceeds in two phases. First, it determines the address of the start of each object file. The order in which object files are to be loaded is given by the user,



to be easily updated. However, it does introduce a delay before the program starts executing.

### 5.4.3 Object code design

We have to take several issues into account when designing object code. In a timesharing system, many of these details are taken care of for us. When designing an embedded system, we may need to handle some of them ourselves.

#### Memory map design

As we have seen, the linker allows us to control where object code modules are placed in memory. We may need to control the placement of several types of data:

- Interrupt vectors and other information for I/O devices must be placed in specific locations.
- Memory management tables must be set up.
- Global variables that are used for communication between processes must be put in locations that are accessible to all users of those data.

We can give these locations symbolic names so that, for example, the same software can work on different processors that put these items at different addresses. However, the linker must be given the proper absolute addresses to configure the program's memory.

#### Reentrancy

Many programs should be designed to be **reentrant**. A program is reentrant if can be interrupted by another call to the function without changing the results of either call. If the program changes the values of global variables, it may give a different answer when it is called recursively. Consider this code:

```
int foo = 1;

int task1() {
    foo = foo + 1;
    return foo;
}
```

In this simple example, the variable `foo` is modified; thus, `task1()` gives a different answer on every invocation. We can avoid this problem by passing `foo` in as an argument:

```
int task1(int foo) {
    return foo+1;
}
```

---

## 5.5 Compilation techniques

Although we don't write our own assembly code in most cases, we still care about the characteristics of the code that our compiler generates: its speed, its size, and its power consumption. Understanding how a compiler works will help us to write code and



direct the compiler to get the assembly language implementation that we want. We will start with an overview of the compilation process, followed by some basic compilation methods, and conclude with some more advanced optimizations.

### 5.5.1 The compilation process

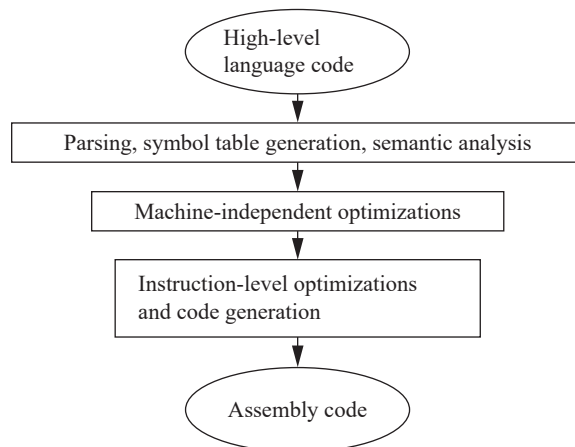
It is useful to understand how a high-level language program is translated into instructions, interrupt handling instructions, place data and instructions in memory, and so on. Understanding how the compiler works can help you to know when you cannot rely on the compiler. Next, because many applications are also performance sensitive, understanding how code is generated can help you to meet your performance goals, either by writing high-level code that gets compiled into the instructions you want or by recognizing when you must write your own assembly code.

We can summarize the compilation process with a formula:

$$\text{compilation} = \text{translation} + \text{optimization}$$

The high-level language program is translated into the lower-level form of instructions; optimizations try to generate better instruction sequences than would be possible if the brute force technique of independently translating source code statements were used. Optimization techniques focus on more of the program to ensure that compilation decisions that appear to be good for one statement are not unnecessarily problematic for other parts of the program.

The compilation process is outlined in [Fig. 5.13](#). Compilation begins with high-level language code such as C or C++ and generally produces assembly code. Directly producing object code simply duplicates the functions of an assembler, which is a very desirable standalone program to have. The high-level language program is



**FIGURE 5.13**

The compilation process.

parsed to break it into statements and expressions. In addition, a symbol table is generated, which includes all of the named objects in the program. Some compilers may then perform higher-level optimizations that can be viewed as modifying the high-level language program input without reference to instructions.

Simplifying arithmetic expressions is one example of a machine-independent optimization. Not all compilers do such optimizations, and compilers can vary widely regarding which combinations of machine-independent optimizations they do perform. Instruction-level optimizations are aimed at generating code. They may work directly on real instructions or on a pseudo-instruction format that is later mapped onto the instructions of the target CPU. This level of optimization also helps to modularize the compiler by allowing code generation to create simpler code that is later optimized. For example, consider this array access code:

```
x[i] = c*x[i];
```

A simple code generator would generate the address for `x[i]` twice: once for each appearance in the statement. The later optimization phases can recognize this as an example of common expressions that need not be duplicated. While it would be possible to create a code generator that never generated the redundant expression in this simple case, taking into account every such optimization at code generation time is very difficult. We get better code and more reliable compilers by generating simple code first, and then, optimizing it.

## 5.5.2 Basic compilation methods

### Statement translation

In this section, we consider the basic job of translating a high-level language program with little or no optimization.

### Procedures

Procedures (or functions, as they are known in C) require specialized code: the code that is used to call and pass parameters is known as **procedure linkage**; procedures also need a way to store local variables. Generating code for procedures is relatively straightforward once we know the procedure linkage that is appropriate for the CPU. At the procedure definition, we generate the code to handle the procedure call and return. At each call of the procedure, we set up the procedure parameters and make the call.

The CPU's subroutine call mechanism is usually not sufficient to support procedures directly in modern programming languages. We introduced the procedure stack and procedure linkages in Section 2.3.3. The linkage mechanism provides a way for the program to pass parameters into the program and for the procedure to return a value. It also provides help in restoring the values of registers that the procedure has modified. All procedures in a given programming language use the same linkage mechanism (although different languages may use different linkages). The mechanism can also be used to call handwritten assembly language routines from compiled code.

The information for a call to a procedure is known as a **frame**; the frames are stored on a stack to keep track of the order in which the procedures have been called.

Procedure stacks are typically built to grow down from high addresses. A **stack pointer** (sp) defines the end of the current frame, whereas a **frame pointer** (fp) defines the end of the last frame. The fp is technically necessary only if the stack frame can be grown by the procedure during execution. The procedure can refer to an element in the frame by addressing relative to sp. When a new procedure is called, the sp and fp are modified to push another frame onto the stack. In addition to allowing parameters and return values to be passed, the frame also holds the locally declared variables. When accessing a local variable, the compiled code must do so by referencing a location within the frame, which requires it to perform address arithmetic.

As we saw in [Chapter 2](#), the ARM Procedure Call Standard (APCS) [Slo04] is the recommended procedure linkage for ARM processors. r0–r3 are used to pass the first four parameters into the procedure. r0 is also used to hold the return value.

The next example looks at compiler-generated procedure linkage code.

---

### Programming Example 5.6: Procedure Linkage in C

Here is a procedure definition:

```
int pl(int a, int b, int c, int d, int e) {
    return a + e;
}
```

This procedure has five parameters, so we would expect that one of them would be passed through the stack while the rest are passed through registers. It also returns an integer value, which should be returned in r0. Here is the code for the procedure generated by the ARM gcc compiler with some handwritten comments:

```
mov    ip, sp                ; procedure entry
stmfd  sp!, {fp, ip, lr, pc}
sub    fp, ip, #4
sub    sp, sp, #16
str    r0, [fp, #-16]        ; put first four args on stack
str    r1, [fp, #-20]
str    r2, [fp, #-24]
str    r3, [fp, #-28]
ldr    r2, [fp, #-16]        ; load a
ldr    r3, [fp, #4]          ; load e
add    r3, r2, r3            ; compute a + e
mov    r0, r3                ; put the result into r0 for return
ldmea  fp, {fp, sp, pc}     ; return
```

Here is a call to that procedure:

```
y = pl(a,b,c,d,x);
```

Here is the ARM gcc code with handwritten comments:

```
ldr    r3, [fp, #-32]        ; get e
str    r3, [sp, #0]          ; put into pl()'s stack frame
ldr    r0, [fp, #-16]        ; put a into r0
ldr    r1, [fp, #-20]        ; put b into r1
```

```

ldr    r2, [fp, #-24]    ; put c into r2
ldr    r3, [fp, #-28]    ; put d into r3
bl     p1 ; call p1()
mov    r3, r0            ; move return value into r3
str    r3, [fp, #-36]    ; store into y in stack frame

```

We can see that the compiler sometimes makes additional register moves but it does follow the APCS standard.

A large amount of the code in a typical application consists of arithmetic and logical expressions. Understanding how to compile a single expression, as described in the next example, is a good first step in understanding the entire compilation process.


### Example 5.2: Compiling an Arithmetic Expression

Consider this arithmetic expression:


$$x = a * b + 5 * (c - d)$$

The expression is written in terms of program variables. In some machines, we may be able to perform memory-to-memory arithmetic directly on the locations corresponding to those variables. However, in many machines, such as the ARM, we must first load the variables into registers. This requires selecting which registers receive not only the named variables, but also intermediate results such as  $(c - d)$ .

The code for the expression can be built by walking the data flow graph. Here is the data flow graph for the expression:



The temporary variables for the intermediate values and final result have been named  $w$ ,  $x$ ,  $y$ , and  $z$ . To generate code, let's walk from the tree's root (where  $z$ , the final result, is generated) by traversing the nodes in post order. During the walk, we generate instructions to cover the operation at every node. Here is the path:



The nodes are numbered in the order in which code is generated. Because every node in the data flow graph corresponds to an operation that is directly supported by the instruction set, we simply generate an instruction at every node. Because we are making an arbitrary register assignment, we can use up the registers in order, starting with r1. Here is the resulting ARM code:

```

; operator 1 (*)
ADR r4,a           ; get address for a
MOV r1,[r4]        ; load a
ADR r4,b           ; get address for b
MOV r2,[r4]        ; load b
ADD r3,r1,r2       ; put w into r3
; operator 2 (-)
ADR r4,c           ; get address for c
MOV r4,[r4]        ; load c
ADR r4,d           ; get address for d
MOV r5,[r4]        ; load d
SUB r6,r4,r5       ; put z into r6
; operator 3 (*)
MUL r7,r6,#5       ; operator 3, puts y into r7
; operator 4 (+)
ADD r8,r7,r3       ; operator 4, puts x into r8
; assign to x
ADR r1,x
STR r8,[r1]        ; assigns to x location

```

One obvious optimization is to reuse a register whose value is no longer needed. In the case of the intermediate values w, y, and z, we know that they cannot be used after the end of the expression (e.g., in another expression) because they have no name in the C program. However, the final result z may in fact be used in a C assignment, and the value can be reused later in the program. In this case, we would need to know when the register is no longer needed to determine its best use.


For comparison, here is the code generated by the ARM gcc compiler with handwritten comments:

```
ldr r2, [fp, #-16]
ldr r3, [fp, #-20]
mul r1, r3, r2      ; multiply
ldr r2, [fp, #-24]
ldr r3, [fp, #-28]
rsb r2, r3, r2      ; subtract
mov r3, r2
mov r3, r3, asl #2
add r3, r3, r2      ; add
add r3, r1, r3      ; add
str r3, [fp, #-32]  ; assign
```

In the previous example, we made an arbitrary allocation of variables to registers for simplicity. When we have large programs with multiple expressions, we must allocate registers more carefully, because CPUs have a limited number of registers. We will consider register allocation in more detail below.

We also need to be able to translate control structures. Because conditionals are controlled by expressions, the code generation techniques of the last example can be used for those expressions, leaving us with the task of generating code for the flow of control itself. Fig. 5.14 shows a simple example of changing the flow of control in C—an `if` statement, in which the condition controls whether the true or false branch of the `if` is taken. Fig. 5.14 also shows the control flow diagram for the `if` statement.

The next example illustrates how to implement conditionals in assembly language.



**FIGURE 5.14**


Flow of control in C and control flow diagrams.

**Example 5.3: Generating Code for a Conditional**


Consider this C statement:

```
if (a + b > 0)
    x = 5;
else
    x = 7;
```

The CDFG for the statement appears below.



We know how to generate the code for the expressions. We can generate the control flow code by walking the CDFG. One ordered walk through the CDFG follows:



To generate code, we must assign a label to the first instruction at the end of a directed edge and create a branch for each edge that does not go to the immediately following instruction. The exact steps to be taken at the branch points depend on the target architecture. On some machines, evaluating expressions generates condition codes that we can test in subsequent branches, and on other machines, we must use test-and-branch instructions. ARM allows us to test condition codes, so we get the following ARM code for the 1–2–3 walk:

```

ADR r5,a      ; get address for a
LDR r1,[r5]   ; load a
ADR r5,b      ; get address for b
LDR r2,b      ; load b
ADD r3,r1,r2
BLE label3    ; true condition falls through branch
; true case
LDR r3,#5     ; load constant
ADR r5,x
STR r3,[r5]   ; store value into x
B stmtend     ; done with the true case
```

```

; false case
label13 LDR r3,#7    ; load constant
        ADR r5,x     ; get address of x
        STR r3,[r5]  ; store value into x
stmtend

```

The 1–2 and 3–4 edges do not require a branch and label because they are straight-line code. In contrast, the 1–3 and 2–4 edges require a branch and a label for the target.

For comparison, here is the code generated by the ARM gcc compiler with some hand-written comments:

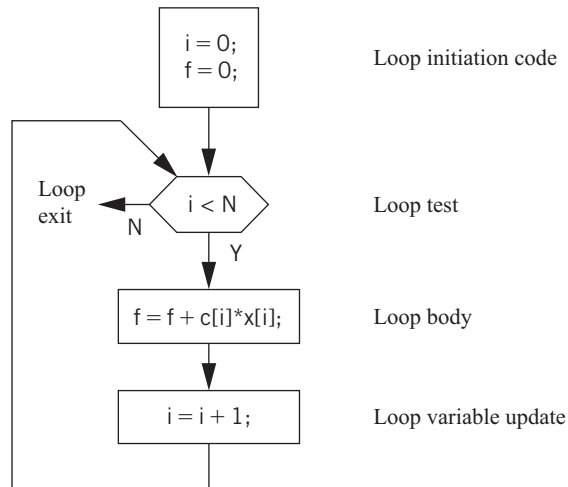
```

ldr r2, [fp, #-16]
ldr r3, [fp, #-20]
add r3, r2, r3
cmp r3, #0           ; test the branch condition
ble .L3             ; branch to false block if <=
mov r3, #5           ; true block
str r3, [fp, #-32]
b .L4               ; go to end of if statement
.L3:                ; false block
mov r3, #7
str r3, [fp, #-32]
.L4:

```

Because expressions are generally created as straight-line code, they typically require careful consideration of the order in which the operations are executed. We have much more freedom when generating conditional code, because the branches ensure that the flow of control goes to the correct block of code. If we walk the CDFG in a different order and lay out the code blocks in a different order in memory, we still get valid code as long as we place the branches properly.

Drawing a control flow graph based on the `while` form of the loop helps us to understand how to translate it into instructions:





C compilers can generate assembler source (using the `-s` flag), which some compilers intersperse with the C code. Such code is a very good way to learn about both assembly language programming and compilation.

## Data structures

The compiler must also translate references to data structures into references to raw memories. In general, this requires address computations. Some of these computations can be done at compile time, whereas others must be done at run time.

Arrays are interesting because the address of an array element must be computed at run time in general, because the array index may change. Let us first consider a one-dimensional array:

```
a[i]
```

The layout of the array in memory is shown in Fig. 5.15: the zeroth element is stored as the first element of the array, the first element directly below, and so on. We can create a pointer for the array that points to the array's head, namely `a[0]`. If we call that pointer `aptr` for convenience, we can rewrite the reading of `a[i]` as

```
*(aptr + i)
```


Two-dimensional arrays are more challenging. There are multiple possible ways to lay out a two-dimensional array in memory, as shown in Fig. 5.16. In this form, which is known as **row major**, the rightmost variable of the array (`j` in `a[i][j]`) varies the most quickly. Fortran uses the opposite organization, known as *column major*, in which the leftmost array index varies the fastest. Two-dimensional arrays also require more sophisticated addressing; in particular, we must know the size of the array. Let us consider the row major form. If the `a[][]` array is of size  $M \times N$ , we can change the two-dimensional array access into a one-dimensional array access. Thus,

```
a[i][j]
```

becomes

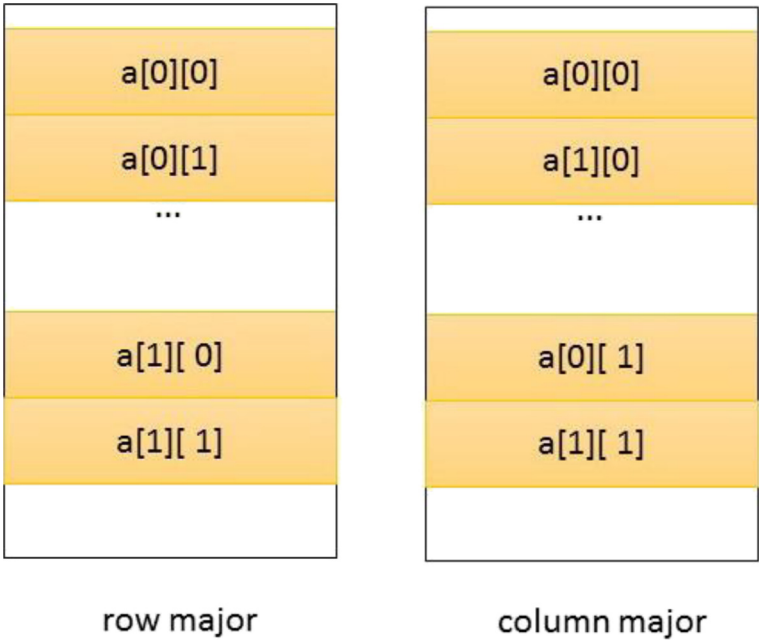
```
a[i*M + j]
```

A C struct is easier to address. As shown in Fig. 5.16, a structure is implemented as a contiguous block of memory. Fields in the structure can be accessed using constant



**FIGURE 5.15**

Layout of a one-dimensional array in memory.



**FIGURE 5.16**  
Memory layout for two-dimensional arrays.

offsets to the base address of the structure. In this example, if `field1` is four bytes long, `field2` can be accessed as

```
*(aptr + 4)
```

This addition can usually be done at compile time, requiring only the indirection itself to fetch the memory location during execution.

**5.5.3 Compiler optimizations**

Basic compilation techniques can generate inefficient code. Compilers use a wide range of algorithms to optimize the code that they generate.

**Inlining**

**Function inlining** replaces a subroutine call to a function with equivalent code to the function body. By substituting the function call’s parameters into the body, the compiler can generate a copy of the code that performs the same operations, but without the subroutine overhead. C++ provides an *inline* qualifier that allows the compiler to substitute an inline version of the function. In C, programmers can perform inlining manually or by using a preprocessor macro to define the code body.

**Outlining** is the opposite of inlining: a set of similar sections of code is replaced with calls to an equivalent function. Although inlining eliminates the function call overhead, it also increases the program size. Inlining also inhibits sharing of the

## Loop transformations

function code in the cache; because the inlined copies are distinct pieces of code, they cannot be represented by the same code in the cache. Outlining is sometimes useful to improve the cache behavior of common functions.

Loops are important program structures; although they are compactly described in the source code, they often use a large fraction of the computation time. Many techniques have been designed to optimize loops.

A simple but useful transformation is known as **loop unrolling**, as illustrated in the next example. Loop unrolling is important because it helps to expose parallelism that can be used by later stages of the compiler.

---

### Example 5.4: Loop Unrolling

Here is a simple C loop:

```
for (i = 0; i < N; i++) {
    a[i] = b[i] * c[i];
}
```

This loop is executed a fixed number of times, namely  $N$ . A straightforward implementation of the loop would create and initialize the loop variable  $i$ , update its value at every iteration, and test it to see whether to exit the loop. However, because the loop is executed a fixed number of times, we can generate more direct code.

If we let  $N = 4$ , then we can substitute this straight-line code for the loop:

```
a[0] = b[0] * c[0];
a[1] = b[1] * c[1];
a[2] = b[2] * c[2];
a[3] = b[3] * c[3];
```

This unrolled code has no loop overhead code at all; that is, no iteration variable and no tests. However, the unrolled loop has the same problems as the inlined procedure; it may interfere with the cache and expands the amount of code required.

We do not, of course, have to unroll loops fully. Rather than unroll the above loop four times, we could unroll it twice. Unrolling produces this code:

```
for (i = 0; i < 2; i++) {
    a[i*2] = b[i*2] * c[i*2];
    a[i*2 + 1] = b[i*2 + 1] * c[i*2 + 1];
}
```

In this case, because all operations in the two lines of the loop body are independent, later stages of the compiler may be able to generate code that allows them to be executed efficiently on the CPU's pipeline.

---

**Loop fusion** combines two or more loops into a single loop. For this transformation to be legal, two conditions must be satisfied. First, the loops must iterate over the same values. Second, the loop bodies must not have dependencies that would be violated if they are executed together; for example, if the second loop's  $i$ th iteration depends on the results of the  $i + 1$ th iteration of the first loop, the two loops cannot

**Dead code elimination**

be combined. **Loop distribution** is the opposite of loop fusion; that is, decomposing a single loop into multiple loops.

**Dead code** is code that can never be executed. Dead code can be generated by programmers, either inadvertently or purposefully. Dead code can also be generated by compilers. Dead code can be identified by **reachability analysis**, finding the other statements or instructions from which it can be reached. If a given piece of code cannot be reached or if it can be reached only by a piece of code that is unreachable from the main program, it can be eliminated. **Dead code elimination** analyzes code for reachability and trims away dead code.

**Register allocation**

**Register allocation** is a very important compilation phase. Given a block of code, we want to select assignments of variables (both declared and temporary) to registers to minimize the total number of required registers.

The next example illustrates the importance of proper register allocation.

---


**Example 5.5: Register Allocation**

To keep the example small, we assume that we can use only four of the ARM's registers. In fact, such a restriction is not unthinkable. Programming conventions can reserve certain registers for special purposes and significantly reduce the number of general-purpose registers that are available.

Consider this C code:

```
w = a + b; /*statement 1 */
x = c + w; /*statement 2 */
y = c + d; /*statement 3 */
```

A naive register allocation, assigning each variable to a separate register, would require seven registers for the seven variables in the above code. However, we can do much better by reusing a register once the value stored in the register is no longer needed. To understand how to do this, we can draw a **lifetime graph** that shows the statements on which each statement is used. Here is a lifetime graph in which the x axis is the statement number in the C code, and the y axis shows the variables:



A horizontal line stretches from the first statement where the variable is used to the last use of the variable; a variable is said to be **live** during this interval. At each statement, we can

determine every variable that is currently in use. The maximum number of variables in use at any statement determines the maximum number of registers that are required. In this case, statement two requires three registers: *c*, *w*, and *x*. This fits within the four-register limitation. By reusing registers once their current values are no longer needed, we can write code that requires no more than four registers. Here is one register assignment:

<i>a</i>	<i>r0</i>
<i>b</i>	<i>r1</i>
<i>c</i>	<i>r2</i>
<i>d</i>	<i>r0</i>
<i>w</i>	<i>r3</i>
<i>x</i>	<i>r0</i>
<i>y</i>	<i>r3</i>

Here is the ARM assembly code that uses the above register assignment:


```
LDR r0,[p_a]      ; load a into r0 using pointer to a (p_a)
LDR r1,[p_b]      ; load b into r1
ADD r3,r0,r1      ; compute a + b
STR r3,[p_w]      ; w = a + b
LDR r2,[p_c]      ; load c into r2
ADD r0,r2,r3      ; compute c + w, reusing r0 for x
STR r0,[p_x]      ; x = c + w
LDR r0,[p_d]      ; load d into r0
ADD r3,r2,r0      ; compute c + d, reusing r3 for y
STR r3,[p_y]      ; y = c + d
```

If a section of code requires more registers than are available, we must **spill** some of the values out to memory temporarily. After computing some values, we write the values to temporary memory locations, reuse those registers in other computations, and then, reread the old values from the temporary locations to resume work. Spilling registers is problematic in several respects: it requires extra CPU time, and uses up both instruction and data memory. Putting effort into register allocation to avoid unnecessary register spills is worth your time.

We can solve register allocation problems by building a **conflict graph** and solving a graph coloring problem. As shown in Fig. 5.17, each variable in the high-level language code is represented by a node. An edge is added between two nodes if they are both live at the same time. The graph coloring problem is to use the smallest number of distinct colors to color all of the nodes, such that no two nodes are directly connected by an edge of the same color. The figure shows a satisfying coloring that uses three colors. Graph coloring is NP-complete, but there are efficient heuristic algorithms that can give good results on typical register allocation problems.

Lifetime analysis assumes that we have already determined the order in which we will evaluate operations. In many cases, we have freedom in the order in which we do things. Consider this expression:

$$(a + b) * (c - d)$$

**FIGURE 5.17**

Using graph coloring to solve the problem of Example 5.5.

We must perform the multiplication last, but we can do either the addition or the subtraction first. Different orders of loads, stores, and arithmetic operations may also result in different execution times on pipelined machines. If we can keep values in registers without having to reread them from main memory, we can save on execution time and reduce the code size as well.


The next example shows how proper **operator scheduling** can improve register allocation.

### Example 5.6: Operator Scheduling for Register Allocation

Here is a sample C code fragment:

```
w = a + b; /* statement 1 */
x = c + d; /* statement 2 */
y = x + e; /* statement 3 */
z = a - b; /* statement 4 */
```


If we compile the statements in the order in which they were written, we get this register graph:



Because *w* is needed until the last statement, we need five registers at statement 3, even though only three registers are needed for the statement at line 3. If we swap statements 3 and 4 (renumbering them 39 and 49), we reduce our requirements to three registers. Here is the modified C code:

```
w = a + b; /*statement 1 */
z = a - b; /* statement 29 */
x = c + d; /*statement 39 */
y = x + e; /*statement 49 */
```

Additionally, here is the lifetime graph for the new code:



Compare the ARM assembly code for the two code fragments. We have written both assuming that we have only four free registers. In the *before* version, we do not have to write out any values, but we must read *a* and *b* twice. The *after* version allows us to retain all values in registers for as long as we need them.

#### *Before version*

```
LDR r0,a
LDR r1,b
ADD r2,r0,r1
STR r2,w ; w = a + b
LDR r0,c
LDR r1,d
ADD r2,r0,r1
STR r2,x ; x = c + d
LDR r1,e
ADD r0,r1,r2
STR r0,y ; y = x + e
LDR r0,a ; reload a
LDR r1,b ; reload b
SUB r2,r1,r0
STR r2,z ; z = a - b
```

#### *After version*

```
LDR r0,a
LDR r1,b
ADD r2,r1,r0
STR r2,w ; w = a + b
SUB r2,r0,r1
STR r2,z ; z = a - b
LDR r0,c
LDR r1,d
ADD r2,r1,r0
STR r2,x ; x = c + d
LDR r1,e
ADD r0,r1,r2
STR r0,y ; y = x + e
```

Scheduling

We have some freedom to choose the order in which operations will be performed. We can use this to our advantage; for example, we may be able to improve the register allocation by changing the order in which operations are performed, thereby changing the lifetimes of the variables.

We can solve scheduling problems by keeping track of resource utilization over time. We do not have to know the exact microarchitecture of the CPU. All we have to know is that, for example, instruction types 1 and 2 both use resource A, while instruction types 3 and 4 use resource B. CPU manufacturers generally disclose enough information about the microarchitecture to allow us to schedule instructions even when they do not provide a detailed description of the CPU’s internals.

We can keep track of CPU resources during instruction scheduling using a **reservation table** [Kog81]. As illustrated in Fig. 5.18, the rows in the table represent the instruction execution time slots and the columns represent resources that must be scheduled. Before scheduling an instruction to be executed at a particular time, we check the reservation table to determine whether all resources that are needed by the instruction are available at that time. Upon scheduling the instruction, we update the table to note all resources that are used by that instruction. Various algorithms can be used for the scheduling itself, depending on the types of resources and instructions involved, but the reservation table provides a good summary of the state of an instruction scheduling problem that is in progress.

We can also schedule instructions to maximize performance. As we know from Section 3.6, when an instruction that takes more cycles than normal to finish is in the pipeline, pipeline bubbles appear that reduce the performance. **Software pipelining** is a technique for reordering instructions across several loop iterations to reduce pipeline bubbles. Some instructions take several cycles to complete; if the value that is produced by one of these instructions is needed by other instructions in the loop iteration, they must wait for that value to be produced. Rather than pad the loop with no-ops, we can start instructions from the next iteration. The loop body then contains instructions that manipulate values from several different loop iterations; some of the instructions are working on the early part of iteration  $n + 1$ , others are working on iteration  $n$ , and still others are finishing iteration  $n - 1$ .

Instruction selection

Selecting the instructions to use to implement each operation is not trivial. There may be several different instructions that can be used to accomplish the same goal, but

Time	Resource A	Resource B
t	X	
t + 1	X	X
t + 2	X	
t + 3		X


FIGURE 5.18

A reservation table for instruction scheduling.



they may have different execution times. Moreover, using one instruction for one part of the program may affect the instructions that can be used in adjacent code. Although we can't discuss all of the problems and methods for code generation here, a little bit of knowledge helps us to envision what the compiler is doing.

One useful technique for generating code is **template matching**, as illustrated in Fig. 5.19. We have a directed acyclic graph (DAG) that represents the expression for which we want to generate code. To help to match up instructions and operations, we represent instructions using the same DAG representation. We shade the instruction template nodes to distinguish them from code nodes. Each node has a cost, which may be simply the execution time of the instruction, or may include factors for size, power consumption, and so on. In this case, we have shown that each instruction takes the same amount of time, and thus, all have a cost of 1. Our goal is to cover all nodes in the code DAG with instruction DAGs; until we have covered the code DAG, we haven't generated code for all of the operations in the expression. In this case, the lowest-cost covering uses the multiply–add instruction to cover both nodes. If we first tried to cover the bottom node with the multiply instruction,



**FIGURE 5.19**

Code generation by template matching.

**Understanding your compiler**

we would find ourselves blocked from using the multiply—add instruction. Dynamic programming can be used to find the lowest-cost covering of trees efficiently and heuristics can extend the technique to DAGs.

Clearly, the compiler can vastly transform your program during the creation of assembly language. However, compilers are also substantially different in terms of the optimizations they perform. Understanding your compiler can help you get the best code out of it.

Studying the assembly language output of the compiler is a good way to learn about what the compiler does. Some compilers will annotate sections of code to help you to make the correspondence between the source and assembler output. Starting with small examples that exercise only a few types of statements will help. You can experiment with different optimization levels (the `-O` flag on most C compilers). You can also try writing the same algorithm in several ways to see how the compiler's output changes.

If you can't get your compiler to generate the code that you want, you may need to write your own assembly language. You can do this by writing it from scratch or by modifying the output of the compiler. If you write your own assembly code, you must ensure that it conforms to all compiler conventions, such as procedure call linkage. If you modify the compiler output, you should be sure that you have the algorithm correct before you start writing code so that you don't have to edit the compiler's assembly language output repeatedly. You also need to document the fact that the high-level language source is, in fact, not the code used in the system.

---


## 5.6 Program-level performance analysis

Because embedded systems must perform functions in real time, we often need to know how fast a program runs. The techniques that we use to analyze program execution time are also helpful in analyzing properties such as power consumption. In this section, we study how to analyze programs to estimate their run times. We also examine how to optimize programs to improve their execution times; of course, optimization relies on analysis.

It is important to keep in mind that CPU performance is not judged in the same way as program performance. Certainly, the CPU clock rate is a very unreliable metric for program performance. However, more importantly, the fact that the CPU executes part of our program quickly doesn't mean that it will execute the entire program at the desired rate. As illustrated in [Fig. 5.20](#), the CPU pipeline and cache act as windows into our program. To understand the total execution time of our program, we must look at execution paths, which are far longer than the pipeline and cache windows in general. The pipeline and cache influence the execution time, but the execution time is a global property of the program.

While we might hope that the execution time of programs could be precisely determined, this is in fact difficult to do in practice:

- The execution time of a program often varies with the input data values, because those values select different execution paths in the program. For example, loops

**FIGURE 5.20**

Execution time is a global property of a program.

may be executed a varying number of times and different branches may execute blocks of varying complexity.

- The cache has a major effect on the program performance, and once again, the cache's behavior depends in part on the data values that are input into the program.
- Execution times may vary even at the instruction level. Floating-point operations are the most sensitive to data values, but the normal integer execution pipeline can also introduce data-dependent variations. In general, the execution time of an instruction in a pipeline depends not only on that instruction, but also on the instructions around it in the pipeline.

#### Measuring execution speed

We can measure program performance in several ways:

- Some microprocessor manufacturers supply simulators for their CPUs. The simulator runs on a workstation or PC, takes an executable as input for the microprocessor along with input data, and simulates the execution of that program. Some of these simulators go beyond functional simulation to measure the execution time of the program. Simulation is clearly slower than executing the program on the actual microprocessor, but it also provides much greater visibility during execution. Be careful: some microprocessor performance simulators are not 100% accurate, and simulation of I/O-intensive code may be difficult.
- A timer that is connected to the microprocessor bus can be used to measure the performance of executing sections of code. The code to be measured would reset and start the timer at its start and stop the timer at the end of execution. The length of the program that can be measured is limited by the accuracy of the timer.
- A logic analyzer can be connected to the microprocessor bus to measure the start and stop times of a code segment. This technique relies on the code being able to produce identifiable events on the bus to identify the start and stop of execution.

The length of code that can be measured is limited by the size of the logic analyzer's buffer.

We are interested in the following three different types of performance measures on programs:

- **Average-case execution time:** This is the typical execution time that we would expect for typical data. Clearly, the first challenge is defining typical inputs.
- **Worst-case execution time:** The longest time that the program can spend on any input sequence is clearly important for systems that must meet deadlines. In some cases, the input set that causes the worst-case execution time is obvious, but in many cases, it is not.
- **Best-case execution time:** This measure can be important in multirate real-time systems, as seen in [Chapter 6](#).

First, we look at the fundamentals of program performance in more detail. We then consider trace-driven performance based on executing the program and observing its behavior.

### 5.6.1 Analysis of program performance

The key to evaluating the execution time is breaking the performance problem into parts. The program execution time [Sha89] can be seen as

$$\text{execution time} = \text{program path} + \text{instruction timing}$$

The path is the sequence of instructions executed by the program (or its equivalent in the high-level language representation of the program). The instruction timing is determined based on the sequence of instructions traced by the program path, which takes into account data dependencies, pipeline behavior, and caching. Luckily, these two problems can be solved relatively independently.

Although we can trace the execution path of a program through its high-level language specification, it is difficult to get accurate estimates of the total execution time from a high-level language program. This is because there is no direct correspondence between program statements and instructions. The number of memory locations and variables must be estimated, and results may be either saved for reuse or recomputed on the fly, among other effects. These problems become more challenging as the compiler puts more and more effort into optimizing the program. However, some aspects of program performance can be estimated by looking directly at the C program. For example, if a program contains a loop with a large, fixed-iteration bound or if one branch of a conditional is much longer than another, we can get at least a rough idea that these are more time consuming segments of the program.

Of course, a precise estimate of performance also relies on the instructions to be executed, because different instructions take different amounts of time. In addition, to make life even more difficult, the execution time of one instruction can depend on the instructions executed before and after it.

The next example illustrates data-dependent program paths.

### Example 5.5: Data-dependent Paths in *if* Statements

Here is a pair of nested *if* statements:

```

if (a || b) { /* test 1 */
    if (c) /*test 2 */
        { x = r * s + t; /* assignment 1 */ }
    else { y = r + s; /*assignment 2 */ }
    z = r + s + u; /*assignment 3 */
} else {
    if (c) /*test 3 */
        { y = r - t; /* assignment 4 */ }
}

```

The conditional tests and assignments are labeled within each *if* statement to make it easier to identify paths. Which execution paths may be exercised? One way to enumerate all of the paths is to create a truth table-like structure. The paths are controlled by the variables in the *if* conditions, namely, *a*, *b*, and *c*. For any given combination of values of those variables, we can trace through the program to see which branch is taken at each *if* and which assignments are performed. For example, when *a* = 1, *b* = 0, and *c* = 1, test 1 is true and test 2 is true. This means that we first perform assignment 1, and then, assignment 3.

Here are the results for all of the controlling variable values:

a	b	c	Path
0	0	0	test 1 false, test 3 false: no assignments
0	0	1	test 1 false, test 3 true: assignment 4
0	1	0	test 1 true, test 2 false: assignments 2, 3
0	1	1	test 1 true, test 2 true: assignments 1, 3
1	0	0	test 1 true, test 2 false: assignments 2, 3
1	0	1	test 1 true, test 2 true: assignments 1, 3
1	1	0	test 1 true, test 2 false: assignments 2, 3
1	1	1	test 1 true, test 2 true: assignments 1, 3

Note that there are only four distinct cases: no assignment, assignment 4, assignments 2 and 3, or assignments 1 and 3. These correspond to the possible paths through the nested *ifs*; the table adds value by telling us which variable values exercise each of these paths.

Enumerating the paths through a fixed-iteration *for* loop is seemingly simple. In the code,

```

for (i = 0; i < N; i++)
    a[i] = b[i] * c[i];

```

the assignment in the loop is performed exactly  $N$  times. However, we can't forget the code that is executed to set up the loop and to test the iteration variable.

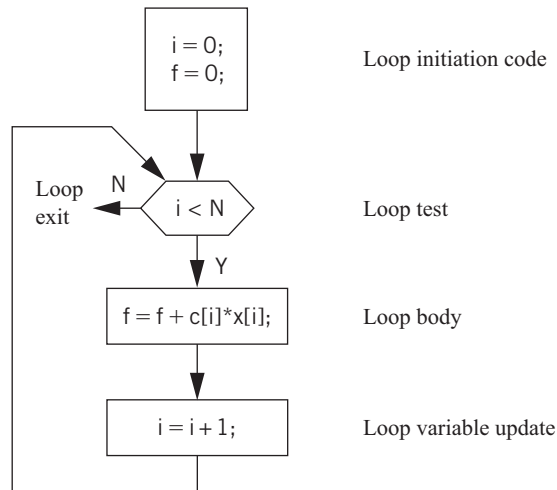
Example 5.6 illustrates how to determine the path through a loop.

### Example 5.6: Paths in a Loop

Here is the loop code for the FIR filter of Application Example 2.1:

```
for (i = 0, f = 0; i < N; i++)
    f = f + c[i] * x[i];
```

By examining the CDFG for the code, we can more easily determine how many times various statements are executed. Here is the CDFG once again:



The CDFG makes it clear that the loop initiation block is executed once, the test is executed  $N + 1$  times, and the body and loop variable update are each executed  $N$  times.

Example 5.6 has very simple behavior: the loop executes for a fixed number of iterations and the loop body contains no conditionals. The next example makes a small change to the FIR filter that results in substantially more complex behavior.

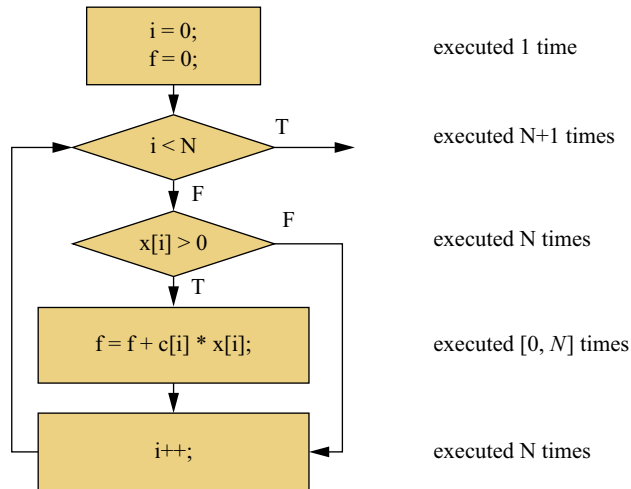
### Example 5.7: Conditional Behavior in a Loop

Here is a slightly more complex version of the FIR filter:

```
for (i=0, f=0; i<N; i++) {
    if (x[i] > 0)
        f = f + c[i] * x[i];
}
```

The result is updated conditionally based on the value of  $x[i]$ .

The CDFG for this code is more complex than that for the basic FIR filter:



The body of the *if* statement that tests  $x[i] > 0$  is executed a variable number of times. We can say that it is executed no more than  $N$  times, as it can be executed at most once per loop iteration. However, the exact number of times that it will be executed depends on the data values that are input into the program. Without knowing something about the behavior of the data values, we cannot bound the number of times that this statement is executed more precisely. As a result, we can only write the execution time of the code as being bounded by its best case (no values of  $x[i]$  are greater than zero) and its worst case (all values of  $x[i]$  are greater than zero).

### Instruction timing

Once we know the execution path of the program, we have to measure the execution time of the instructions that are executed along that path. The simplest estimate is to assume that every instruction takes the same number of clock cycles, which means that we need only count the instructions and multiply by the per-instruction execution time to obtain the program's total execution time. However, even when ignoring cache effects, this technique is simplistic for the reasons summarized below.

- *Not all instructions take the same amount of time.* RISC architectures tend to provide uniform instruction execution times to keep the CPU's pipeline full. However, as we saw in [Chapter 3](#), even very simple RISC architectures like the PIC16F take different amounts of time to execute certain instructions. Floating-point instructions show especially wide variations in execution time; while basic multiply and add operations are fast, some transcendental functions can take thousands of cycles to execute.
- *The execution times of instructions are not independent.* The execution time of one instruction depends on the instructions around it. For example, many CPUs use register bypassing to speed up instruction sequences when the result of one

instruction is used in the next instruction. As a result, the execution time of an instruction may depend on whether its destination register is used as a source for the next operation (or vice versa).

- *The execution time of an instruction may depend on operand values.* This is clearly true of floating-point instructions in which a different number of iterations may be required to calculate the result. Other specialized instructions can, for example, perform a data-dependent number of integer operations.

We can handle the first two problems more easily than the third. We can look up instruction execution times in a table; the table will be indexed by opcode and possibly by other parameter values such as the registers used. To handle interdependent execution times, we can add columns to the table to consider the effects of nearby instructions. Because these effects are generally limited by the size of the CPU pipeline, we know that we need to consider a relatively small window of instructions to handle such effects. Handling variations due to operand values is difficult without executing the program using a variety of data values, given the large number of factors that can affect value-dependent instruction timing. Luckily, these effects are often small. Even in floating-point programs, most of the operations are typically additions and multiplications, whose execution times have small variances.

**Caching effects**

Thus far, we have not considered the effect of the cache. Because the access time for main memory can be 10–100 times larger than the cache access time, caching can have huge effects on the instruction execution time by changing both the instruction and data access times. Caching performance is inherently dependent on the program’s execution path because the cache’s contents depend on the history of accesses.

The next example studies the effects of caching on the FIR filter.

---

**Example 5.8: Caching Effects on the FIR Filter**

Here is the loop code for our FIR filter:

```
for (i = 0, f = 0; i < N; i++)
    f = f + c[i] * x[i];
```

The value of *f* is kept in a register so it doesn’t have to be fetched from memory. However, *c*[*i*] and *x*[*i*] must be fetched during every iteration. For simplicity, let’s assume that the cache has *L* = 4 words per line:

Line 0	Word 0	Word 1	Word 2	Word 3
Line 1	Word 4	Word 5	Word 6	Word 7

We also assume that the *c* and *x* arrays are placed in memory such that they do not interfere with each other in the cache. In this situation, the time required to read the next value of *c* or *x* depends on that word’s position in the cache line. An array entry corresponding to the first entry in the cache line will result in a cache miss and will require *t*<sub>miss</sub> cycles. The other entries will result in cache hits and require *t*<sub>hit</sub> cycles. Remember that the loop body has four instructions, two of which are loads. As the loop body is executed *N* times, we can write the total execution time for all *N* iterations as



$$t_{loop} = 2N + \frac{N}{L}t_{miss} + N\left(1 - \frac{1}{L}\right)t_{hit}$$

This formula assumes that the number of words in the cache line evenly divides the number of loop iterations; it is slightly more cumbersome in the general case.

---

### 5.6.2 Measurement-driven performance analysis

The most direct way to determine the execution time of a program is by measuring it. This approach is appealing, but it does have some drawbacks. First, to cause the program to execute its worst-case execution path, we must provide the proper inputs for it. Determining the set of inputs that will guarantee the worst-case execution path is infeasible. Furthermore, to measure the program's performance on a particular type of CPU, we need the CPU or its simulator.

Despite these problems, measurement is the most commonly used way to determine the execution time of embedded software. Worst-case execution time analysis algorithms have been used successfully in some areas, such as flight control software, but many system design projects determine the execution time of their programs by measurement.

#### Program traces

Most methods of measuring program performance combine the determination of the execution path and the timing of that path: as the program executes, it chooses a path and we observe the execution time along that path. We refer to the record of the execution path of a program as a **program trace** (or more succinctly, a **trace**). Traces can be valuable for other purposes, such as analyzing the cache behavior of the program.

#### Measurement issues

Perhaps the biggest problem in measuring program performance is determining a useful set of inputs to give the program. This problem has two aspects. First, we have to determine the actual input values. We may be able to use benchmark data sets or data that are captured from a running system to help us to generate typical values. For simple programs, we may be able to analyze the algorithm to determine the inputs that cause the worst-case execution time. The software testing methods of [Section 5.10](#) can help us to generate some test values and determine how thoroughly we have exercised the program.

The other problem with input data is the **software scaffolding** that we may need to feed data into the program and get data out. When we are designing a large system, it may be difficult to extract out part of the software and test it independently of the other parts of the system. We may need to add new testing modules to the system software to help us to introduce testing values and to observe the testing outputs.

We can measure program performance either directly on the hardware or by using a simulator. Each method has its advantages and disadvantages.

#### Profiling

**Profiling** is a simple method for analyzing software performance. A profiler does not measure the execution time; instead it counts the number of times that procedures or basic blocks in the program are executed. There are two major ways to profile a

**Physical performance  
measurement**

program: we can modify the executable program by adding instructions that increment a location every time the program passes that point in the program, or we can sample the program counter during execution and keep track of the distribution of PC values. Profiling adds relatively little overhead to the program and it gives us some useful information about where the program spends most of its time.

Physical measurement requires some sort of hardware instrumentation. The most direct method of measuring the performance of a program would be to watch the program counter's value: start a timer when the PC reaches the program's start and stop the timer when it reaches the program's end. Unfortunately, it generally isn't possible to observe the program counter directly. However, in many cases, it is possible to modify the program so that it starts a timer at the beginning of execution and stops the timer at the end. While this doesn't give us direct information about the program trace, it does give us the execution time. If we have several timers available, we can use them to measure the execution time of different parts of the program.

A logic analyzer or an oscilloscope can be used to watch for signals that mark various points in the execution of the program. However, because logic analyzers have a limited amount of memory, this approach doesn't work well for programs with extremely long execution times.

**Hardware traces**

Some CPUs have hardware facilities for automatically generating trace information. For example, the Pentium family microprocessors generate a special bus cycle, namely a branch trace message, that shows the source and/or destination address of a branch [Col97]. If we record only traces, we can reconstruct the instructions that are executed within the basic blocks, while greatly reducing the amount of memory required to hold the trace.

**Simulation-based  
performance  
measurement**

The alternative to physical measurement of the execution time is simulation. A CPU simulator is a program that takes a memory image for a CPU as input and performs the operations on that memory image that the actual CPU would perform, leaving the results in the modified memory image. For purposes of performance analysis, the most important type of CPU simulator is the **cycle-accurate simulator**, which performs a sufficiently detailed simulation of the processor's internals that it can determine the exact number of clock cycles required for execution. A cycle-accurate simulator is built with detailed knowledge of how the processor works, so that it can take into account all of the possible behaviors of the microarchitecture that may affect the execution time. Cycle-accurate simulators are slower than the processor itself, but a variety of techniques can be used to make them surprisingly fast, running only hundreds of times slower than the hardware itself. A simulator that functionally simulates instructions but does not provide timing information is known as an **instruction-level simulator**.

A cycle-accurate simulator has a complete model of the processor, including the cache. It can therefore provide valuable information about why the program runs too slowly. The next example discusses a simulator that can be used to model many different processors.

### Example 5.9: Cycle-accurate Simulation

SimpleScalar (<http://www.simplescalar.com>) is a framework for building cycle-accurate CPU models. Some aspects of the processor can be configured easily at run time. For more complex changes, we can use the SimpleScalar toolkit to write our own simulator.

We can use SimpleScalar to simulate the FIR filter code. SimpleScalar can model a number of different processors; we will use a standard ARM model here.

We want to include the data as part of the program so that the execution time doesn't include file I/O. File I/O is slow and the time that it takes to read or write data can change substantially from one execution to another. We get around this problem by setting up an array that holds the FIR data. Furthermore, because the test program will include some initialization and other miscellaneous code, we execute the FIR filter many times in a row using a simple loop. Here is the complete test program:

```
#define COUNT 100
#define N 12

int x[N] = {8,17,3,122,5,93,44,2,201,11,74,75};
int c[N] = {1,2,4,7,3,4,2,2,5,8,5,1};

main() {
    int i, k, f;
    for (k=0; k<COUNT; k++) { /* run the filter */
        for (i=0; i<N; i++)
            f += c[i]*x[i];
    }
}
```

To start the simulation process, we compile our test program using a special compiler:

```
% arm-linux-gcc firtest.c
```

This gives us an executable program (by default, a.out) that we use to simulate our program:

```
% arm-outorder a.out
```

SimpleScalar produces a large output file with a great deal of information about the program's execution. Because this is a simple example, the most useful piece of data is the total number of simulated clock cycles that are required to execute the program:

```
sim_cycle    25854 x total simulation time in cycles
```

To ensure that we can ignore the effects of program overhead, we will execute the FIR filter for several different values of COUNT and compare. This run uses COUNT = 100; when we also run COUNT = 1,000 and COUNT = 10,000, we get these results:

COUNT	Total simulation time in cycles	Simulation time for one filter execution
100	25,854	259
1000	155,759	156
10,000	1,451,840	145

Because the FIR filter is so simple and runs in so few cycles, we have to execute it a number of times to wash out all the other overhead of program execution. However, the times for 1000 and 10,000 filter executions are within 10% of each other, so these values are reasonably close to the actual execution time of the FIR filter itself.

---

#### Performance counters

Some CPUs provide hardware **performance counters** to keep track of performance-related events during execution. These registers can be used to count the number of events such as cache and pipeline operations. The ARM Performance Monitoring Unit (PMU) is an example of a performance counter system [Arm21].

---

## 5.7 Software performance optimization

In this section, we will look at several techniques for optimizing software performance, including basic loop and cache-oriented loop optimizations as well as more generic strategies.

### 5.7.1 Basic loop optimizations

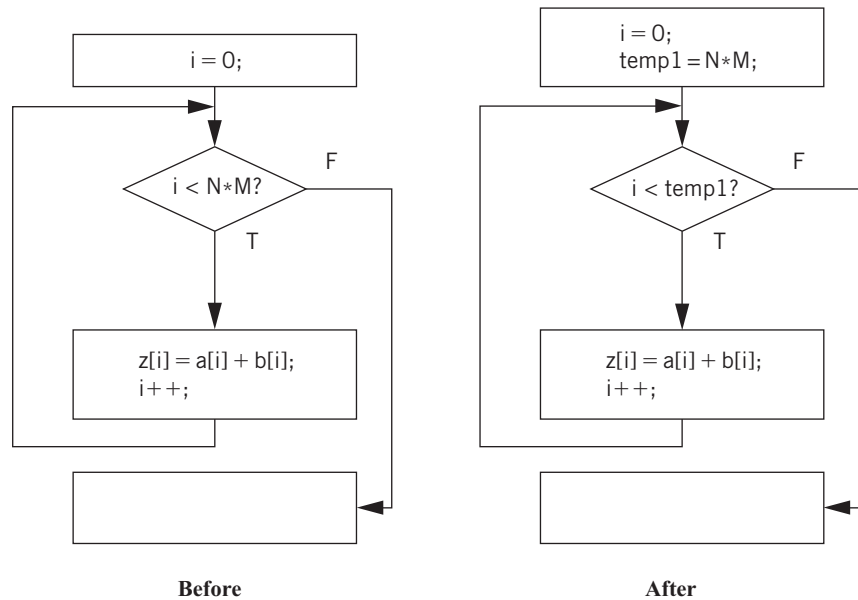
Loops are important targets for optimization, because programs with loops tend to spend a lot of time executing those loops. There are three important techniques in optimizing loops: **code motion**, **induction variable elimination**, and **strength reduction**.

Code motion lets us move unnecessary code out of a loop. If a computation's result does not depend on the operations that are performed in the loop body, we can safely move it out of the loop. Code motion opportunities can arise because programmers may find some computations clearer and more concise when put in the loop body, even though they are not strictly dependent on the loop iterations. A simple example of code motion is also common. Consider this loop:

```
for (i = 0; i < N*M; i++) {
    z[i] = a[i] + b[i];
}
```

The code motion opportunity becomes more obvious when we draw the loop's CDFG, as shown in Fig. 5.21. The loop bound computation is performed on every iteration during the loop test, even though the result never changes. We can avoid  $N \times M - 1$  unnecessary executions of this statement by moving it before the loop, as shown in the figure.

An **induction variable** is a variable whose value is derived from the loop iteration variable's value. The compiler often introduces induction variables to help it to implement the loop. When properly transformed, we may be able to eliminate some variables and apply strength reduction to others.

**FIGURE 5.21**

Code motion in a loop.

A nested loop is a good example of the use of induction variables. Here is a simple nested loop:

```
for (i = 0; i < N; i++)
    for (j = 0; j < M; j++)
        z[i][j] = b[i][j];
```

The compiler uses induction variables to help it to address the arrays. Let us rewrite the loop in C using induction variables and pointers. Later, we use a common induction variable for the two arrays, even though the compiler would probably introduce separate induction variables, and then, merge them.

```
for (i = 0; i < N; i++)
    for (j = 0; j < M; j++) {
        zbinduct = i*M + j;
        *(zptr + zbinduct) = *(bptr + zbinduct);
    }
```

In the above code, `zptr` and `bptr` are pointers to the heads of the `z` and `b` arrays, and `zbinduct` is the shared induction variable. However, we do not need to compute

zbinduct afresh each time. Because we are stepping through the arrays sequentially, we can simply add the update value to the induction variable:

```
zbinduct = 0;
for (i = 0; i < N; i++) {
    for (j = 0; j < M; j++) {
        *(zptr + zbinduct) = *(bptr + zbinduct);
        zbinduct++;
    }
}
```

This is a form of strength reduction because we have eliminated the multiplication from the induction variable computation.

Strength reduction helps us to reduce the cost of a loop iteration. Consider this assignment is

```
y = x * 2;
```

In integer arithmetic, we can use a left shift rather than a multiplication by 2 (as long as we properly keep track of overflows). If the shift is faster than the multiply, we probably want to perform the substitution. This optimization can often be used with induction variables because loops are often indexed with simple expressions. Strength reduction can often be performed with simple substitution rules because there are relatively few interactions between the possible substitutions.

### 5.7.2 Cache-oriented loop optimizations

A **loop nest** is a set of loops, one inside the other. Loop nests occur when we process arrays. A large body of techniques has been developed for optimizing loop nests. Many of these methods are designed to improve the cache performance. Rewriting a loop nest changes the order in which array elements are accessed. This can expose new parallelism opportunities that can be exploited by later stages of the compiler, and it can also improve the cache performance. In this section, we concentrate on the analysis of loop nests for cache performance.

The next example looks at two cache-oriented loop nest optimizations.

---


#### Programming Example 5.7: Data Realignment and Array Padding

Assume that we want to optimize the cache behavior of the following code:

```
for (j = 0; j < M; j++)
    for (i = 0; i < N; i++)
        a[j][i] = b[j][i] * c;
```

Let us also assume that the *a* and *b* arrays are sized with *M* at 265 and *N* at 4 and a 256-line, four-way set-associative cache with four words per line. Even though this code does not reuse any data elements, cache conflicts can cause serious performance problems because they interfere with spatial reuse at the cache line level.

Assume that the starting location for `a[]` is 1024 and the starting location for `b[]` is 4099. Although `a[0][0]` and `b[0][0]` do not map to the same word in the cache, they do map to the same block.



As a result, we see the following scenario in execution:

- The access to `a[0][0]` brings in the first four words of `a[]`.
- The access to `b[0][0]` replaces `a[0][0]` through `a[0][3]` with `b[0][3]` and the contents of the three locations before `b[]`.
- When `a[0][1]` is accessed, the same cache line is again replaced with the first four elements of `a[]`.

Once the `a[0][1]` access brings that line into the cache, it remains there for the `a[0][2]` and `a[0][3]` accesses, because the `b[]` accesses are now on the next line. However, the scenario repeats itself at `a[1][0]` and every four iterations of the cache.

One way to eliminate the cache conflicts is to move one of the arrays. We do not have to move it far. If we move `b`'s start to 4100, we eliminate the cache conflicts.

However, this fix won't work in more complex situations. Moving one array may only introduce cache conflicts with another array. In such cases, we can use another technique called padding. If we extend each of the rows of the arrays to have four elements rather than three, with the padding word placed at the beginning of the row, we eliminate the cache conflicts. In this case, `b[0][0]` is located at 4100 by the padding. Although padding wastes memory, it substantially improves memory performance. In complex situations with multiple arrays and sophisticated access patterns, we must use a combination of techniques, namely relocating and padding arrays, to be able to minimize cache conflicts.

**Loop tiling** breaks a loop up into a set of nested loops, with each inner loop performing the operations on a subset of the data. Loop tiling changes the order in which array elements are accessed, thereby allowing us to control the behavior of the cache better during loop execution. The next example illustrates the use of loop tiling.

---

**Programming Example 5.8: Loop Tiling**

Here is a nest of two loops, with each loop walking through an array:

```
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        z[i][j] = x[i] * y[j];
    }
}
```

Every iteration of the outer *i* loop makes use of every value of *y*[]. The values of *y*[] may conflict with *x*[][] in the cache.

We can improve the cache behavior of *y*[] by dividing it into tiles of size *TILE*, each of which can fit into the cache:

```
for (j=0; j < N; j += TILE) {
    for (i=0; i<N; i++) {
        for (jj=0; jj<TILE; jj++) {
            z[i][j + jj] = x[i] * y[j + jj];
        }
    }
}
```

In this code, each iteration of the *i* loop makes use of one tile of *y*[]. The new, outermost loop iterates over the tiles to ensure that the entire array is covered.

---

**5.7.3 Performance optimization strategies**

Let's look more generally at how to improve the program execution time. First, we must make sure that the code really needs to run faster. Performance analysis and measurement will give you a baseline for the execution time of the program. Knowing the overall execution time tells you how much it needs to be improved. Knowing the execution time of various pieces of the program helps you to identify the correct locations for changes to the program.

You may be able to redesign your algorithm to improve efficiency. Examining asymptotic performance is often a good guide to efficiency. Doing fewer operations is usually the key to performance. However, in a few cases, brute force may provide a better implementation. A seemingly simple high-level-language statement may in fact hide a very long sequence of operations that slows down the algorithm. Using dynamically allocated memory is one example, because managing the heap takes time but it is hidden from the programmer. For example, a sophisticated algorithm that uses dynamic storage may be slower in practice than an algorithm that performs more operations on statically allocated memory.

Finally, you can look at the implementation of the program itself. Here are a few hints on program implementation:

- Try to use registers efficiently. Group accesses to a value together so that the value can be brought into a register and kept there.



- Make use of page mode accesses in the memory system whenever possible. Page mode reads and writes eliminate one step in the memory access. You can increase the use of page mode by rearranging your variables so that more can be referenced contiguously.
- Analyze the cache behavior to find major cache conflicts. Restructure the code to eliminate as many of these as you can, as follows:
  - For instruction conflicts, if the offending code segment is small, try to rewrite the segment to make it as small as possible so that it better fits into the cache. Writing in assembly language may be necessary. For conflicts across larger spans of code, try moving the instructions or padding with NOPs.
  - For scalar data conflicts, move the data values to different locations to reduce conflicts.
  - For array data conflicts, consider either moving the arrays or changing your array access patterns to reduce conflicts.

---


## 5.8 Program-level energy, and power analysis and optimization

Power consumption is a particularly important design metric for battery-powered systems because the battery has a very limited lifetime. However, power consumption is increasingly important in systems that run off the power grid. Fast chips run hot and controlling power consumption is an important element of increasing reliability and reducing the system cost.

How much control do we have over power consumption? Ultimately, we must consume the energy that is required to perform necessary computations. However, there are opportunities for saving power:

- We may be able to replace the algorithms with others that do things in clever ways that consume less power.
- Memory accesses are a major component of power consumption in many applications. By optimizing memory accesses, we may be able to reduce the power significantly.
- We may be able to turn off parts of the system—such as subsystems of the CPU, chips in the system, and so on—when we don’t need them in order to save power.

The first step in optimizing a program’s energy consumption is knowing how much energy the program consumes. It is possible to measure the power consumption for an instruction or a small code fragment [Tiw94]. The technique, which is illustrated in [Fig. 5.22](#), executes the code under test repeatedly in a loop. By measuring the current flowing into the CPU, we are measuring the power consumption of the complete loop, including both the body and other code. By separately measuring the power consumption of a loop with no body (ensuring, of course, that the compiler hasn’t optimized away the empty loop), we can calculate the power consumption of the loop body code as the difference between the full loop and the bare loop energy cost of an instruction.

**FIGURE 5.22**

Measuring energy consumption for a piece of code.

Several factors contribute to the energy consumption of the program:

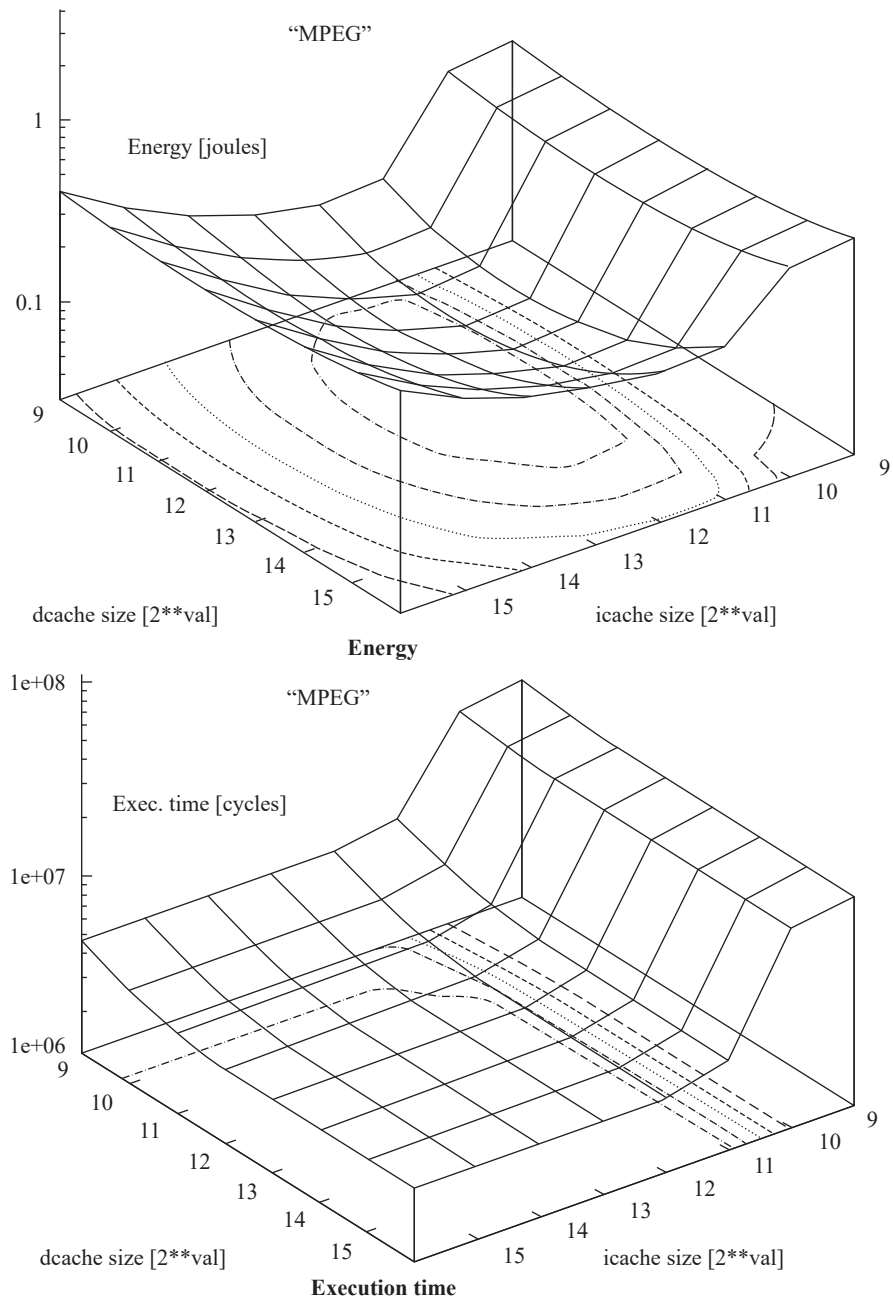
- Energy consumption varies somewhat from instruction to instruction.
- The sequence of instructions has some influence.
- The opcode and locations of the operands also matter.

Choosing which instructions to use can make some difference in a program's energy consumption, but concentrating on the instruction opcodes has limited payoffs in most CPUs. The program must do a certain amount of computation to perform its function. While there may be some clever ways to perform that computation, the energy cost of the basic computation will change by only a small amount compared to the total system energy consumption, and usually only after a great deal of effort. We are further hampered in our ability to optimize instruction-level energy consumption because most manufacturers do not provide detailed, instruction-level energy consumption figures for their processors.

#### Memory effects

In many applications, the biggest payoff in energy reduction for a given amount of designer effort comes from concentrating on the memory system. Memory transfers are by far the most expensive type of operation performed by a CPU [Cat98]; a memory transfer requires tens or hundreds of times more energy than an arithmetic operation. As a result, the biggest payoffs in energy optimization come from properly organizing the instructions and data in memory. Accesses to registers are the most energy efficient; cache accesses are more energy efficient than main memory accesses.

Caches are an important factor in energy consumption. A cache hit saves a costly main memory access and the cache itself is relatively power hungry because it is built from SRAM, not DRAM. If we can control the size of the cache, we want to choose the smallest cache that provides us with the necessary performance. Li and Henkel [Li98] measured the influence of caches on energy consumption in detail. Fig. 5.23 breaks down the energy consumption of a computer running MPEG (a video encoder)

**FIGURE 5.23**

Energy and execution time vs. instruction/data cache size for a benchmark program [Li98].

into several components: the software running on the CPU, main memory, data cache, and instruction cache.

As the instruction cache size increases, the energy cost of the software on the CPU decreases, but the instruction cache starts to dominate the energy consumption. Experiments like this on several benchmarks show that many programs have sweet spots in energy consumption. If the cache is too small, the program runs slowly and the system consumes a lot of power owing to the high cost of main memory accesses. If the cache is too large, the power consumption is high, without a corresponding payoff in performance. At intermediate values, the execution time and power consumption are both good.

#### Energy optimization

How can we optimize a program for low power consumption? The best overall advice is that *high performance = low power*. In general, making the program run faster also reduces the energy consumption.

Clearly, the biggest factor that can be reasonably well controlled by the programmer is the memory access patterns. If the program can be modified to reduce instruction or data cache conflicts, for example, the energy required by the memory system can be significantly reduced. The effectiveness of changes such as reordering instructions or selecting different instructions depends on the processor involved, but they are generally less effective than cache optimizations.

A few optimizations mentioned previously for performance are also often useful for improving energy consumption:

- Try to use registers efficiently. Group accesses to a value together so that the value can be brought into a register and kept there.
- Analyze cache behavior to find major cache conflicts. Restructure the code to eliminate as many of these as you can:
  - For instruction conflicts, if the offending code segment is small, try to rewrite the segment to make it as small as possible so that it better fits into the cache. Writing in assembly language may be necessary. For conflicts across larger spans of code, try moving the instructions or padding with NOPs.
  - For scalar data conflicts, move the data values to different locations to reduce conflicts.
  - For array data conflicts, consider either moving the arrays or changing your array access patterns to reduce conflicts.
- Make use of page mode accesses in the memory system whenever possible. Page mode reads and writes eliminate one step in the memory access, saving a considerable amount of power.

Metha et al. [Met97] present some additional observations about energy optimization:

- Moderate loop unrolling eliminates some loop control overhead. However, when the loop is unrolled too much, the power increases owing to the lower hit rates of straight-line code.

- Software pipelining reduces pipeline stalls, thereby reducing the average energy per-instruction.
- Eliminating recursive procedure calls where possible saves power by getting rid of function call overhead. Tail recursion can often be eliminated; some compilers do this automatically.

---

## 5.9 Analysis and optimization of program size

The memory footprint of a program is determined by the size of its data and instructions. Both must be considered to minimize the program size.

Data provide an excellent opportunity for minimizing size because the data are most highly dependent on the programming style. Because inefficient programs often keep several copies of data, identifying and eliminating duplications can lead to significant memory savings, usually with little performance penalty. Buffers should be sized carefully; rather than defining a data array to a large size that the program will never attain, determine the actual maximum amount of data that is held in the buffer and allocate the array accordingly. Data can sometimes be packed, such as by storing several flags in a single word and extracting them by using bit-level operations.

A very low-level technique for minimizing data is to reuse values. For instance, if several constants happen to have the same value, they can be mapped to the same location. Data buffers can often be reused at several different points in the program. However, this technique must be used with extreme caution, because subsequent versions of the program may not use the same values for the constants. A more generally applicable technique is to generate data on the fly rather than to store it. Of course, the code required to generate the data takes up space in the program, but when complex data structures are involved, there may be some net space savings from using code to generate data.

Minimizing the size of the instruction text of a program requires a mix of high-level program transformations and careful instruction selection. Encapsulating functions in subroutines can reduce the program size when done carefully. Because subroutines have overhead for parameter passing that is not obvious from the high-level language code, there is a minimum-sized function body for which a subroutine makes sense. Architectures that have variable-sized instruction lengths are particularly good candidates for careful coding to minimize the program size, which may require assembly language coding of key program segments. There may also be cases in which one or a sequence of instructions is much smaller than alternative implementations; for example, a multiply—accumulate instruction may be both smaller and faster than separate arithmetic operations.

When reducing the number of instructions in a program, one important technique is the proper use of subroutines. If the program performs identical operations repeatedly, these operations are natural candidates for subroutines. Even if the operations vary somewhat, you may be able to construct a properly parameterized subroutine

that saves space. Of course, when considering the code size savings, the subroutine linkage code must be considered in the equation. There is extra code not only in the subroutine body, but also in each call to the subroutine that handles parameters. In some cases, proper instruction selection may reduce the code size; this is particularly true in CPUs that use variable-length instructions.

Some microprocessor architectures support **dense instruction sets**, which are specially designed instruction sets that use shorter instruction formats to encode the instructions. The ARM Thumb instruction set and the MIPS-16 instruction set for the MIPS architecture are two examples of this type of instruction set. In many cases, a microprocessor that supports the dense instruction set also supports the normal instruction set, although it is possible to build a microprocessor that executes only the dense instruction set. Special compilation modes produce the program in terms of the dense instruction set. Of course, the program size varies with the type of program, but programs using the dense instruction set are often 70%–80% of the size of the standard instruction set equivalents.

---

## 5.10 Program validation and testing

Complex systems need testing to ensure that they work as they are intended. However, bugs can be subtle, particularly in embedded systems, where specialized hardware and real-time responsiveness make programming more challenging. Fortunately, there are many available techniques for software testing that can help us generate a comprehensive set of tests to ensure that our system works properly. We examine the role of validation in the overall design methodology in [Section 9.6](#). In this section, we concentrate on nuts-and-bolts techniques for creating a good set of tests for a given program.

The first question we must ask ourselves is how much testing is enough. Clearly, we cannot test the program for every possible combination of inputs. Because we cannot implement an infinite number of tests, we naturally ask ourselves what a reasonable standard of thoroughness is. One of the major contributions of software testing is to provide us with standards of thoroughness that make sense. Following these standards does not guarantee that we will find all bugs. However, by breaking the testing problem into subproblems and analyzing each subproblem, we can identify testing methods that provide reasonable amounts of testing, while keeping the testing time within reasonable bounds.

We can use various combinations of two major types of testing strategies:

- **Black-box** methods generate tests without looking at the internal structure of the program.
- **Clear-box** (also known as **white-box**) methods generate tests based on the program structure.

In this section, we cover both types of tests, which complement each other by exercising programs in very different ways.

### 5.10.1 Clear-box testing

The control/data flow graph extracted from a program's source code is an important tool in developing clear-box tests for the program. To test the program adequately, we must exercise both its control and data operations.

In order to execute and evaluate these tests, we must be able to control the variables in the program and observe the results of computations, much as in manufacturing testing. In general, we may need to modify the program to make it more testable. By adding new inputs and outputs, we can usually substantially reduce the effort required to find and execute the test. No matter what we are testing, we must accomplish the following three things in a test:

- Provide the program with inputs that exercise the test that we are interested in.
- Execute the program to perform the test.
- Examine the outputs to determine whether the test was successful.

Example 5.10 illustrates the importance of observability and controllability in software testing.

---

#### Example 5.10: Controlling and Observing Programs

Let's first consider controllability by examining the following FIR filter with a limiter:

```
firout = 0.0; /*initialize filter output */
/*compute buff*c in bottom part of circular buffer */
for (j = curr, k = 0; j < N; j++, k++)
    firout += buff[j] * c[k];
/*compute buff*c in top part of circular buffer */
for (j = 0; j < curr; j++, k++)
    firout += buff[j] * c[k];
/*limit output value */
if (firout > 100.0) firout = 100.0;
if (firout < -100.0) firout = -100.0;
```

The above code computes the output of a FIR filter from a circular buffer of values, and then, limits the maximum filter output (much as an overloaded speaker will hit a range limit). If we want to test whether the limiting code works, we must be able to generate two out-of-range values for `firout`: positive and negative. To do this, we must fill the FIR filter's circular buffer with `N` values in the proper range. Although there are many sets of values that will work, it will still take time for us to set up the filter output for each test properly.

This code also illustrates an observability problem. If we want to test the FIR filter itself, we look at the value of `firout` before the limiting code executes. We could use a debugger to set breakpoints in the code, but this is an awkward way to perform a large number of tests. If we want to test the FIR code independently of the limiting code, we need to add a mechanism to observe `firout` independently.

---

**Execution paths**

Being able to perform this process for many tests entails some amount of drudgery, which can be alleviated with good program design that simplifies controllability and observability.

The next task is to determine the set of tests to be performed. We need to perform many different types of tests to be confident that we have identified a large fraction of the existing bugs. Even if we thoroughly test the program using one criterion, that criterion ignores other aspects of the program. Over the next few pages, we will describe several very different criteria for program testing.

The most fundamental concept in clear-box testing is the path of execution through a program. Previously, we considered paths for performance analysis; we are now concerned with making sure that a path is covered and determining how to ensure that the path is in fact executed. We want to test the program by forcing the program to execute along the chosen paths. We force the execution of a path by giving it inputs that cause it to take the appropriate branches. The execution of a path exercises both the control and data aspects of the program. The control is exercised as we take branches; both the computations leading up to the branch decision and other computations that are performed along the path exercise the data aspects.

Is it possible to execute every complete path in an arbitrary program? The answer is no, because the program may contain a `while` loop that is not guaranteed to terminate. The same is true for any program that operates on a continuous stream of data, because we cannot arbitrarily define the beginning and end of the data stream. If the program always terminates, there are indeed a finite number of complete paths that can be enumerated from the path graph. This leads us to the next question: Does it make sense to exercise every path? The answer to this question is *no* for most programs, because the number of paths, especially for any program with a loop, is extremely large. However, the choice of an appropriate subset of paths to test requires some thought.

Example 5.11 illustrates the consequences of two different choices of testing strategies.

---

**Example 5.11: Choosing the Paths to Test**

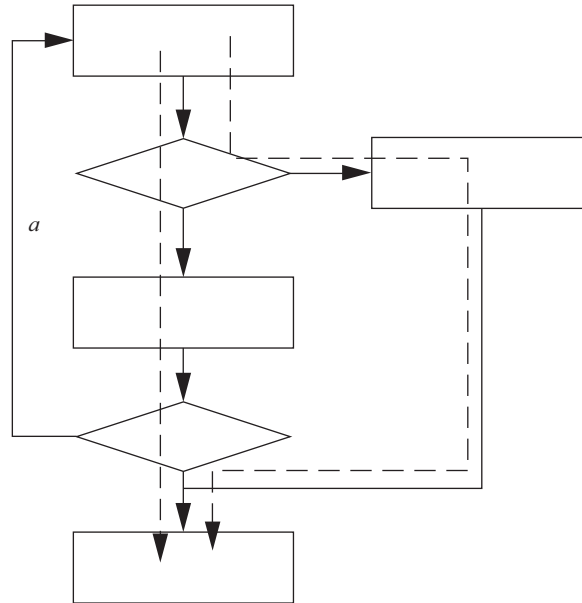
We have at least two reasonable ways to choose a set of paths in a program to test:

- Execute every statement at least once.
- Execute every direction of a branch at least once.

These conditions are equivalent for structured programming languages without `gotos`, but are not the same for unstructured code. Most assembly language is unstructured, and state machines may be coded in high-level languages with `gotos`.



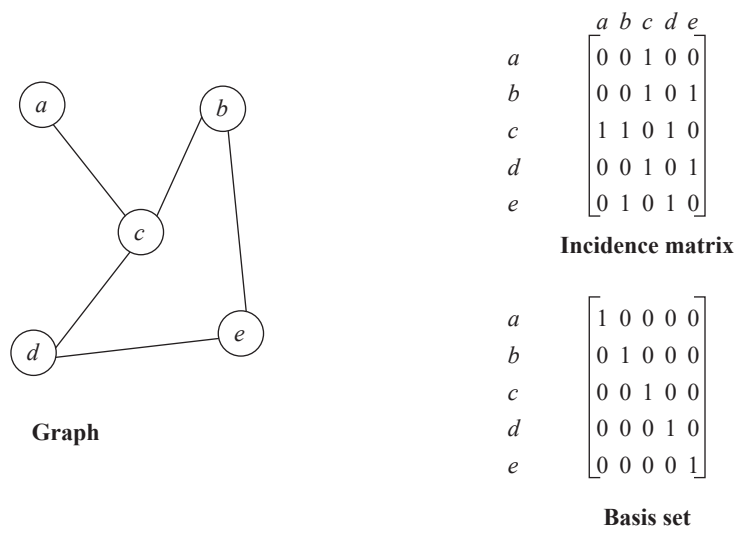
To understand the difference between statement and branch coverage, consider this CDFG:



We can execute every statement at least once by executing the program along two distinct paths. However, this leaves branch *a* out of the lower conditional uncovered. To ensure that we have executed along every edge in the CDFG, we must execute a third path through the program. This path does not test any new statements, but it does cause *a* to be exercised.

How do we choose a set of paths that adequately covers the program's behavior? Intuition tells us that a relatively small number of paths should be able to cover most practical programs. Graph theory helps us to get a quantitative handle on the different paths required. In an undirected graph, we can form any path through the graph from combinations of **basis paths**. Unfortunately, this property does not strictly hold for directed graphs such as CDFGs, but this formulation still helps us to understand the nature of selecting a set of covering paths through a program. The term "basis set" comes from linear algebra. Fig. 5.24 shows how to evaluate the basis set of a graph. The graph is represented as an **incidence matrix**. Each row and column represents a node; 1 is entered for each node pair connected by an edge. We can use standard linear algebra techniques to identify the basis set of the graph. Each vector in the basis set represents a primitive path. We can form new paths by adding the vectors modulo 2. In general, there is more than one basis set for a graph.

The basis set property provides a metric for test coverage. If we cover all of the basis paths, we can consider the control flow to be adequately covered. Although the basis set measure is not entirely accurate because the directed edges of the



**FIGURE 5.24**

The matrix representation of a graph and its basis set.

CDFG may make some combinations of paths infeasible, it does provide a reasonable and justifiable measure of the test coverage.

A simple measure known as **cyclomatic complexity** [McC76] allows us to measure the control complexity of a program. Cyclomatic complexity is an upper bound on the size of the basis set. If  $e$  is the number of edges in the flow graph,  $n$  is the number of nodes, and  $p$  is the number of components in the graph, the cyclomatic complexity is given by

$$M = e - n + 2p \tag{5.1}$$

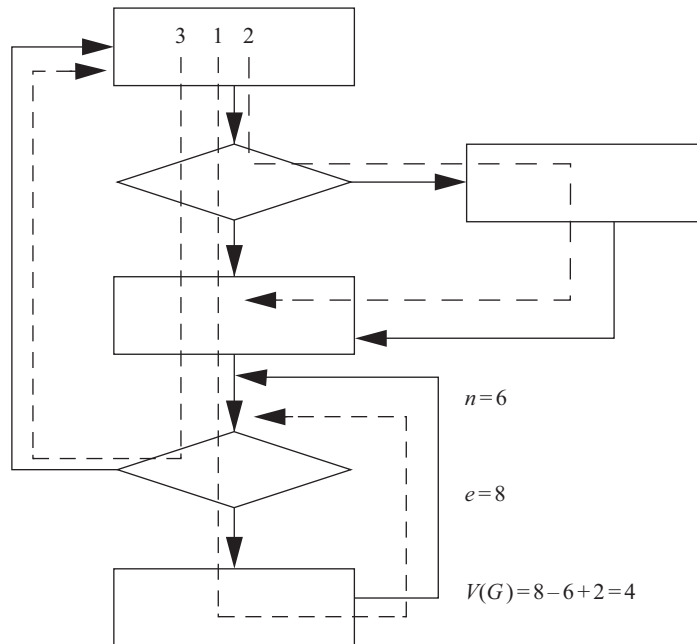
For a structured program,  $M$  can be computed by counting the number of binary decisions in the flow graph and adding 1. If the CDFG has higher-order branch nodes,  $b - 1$  is added for each  $b$ -way branch. In the example of Fig. 5.25, the cyclomatic complexity evaluates to 4. Because there are actually only three distinct paths in the graph, the cyclomatic complexity in this case is an overly conservative bound.

Cyclomatic complexity can be used to identify code that will be difficult to test. A maximum cyclomatic complexity of 10 is widely used [McC76, Wat96].

Another way of looking at control flow—oriented testing is to analyze the conditions that control the conditional statements. Consider the following `if` statement:

```
if ((a == b) || (c >= d)) { ... }
```

This complex condition can be exercised in several different ways. If we want to exercise the paths through this condition fully, it is prudent to exercise the conditional's elements in ways relating to their own structure and not just the structure

**FIGURE 5.25**

Cyclomatic complexity.

of the paths through them. A simple condition testing strategy is known as **branch testing** [Mye79]. This strategy requires the true and false branches of a conditional and every simple condition in the conditional's expression to be tested at least once.

Example 5.12 illustrates branch testing.

### Example 5.12: Condition Testing with the Branch Testing Strategy

Assume that the code below is what we meant to write.


```
if (a || (b >= c)) { printf("OK\n"); }
```

The code that we mistakenly wrote instead follows:

```
if (a && (b >= c)) { printf("OK\n"); }
```

If we apply branch testing to the code that we wrote, one of the tests will use these values:  $a = 0$ ,  $b = 3$ ,  $c = 2$  (making  $a$  false and  $b \geq c$  true). In this case, the code should print the OK term [  $0 \ || \ (3 \geq 2)$  is true], but instead doesn't print [  $0 \ \&\& \ (3 \geq 2)$  evaluates to false]. That test picks up the error.

Another more sophisticated strategy for testing conditionals is known as **domain testing** [How82], as illustrated in Fig. 5.26. Domain testing concentrates on linear


**FIGURE 5.26**

Domain testing for a pair of values.

inequalities. In the figure, the inequality that the program should use for the test is  $j \leq i + 1$ . We test the inequality with three test points—two on the boundary of the valid region and a third outside the region, but between the  $i$  values of the other two points. When we make some common mistakes in typing the inequality, these three tests are sufficient to uncover them, as shown in the figure.

A potential problem with path coverage is that the paths that are chosen to cover the CFG may not have any important relationship to the program's function. Another testing strategy known as **data flow testing** makes use of **def-use analysis** (short for definition-use analysis). It selects paths that have some relationship to the program's function.

The terms "def" and "use" come from compilers, which use def-use analysis for optimization [Aho06]. A variable's value is **defined** when an assignment is made to the variable; it is **used** when it appears on the right side of an assignment (sometimes called a **C-use** for computation use) or in a conditional expression (sometimes called **P-use** for predicate use). A **def-use pair** is a definition of a variable's value and the use of that value. Fig. 5.27 shows a code fragment and all of the def-use pairs for the first assignment to  $a$ . Def-use analysis can be performed on a program using iterative

**FIGURE 5.27**

Definitions and uses of variables.

algorithms. Data flow testing chooses tests that exercise selected def-use pairs. The test first causes a certain value to be assigned at the definition, and then, observes the result at the use point to be sure that the desired value arrived there. Frankl and Weyuker [Fra88] defined criteria for choosing which def-use pairs to exercise to satisfy a well-behaved adequacy criterion.

**Testing loops**

We can write some specialized tests for loops. Because loops are common and often perform important steps in the program, it is worth developing loop-centric testing methods. If the number of iterations is fixed, testing is relatively simple. However, many loops have bounds that are executed at run time.

Consider first the case of a single loop:

```

for (i = 0; i < terminate(); i++)
    proc(i, array);

```

It would be too expensive to evaluate the above loop for all possible termination conditions. However, there are several important cases that we should try at a minimum:

1. skipping the loop entirely [if possible, such as when `terminate()` returns 0 on its first call];
2. one loop iteration;
3. two loop iterations;
4. if there is an upper bound  $n$  on the number of loop iterations (which may come from the maximum size of an array), a value that is significantly below that maximum number of iterations; and
5. tests near the upper bound on the number of loop iterations, that is,  $n - 1$ ,  $n$ , and  $n + 1$ .

We can also have nested loops, like this:

```

for (i = 0; i < terminate1(); i++)
    for (j = 0; j < terminate2(); j++)
        for (k = 0; k < terminate3(); k++)
            proc(i, j, k, array);

```

There are many possible strategies for testing nested loops. One thing to keep in mind is which loops have fixed versus variable numbers of iterations. Beizer [Bei90] suggested an inside-out strategy for testing loops with multiple variable iteration bounds. First, concentrate on testing the innermost loop as above; the outer loops should be controlled to their minimum numbers of iterations. After the inner loop has been thoroughly tested, the next outer loop can be tested more thoroughly, with the inner loop executing a typical number of iterations. This strategy can be repeated until the entire loop nest has been tested. Clearly, nested loops can require a large number of tests. It may be worthwhile to insert testing code to allow greater control over the loop nest for testing.

### 5.10.2 Black-box testing

Black-box tests are generated without knowledge of the code being tested. When used alone, black-box tests have a low probability of finding all of the bugs in a program. However, when used in conjunction with clear-box tests, they help to provide a well-rounded test set, because black-box tests are likely to uncover errors that are unlikely to be found by tests that are extracted from the code structure. Black-box tests can really work. For instance, when asked to test an instrument whose front panel was run by a microcontroller, one acquaintance of the author used his hand to depress all of the buttons simultaneously. The front panel immediately locked up. This situation could occur in practice if the instrument were placed face-down on a table, but the discovery of this bug would be very unlikely via clear-box tests.

One important technique is to take tests directly from the specification for the code under design. The specification should state which outputs are expected for certain inputs. Tests should be created that provide specified outputs and evaluate whether the results also satisfy the inputs.

We can't test every possible input combination, but some rules of thumb help us to select reasonable sets of inputs. When an input can range across a set of values, it is a very good idea to test at the ends of the range. For example, if an input must be between 1 and 10, 0, 1, 10, and 11 are all important values to test. We should be sure to consider tests both within and outside the range, such as testing values within the range and outside the range. We may want to consider tests well outside the valid range as well as boundary condition tests.

#### Random tests

**Random tests** form one category of black-box testing. Random values are generated with a given distribution. The expected values are computed independently of the system, and then, the test inputs are applied. A large number of tests must be applied for the results to be statistically significant, but the tests are easy to generate.

Another scenario is to test certain types of data values. For example, integer-valued inputs can be generated at interesting values such as 0, 1, and values near the maximum end of the data range. Illegal values can be tested as well.

#### Regression tests

**Regression tests** form an extremely important category of tests. When tests are created during earlier stages in the system design or for previous versions of the system, those tests should be saved to apply to the later versions of the system. Clearly,

## Numerical accuracy

unless the system specification has changed, the new system should be able to pass the old tests. In some cases, old bugs can creep back into systems, such as when an old version of a software module is inadvertently installed. In other cases, regression tests simply exercise the code in different ways than would be done for the current version of the code, and therefore, possibly exercise different bugs.

Some embedded systems, particularly digital signal processing systems, lend themselves to numerical analysis. Signal processing algorithms are frequently implemented with limited-range arithmetic to save hardware costs. Aggressive data sets can be generated to stress the numerical accuracy of the system. These tests can often be generated from the original formulas without reference to the source code.

### 5.10.3 Evaluating functional tests

How much testing is enough? Horgan and Mathur [Hor96] evaluated the coverage of two well-known programs, namely *TeX* and *awk*. They used functional tests for these programs that had been developed over several years of extensive testing. Upon applying those functional tests to the programs, they obtained the code coverage statistics shown in Fig. 5.28. The columns refer to various types of test coverage: *block* refers to basic blocks, *decision* refers to conditionals, *P-use* refers to the use of a variable in a predicate (decision), and *C-use* refers to variable use in a nonpredicate computation. These results are at least suggestive that functional testing does not fully exercise the code and that techniques that explicitly generate tests for various pieces of code are necessary to obtain adequate levels of code coverage.

Methodological techniques are important for understanding the quality of your tests. For example, if you keep track of the number of bugs that are tested each day, the data that you collect over time should show you some trends on the number of errors per page of code to expect on average, how many bugs are caught by certain types of tests, and so on. We address methodological approaches to quality control in more detail in Chapter 7.

One interesting method for analyzing the coverage of your tests is **error injection**. First, take your existing code and add bugs to it, keeping track of where the bugs were added. Thereafter, run your existing tests on the modified program. By counting the number of added bugs that your tests found, you can get an idea of how effective the tests are in uncovering the bugs that you haven't yet found. This method assumes that you can deliberately inject bugs that are of similar varieties to those created

	Block	Decision	P-use	C-use
TeX	85%	72%	53%	48%
awk	70%	59%	48%	55%

**FIGURE 5.28**

Code coverage of functional tests for TeX and awk (after Horgan and Mathur [Hor96]).

naturally by programming errors. If the bugs are too easy or too difficult to find, or simply require different types of tests, the bug injection's results will not be relevant. Of course, it is essential that you finally use the correct code and not the code with added bugs.

---

## 5.11 Safety and security

Software testing and eliminating bugs are important aspects of producing secure code. However, some methods that are important for secure software go beyond testing and some security-related bugs are important enough to warrant special attention. A complete discussion of secure program design is beyond our scope, having been the sole topic of several books [Gra03, Che07], but we can identify a few important techniques.

### Buffer overflows

Buffer overflows are a widely available avenue for attacks. If a program reads data from an external source in such a way as to overflow the space that is allocated for the buffer, that external source can change other parts of memory. These changes can be used to insert instructions into the program that are later executed and allow the attacker to take over the program. Transfers into arrays should always check the bounds of the buffer and enforce its limit.

### Buffer initialization

Buffers that have not been initialized to a known value, such as zero, can also result in security leaks. The memory that is used for a buffer is often recycled from some other use during execution. The former values of that memory may contain information that is useful to attackers. If the buffer is not initialized before use, its information may become available to attackers. Graff and van Wyck [Gra03] describe one such incident: a series of bugs caused an uninitialized buffer to be filled with part of the system's password file and for those data to be written out to a software distribution. Although no known problems resulted from this error, it made back-door attacks possible that would allow attackers to intrude into systems.


### Certificates

We often want to install code onto an embedded platform from outside sources. We want to be sure that the code comes from a trusted source, as installing code from a malicious agent would allow malware onto the platform. Establishing the trustworthiness of a source is managed by a **certification authority** [Koh78]. A source requests a **certificate** from the authority; the source may be required to pay the authority for the certificate. The certificate includes an identifier for the source and its encryption key; the certificate may also come with an expiration date. That certificate can then be used to distribute code: the code makes it difficult for an adversary to lie about where the code was produced. The recipient of a certificate can check the source identifier to be sure that it comes from the expected source, and then, use the associated key to decrypt code or messages.

### Code signing

**Code signing** combines the signing certificate with a **digital signature** of the code itself to create a code module with a verifiable source. The code signing process is outlined in Fig. 5.29. First, a **cryptographic hash function** is used to create a **hash** or **digest** of the program. That hash is a compact form of the program generated by a carefully designed function that makes it easy to generate the digest, but difficult



**FIGURE 5.29**


Code signing.

to reconstruct the original input from that digest. The SHA hash algorithms are examples of cryptographic hash functions [NIS15]. That digest is then encrypted using public key cryptography. The signer uses a private key to generate the encrypted version; a public key, which is available to everyone, can be used to decrypt. The original program, its encrypted digest, the public key used to generate the digest, and the signing certificate are then packaged to create a signed program.

The signature of a program can be checked before execution, as shown in Fig. 5.30. The encrypted digest is decrypted using the public key. It should be the same as the hash generated from the program. If the two do not match, the code has been modified after it was signed by the sender and should not be trusted.

## Passwords

Some programs may use passwords to authenticate users; for example, requiring login before allowing some system parameters to be changed. Passwords should be stored only in encrypted form.

**FIGURE 5.30**

Signature verification.


## 5.12 Design example: software modem

In this section, we design a simple modem such as we might use in a basic Internet-of-Things (IoT) node. Before jumping into the modem design itself, we discuss the principles of how to transmit a digital data communications link such as a radio signal or a wire. We will then go through the specification and discuss architecture, module design, and testing.

### 5.12.1 Theory of operation and requirements

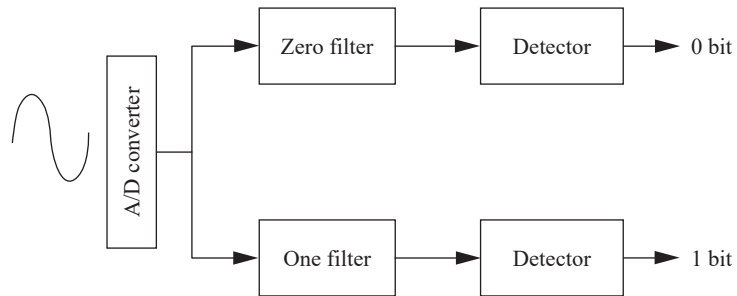
The modem will use **frequency-shift keying (FSK)**, which is a technique used in 1200-baud modems. Keying alludes to Morse code—style keying. As shown in Fig. 5.31, the FSK scheme transmits sinusoidal tones, with 0 and 1 assigned to different frequencies. Sinusoidal tones are much better suited to transmission over analog phone lines than are the traditional high and low voltages of digital circuits. The 01 bit patterns create the chirping sound that is characteristic of modems. Higher-speed modems are backward compatible with the 1200-baud FSK scheme and begin a transmission with a protocol to determine which speed and protocol should be used.

The scheme that is used to translate the audio input into a bit stream is illustrated in Fig. 5.32. The analog input is sampled and the resulting stream is sent to two digital filters (such as a FIR filter). One filter passes frequencies in the range that represents a 0 and rejects the 1-band frequencies, whereas the other filter does the converse. The outputs of the filters are sent to detectors, which compute the average value of the signal over the past  $n$  samples. When the energy goes above a threshold value, the appropriate bit is detected.



**FIGURE 5.31**


Frequency-shift keying.

**FIGURE 5.32**

The FSK detection scheme.

We will send data in units of 8-bit bytes. The transmitting and receiving modems agree in advance on the length of time during which a bit will be transmitted (otherwise known as the baud rate). However, the transmitter and receiver are physically separated and are therefore not synchronized in any way. The receiving modem does not know when the transmitter has started to send a byte. Furthermore, even when the receiver does detect a transmission, the clock rates of the transmitter and receiver may vary somewhat, causing them to fall out of sync. In both cases, we can reduce the chances of error by sending the waveforms for a longer time.

The receiving process is illustrated in [Fig. 5.33](#). The receiver will detect the start of a byte by looking for a start bit, which is always 0. By measuring the length of the start bit, the receiver knows where to look for the start of the first bit. However, because the receiver may have slightly misjudged the start of the bit, it does not immediately try to detect the bit. Instead, it runs the detection algorithm at the predicted middle of the bit.

**FIGURE 5.33**

Receiving bits in the modem.

The modem will not include the physical layer interface such as the radio. We will assume that we have analog audio inputs and outputs for sending and receiving. The modem’s job is to generate a waveform that can be transmitted. We will also run at a much slower bit rate than 1200 baud to simplify the implementation. Furthermore, we will not implement a serial interface to a host, but rather put the transmitter’s message in memory and save the receiver’s result in memory as well. Given those understandings, let’s fill out the requirements table.


Name	Modem
Purpose	A fixed baud rate, frequency-shift keyed modem.
Inputs	Analog sound input, reset button.
Outputs	Analog sound output, LED bit display.
Functions	Transmitter: Sends data stored in microprocessor memory in 8-bit bytes. Sends start bit for each byte equal in length to one bit.  Receiver: Automatically detects bytes and stores results in main memory.
Performance	1200 baud.
Manufacturing cost	Dominated by microprocessor and radio link.
Power	Compatible with IoT node.
Physical size and weight	Compatible with IoT node.

**5.12.2 Specification**

The basic classes for the modem are shown in [Fig. 5.34](#). The classes include physical classes for line-in and line-out plus classes for the receiver and transmitter.


**5.12.3 System architecture**

The modem consists of one small subsystem (the interrupt handlers for the samples) and two major subsystems (transmitter and receiver). Two sample interrupt handlers



**FIGURE 5.34**

Class diagram for the modem.

**FIGURE 5.35**

Waveform generation by table lookup.

are required: one for input and another for output, but they are very simple. The transmitter is simpler, so let's consider its software architecture first.

The best way to generate waveforms that retain the proper shape over long intervals is **table lookup**. Software oscillators can be used to generate periodic signals, but numerical problems limit their accuracy. Fig. 5.35 shows an analog waveform with sample points and the C code for these samples. Table lookup can be combined with interpolation to generate high-resolution waveforms without excessive memory costs, which is more accurate than oscillators because no feedback is involved. The required number of samples for the modem can be found by experimentation with the analog/digital converter and sampling code.

The structure of the receiver is considerably more complex. The filters and detectors of Fig. 5.32 can be implemented with circular buffers. However, that module must feed a state machine that recognizes the bits. The recognizer state machine must use a timer to determine when to start and stop computing the filter output average based on the starting point of the bit. It must then determine the nature of the bit at the proper interval. It must also detect the start bit and measure it using the counter. The receiver sample interrupt handler is a natural candidate to double as the receiver timer, because the receiver's time points are relative to the samples.

The hardware architecture is relatively simple. In addition to the analog/digital and digital/analog converters, a timer is required. The amount of memory that is required to implement the algorithms is relatively small.

#### 5.12.4 Component design and testing

The transmitter and receiver can be tested relatively thoroughly on the host platform because the timing-critical code only delivers data samples. The transmitter's output is quite easy to verify, particularly if the data are plotted. A testbench can be

constructed to feed the receiver code sinusoidal inputs and to test its bit recognition rate. It is a good idea to test the bit detectors first before testing the complete receiver operation. One potential problem in host-based testing of the receiver is encountered when library code is used for the receiver function. If a DSP library for the target processor is used to implement the filters, a substitute must be found or built for the host processor testing. The receiver must then be retested when it is moved to the target system to ensure that it still functions properly with the library code.

Care must be taken to ensure that the receiver does not run too long and miss its deadline. Because the bulk of the computation is in the filters, it is relatively simple to estimate the total computation time early in the implementation process.

### 5.12.5 System integration and testing

There are two ways to test the modem system: by having the modem's transmitter send bits to its receiver and by connecting two different modems. The ultimate test is to connect two different modems, particularly modems that are designed by different people, to be sure that incompatible assumptions or errors were not made. However, single-unit testing, which is called **loop-back** testing in the telecommunications industry, is simpler and a good first step. Loop-back can be performed in two ways. First, a shared variable can be used to pass data directly from the transmitter to the receiver. Second, an audio cable can be used to plug the analog output into the analog input. In this case, it is also possible to inject analog noise to test the resiliency of the detection algorithm.

---

## 5.13 Design example: digital still camera

In this section, we design a simple digital still camera (DSC). Video cameras share some similarities with DSCs but are different in several ways, most notably their emphasis on streaming media. We will study a design example for one subsystem of a video camera in [Chapter 8](#).

### 5.13.1 Theory of operation and requirements

To understand the DSC better, we will first consider the digital photography process. A modern digital camera performs a great many steps:

- Determine the exposure and focus.
- Capture the image.
- Develop the image.
- Compress the image.
- Generate and store the image as a file.

In addition to taking the photo, the camera must perform several other important operations, often simultaneously with the picture-taking process. For example, it

must update the camera's electronic display. It must also listen for button presses from the user, which may command operations that modify the current operation of the camera. The camera should also provide a browser with which the user can review the stored images.

#### Imaging terminology

Much of this terminology comes largely from film photography; for example, the decisions that are made while developing a digital photo are similar to those made in the development of a film image, even though the steps to carry out those decisions are very different. Many variations are possible on all of these steps, but the basic process is common to all DSCs. A few basic terms are useful: the image is divided into **pixels**; a pixel's brightness is referred to as its **luminance**; a color pixel's brightness in a particular color is known as **chrominance**.

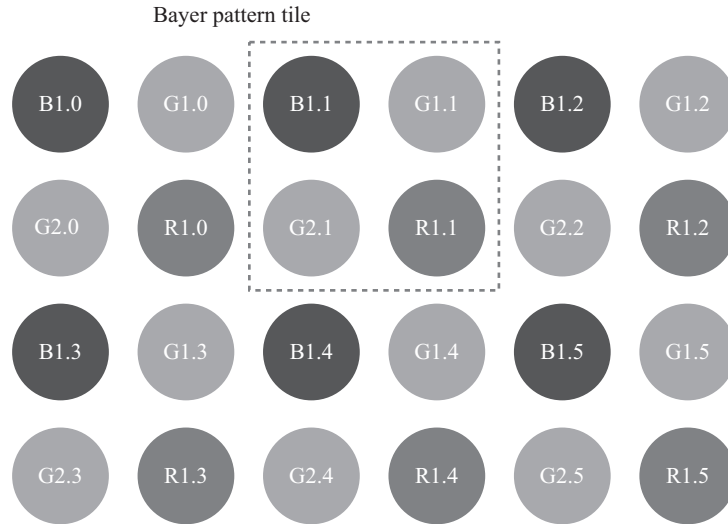
#### Imaging algorithms

The camera can use the image sensor data to drive the exposure-setting process. Several algorithms are possible, but all involve starting with a test exposure and using some search algorithm to select the final exposure. Exposure setting may also use any of several different metrics to judge the exposure. The simplest metric is the average luminance of a pixel. The camera may evaluate some function of several points in the image. It may also evaluate the image's **histogram**. The histogram is composed by sorting the pixels into bins by luminance; 256 bins is a common choice for the resolution of the histogram. The histogram gives us more information than does a single average. For instance, we can tell whether many pixels are overexposed or underexposed by looking at the bins at the extremes of luminance.

Three major approaches are used to determine focus: active rangefinding, phase detection, and contrast detection. Active rangefinding uses a pulse that is sent out, and the time-of-flight of the reflected and returned pulse is measured to determine the distance. Ultrasound was used in one early autofocus system, the Polaroid SX-70, but infrared pulses are more commonly used today. Phase detection compares light from opposite sides of the lens, creating an optical rangefinder. Contrast detection relies on the fact that out-of-focus edges do not display the sharp changes in luminance of in-focus edges. The simplest algorithms for contrast detection evaluate the focus at predetermined points in the image, which may require the photographer to move the camera to place a suitable edge at one of the autofocus spots.

Two major types of image sensors are used in modern cameras [Nak05]: charged-coupled devices (CCDs) and CMOS. For our purposes, the operations of these two in a digital still camera are equivalent.

Developing the image involves both putting the image into a usable form and improving its quality. The most basic operation required for a color image is to interpolate a full color value for each pixel. Most image sensors capture color images using a **color filter array**—color filters that cover a single pixel. The filters usually capture one of the primary colors, namely red, green, and blue, and are arranged in a two-dimensional pattern. The first such color filter array was proposed by Bayer [Bay76]. His pattern is still widely used and is known as the **Bayer pattern**. As shown in Fig. 5.36, the Bayer pattern is a  $2 \times 2$  array with two green, one blue, and one red pixels. More green pixels are used because the eye is most sensitive to green, which Bayer viewed as a simple luminance signal. Because each image

**FIGURE 5.36**

A color filter array arranged in a Bayer pattern.

sensor pixel captures only one of the primaries, the other two primaries must be interpolated for each pixel. This process is known as **Bayer pattern interpolation** or **demosaicing**. The simplest interpolation algorithm is a simple average, which can also be used as a low-pass filter. For example, we can interpolate the missing values from the green pixel G2.1 as

$$R2.1 = (R1.0 + R1.1)/2$$

$$B2.1 = (B1.1 + B1.4)/2$$

We can use more information to interpolate missing green values. For example, the green value that is associated with red pixel R1.1 can be interpolated from the four nearest green pixels:

$$G1.1 = (G1.1 + G2.1 + G2.2 + G1.4)/4$$

However, this simple interpolation algorithm introduces color fringes at edges. More sophisticated interpolation algorithms exist to minimize color fringing while maintaining other aspects of image quality.

Development also determines and corrects for **color temperature** in the scene. Different light sources can be composed of light of different colors; color can be described as the temperature of a black body that will emit radiation of that frequency. The human visual system automatically corrects for color temperature. For example, we do not notice that fluorescent lights emit a greenish cast. However, a photo that is taken without color temperature correction will display a cast from the color of light used to illuminate the scene. A variety of algorithms exists to determine the color



## Image compression

temperature. A set of chrominance histograms is often used to judge color temperature. Once a color correction is determined, it can be applied to the pixels in the image.

Many cameras also apply **sharpening algorithms** to reduce some of the effects of pixilation in digital image capture. Once again, various sharpening algorithms are used, which apply filters to sets of adjacent pixels.


Some cameras offer a RAW mode of image capture. The resulting RAW file holds the pixel values without processing. RAW capture allows the user to develop the image manually and offline using sophisticated algorithms and programs. Although generic RAW file formats exist, most camera manufacturers use proprietary RAW formats. Some uncompressed image formats also exist, such as TIFF [Ado92].

Compression reduces the amount of storage that is required for the image. Some **lossless compression** methods are used that do not throw away information from the image. However, most images are stored using **lossy compression**. A lossy compression algorithm throws away information in the image, so that the decompression process cannot reproduce an exact copy of the original image. A number of compression techniques has been developed to reduce the storage space required for an image substantially, without noticeably affecting the image quality. The most common family of compression algorithms is **JPEG** [CCI92] (JPEG stands for **Joint Photographic Experts Group**). The JPEG standard was extended as JPEG 2000, but classic JPEG is still widely used. The JPEG standard itself contains a large number of options, not all of which need to be implemented to conform to the standard.

As shown in Fig. 5.37, the typical compression process used for JPEG images has five main steps:

- color space conversion
- color downsampling
- block-based discrete cosine transform (DCT)
- quantization
- entropy coding

(We call this typical because the standard allows for several variations.) The first step puts the color image into a form that allows optimizations that are less likely to reduce image quality. Color can be represented by a number of different **color spaces** or combinations of colors that, when combined, form the full range of colors. We saw the red/green/blue (RGB) color space in the color filter array. For compression, we convert into the  $Y' C_R C_B$  color space:  $Y'$  is a luminance channel,  $C_R$  is a red channel,



**FIGURE 5.37**

The typical JPEG compression process.

and  $C_B$  is a blue channel. The conversion between these two color spaces is defined by the JPEG File Interchange Format (JFIF) standard [Ham92]:

$$\begin{aligned} y' &= & + (0.299 * R'_D) & + (0.587 * G'_D) & + (0.114 * B'_D) \\ C_R &= & 128 - (0.168736 * R'_D) & - (0.331264 * R'_D) & + (0.5 * R'_D) \\ C_B &= & 128 - (0.5 * R'_D) & - (0.418688 * R'_D) & - (0.081312 * R'_D) \end{aligned}$$

Once in  $Y'C_R C_B$  form,  $C_R$  and  $C_B$  are generally reduced by downsampling. The **4:2:2** method downsamples both  $C_R$  and  $C_B$  to half of their normal resolution only in the horizontal direction. The **4:2:0** method downsamples both  $C_R$  and  $C_B$  both horizontally and vertically. Downsampling reduces the amount of data that are required, and is justified by the human visual system's lower acuity in chrominance. The three **color channels**,  $Y'$ ,  $C_R$ , and  $C_B$ , are processed separately.

The color channels are next separately broken into  $8 \times 8$  **blocks** (the term block specifically refers to an  $8 \times 8$  array of values in JPEG). The DCT is applied to each block. DCT is a frequency transform that produces an  $8 \times 8$  block of **transform coefficients**. It is reversible, so that the original values can be reconstructed from the transform coefficients. DCT does not itself reduce the amount of information in a block. Many highly optimized algorithms exist to compute the DCT, particularly for an  $8 \times 8$  DCT.

The lossiness of the compression process occurs in the quantization step. This step changes the DCT coefficients with the aim to do so in a way that allows them to be stored in fewer bits. Quantization is applied on the DCT rather than on the pixels because some image characteristics are easier to identify in the DCT representation. Specifically, because the DCT breaks up the block according to **spatial frequencies**, quantization often reduces the high frequency content of the block. This strategy tends to reduce the data set size during compression significantly while causing less noticeable visual artifacts.

The quantization is defined by an  $8 \times 8$  **quantization matrix**  $Q$ . A DCT coefficient  $G_{i,j}$  is quantized to a value  $B_{i,j}$  using the  $Q_{i,j}$  value from the quantization matrix:

$$B_{i,j} = \text{round}\left(\frac{G_{i,j}}{Q_{i,j}}\right)$$

The JPEG standard allows for different quantization matrices but gives a typical matrix that is widely used:

$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

This matrix tends to zeros in the lower right coefficients. Higher spatial frequencies (and therefore, finer detail) are represented by the coefficients in the lower and right parts of the matrix. Putting these to zero eliminates some fine detail from the block.


Quantization does not directly provide for a smaller representation of the image. Entropy coding (lossless encoding) recodes the quantized blocks in a form that requires fewer bits. JPEG allows multiple entropy coding algorithms to be used. The most common algorithm is Huffman coding. The encoding can be represented as a table that maps a fixed number of bits into a variable number of bits. This step encodes the difference between the current and previous coefficients, not the coefficient itself. Several different styles of encoding are possible: **baseline sequential** codes one block at a time; **baseline progressive** encodes corresponding coefficients of every block.

Coefficients are read from the coefficient matrix in the zig-zag pattern, as shown in Fig. 5.38. This pattern reads along diagonals from the upper left to the lower right, which corresponds to reading from the lowest to highest spatial frequency. If fine detail is reduced equally in both the horizontal and vertical dimensions, then zero coefficients are also arranged diagonally. The zig-zag pattern increases the length of sequences of zero coefficients in such cases. Long strings of zeros can be coded in a very small number of bits.

The requirements for the digital still camera are given in Fig. 5.39.

### 5.13.2 Specification

A digital still camera must comply with a number of standards. These standards generally govern the format of the output generated by the camera. Standards in



**FIGURE 5.38**

Zig-zag pattern for reading coefficients.

Name	Digital still camera
Purpose	Digital still camera with JPEG compression
Inputs	Image sensor, shutter button
Outputs	Display, flash memory
Functions	Determine exposure and focus, capture image, perform Bayer pattern interpolation, JPEG compression, store in flash file system
Performance	Take one picture in 2 sec.
Manufacturing cost	Approximately \$75
Power	Two AA batteries
Physical size and weight	Approx 4 in × 4 in × 1 in, less than 4 ounces.

**FIGURE 5.39**

Requirements for the digital still camera.

general allow for some freedom of implementation in certain aspects while still adhering to the standard.

#### File formats


The **Tagged Image File Format (TIFF)** [Ado92] is often used to store uncompressed images, although it also supports several compression methods. Baseline TIFF specifies a basic format that also provides flexibility on the image size, bits per pixel, compression, and other aspects of image storage.

The JPEG standard itself allows many different options that can be followed in various combinations. A set of operations that is used to generate a compressed image is known as a **process**. The JFIF standard [Ham92] is a widely used file interchange format. It is compatible with the JPEG standard but specifies a number of items in more detail.

The **Exchangeable Image File Format (EXIF)** standard is widely used to extend the information stored in an image file further. As shown in Fig. 5.40, an EXIF file holds several types of data:

- The metadata section provides a wide range of information: date, time, location, and so on. The metadata are defined as attribute/value pairs. An EXIF file need not contain all possible attributes.
- A **thumbnail** is a smaller version of a file that is used for quick display. Thumbnails are widely used both by cameras to display the image on a small screen and on desktop computers to display a sample of the image more quickly. Storage of the thumbnail avoids the need to regenerate the thumbnail each time, saving both computation time and energy.
- JPEG compression data include tables, such as entropy coding and quantization tables, that are used to encode the image.
- The compressed JPEG image itself takes up the bulk of the space in the file.

The entire image storage process is defined by yet another standard, namely the Design rule for Camera File (DCF) standard [CIP10]. The DCF specifies three major



**FIGURE 5.40**  
Structure of an EXIF file.

steps: JPEG compression, EXIF file generation, and DOS file allocation table (FAT) image storage. DCF specifies a number of aspects of file storage:

- The DCF image root directory is kept in the root directory and DCIM (for “digital camera images”).
- The directories within DCIM have names of eight characters, the first three of which are numbers between 100 and 999, giving the directory number. The remaining five characters are required to be upper-case alphanumeric.
- File names in DCF are eight characters long. The first four characters are upper-case alphanumeric, followed by a four-digit number between 0001 and 9999.
- Basic files in DCF are in EXIF version 2 format. The standard specifies a number of properties of these EXIF files.

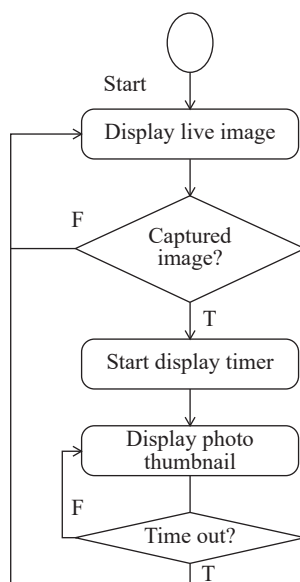
The Digital Print Order Format (DPOF) standard [DPO00] provides a standard way for camera users to order prints of selected photographs. Print orders can be captured in the camera, a home computer, or other devices, and transmitted to a photo-finisher or printer.

**Camera operating modes**

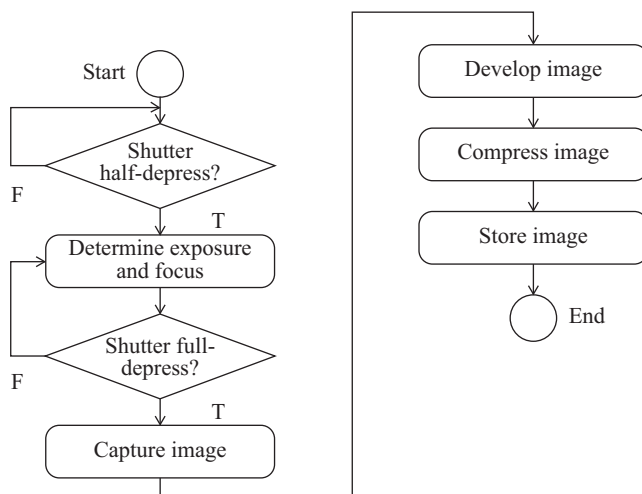
Even a simple point-and-shoot camera provides a number of options and modes. Two basic operations are fundamental: display a live view of the current image and capture an image.

Fig. 5.41 shows a state diagram for the display. In normal operation, the camera repeatedly displays the latest image from the image sensor. After an image is captured, it is briefly shown on the display.

Fig. 5.42 shows a state diagram for the picture-taking process. Depressing the shutter button halfway signals the camera to evaluate the exposure and focus. Once the shutter is fully depressed, the image is captured, developed, compressed, and stored.


**FIGURE 5.41**

State diagram for display operation.


**FIGURE 5.42**

State diagram for picture taking.

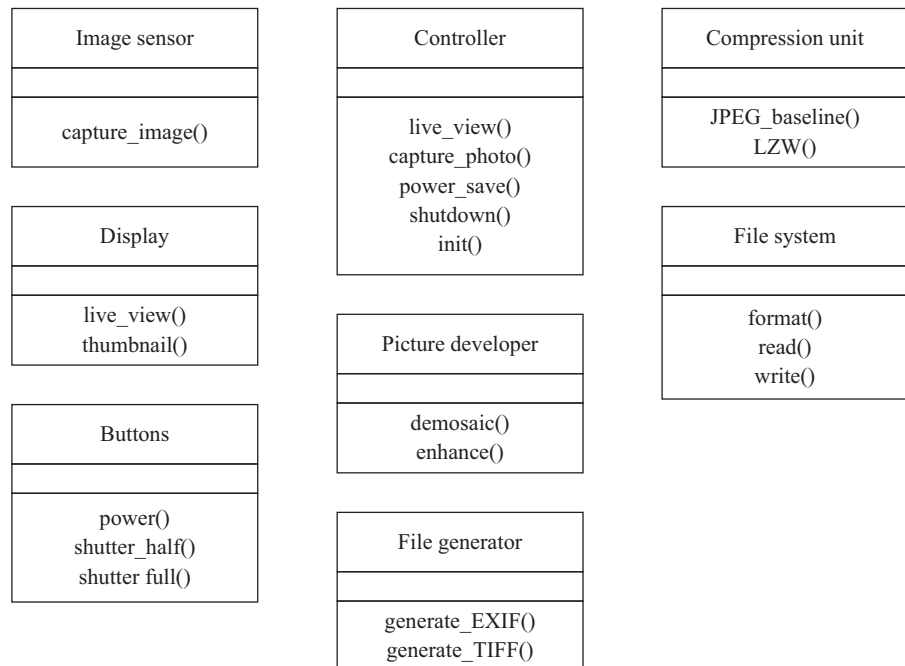
### 5.13.3 System architecture

The basic architecture of a digital still camera uses a controller to provide the basic sequencing for picture taking and camera operation. The controller calls on a number of other units to perform the various steps in each process.

Fig. 5.43 shows the basic classes in a digital still camera. The controller class implements the state diagrams for camera operation. The buttons and display classes provide an abstraction of the physical user interface. The image sensor abstracts the operation of the sensor. The picture developer provides algorithms for mosaicing, sharpening, and so on. The compression unit generates compressed image data. The file generator takes care of aspects of the file generation beyond compression. The file system performs basic file functions. Cameras may also provide other communication ports such as USB or Firewire.

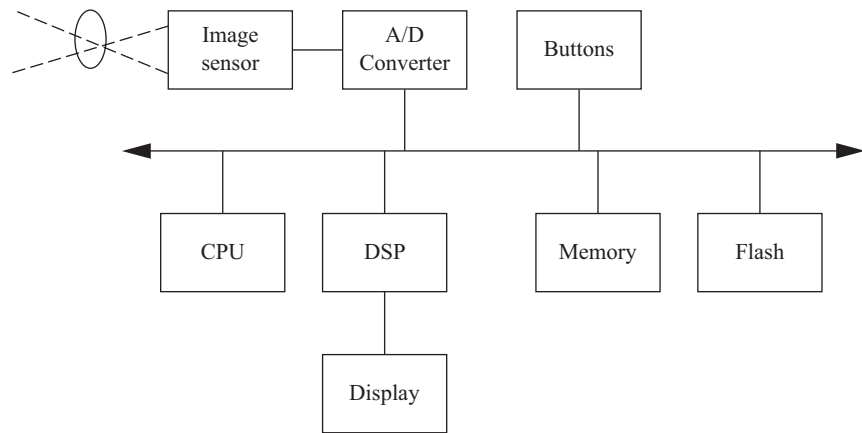
Figure 5.44 shows a typical block diagram for digital still camera hardware. While some cameras will use the same processor for both system control and image processing, many cameras rely on separate processors for these two tasks. The DSP may be programmable or custom hardware.

A sequence diagram for taking a photo is shown in Figure 5.45. This sequence diagram maps the basic operations onto the units in the hardware architecture.



**FIGURE 5.43**

Basic classes in the digital still camera.

**FIGURE 5.44**

Computing platform for a digital still camera.

The design of buffering is very important in a digital still camera. Buffering affects the rate at which pictures can be taken, the energy consumed, and the cost of the camera. The image exists in several versions at different points in the process: the raw image from the image sensor; the developed image; the compressed image data; and the file. Most of these data is buffered in system RAM. However, displays may have their own memory to ensure adequate performance.

#### 5.13.4 Component design and testing

Components based on standards, such as JPEG or FAT, may be implemented using modules developed elsewhere. JPEG compression in particular may be implemented with special-purpose hardware. DCT accelerators are common. Some digital still camera engines use hardware units that produce a complete JFIF file from an image.

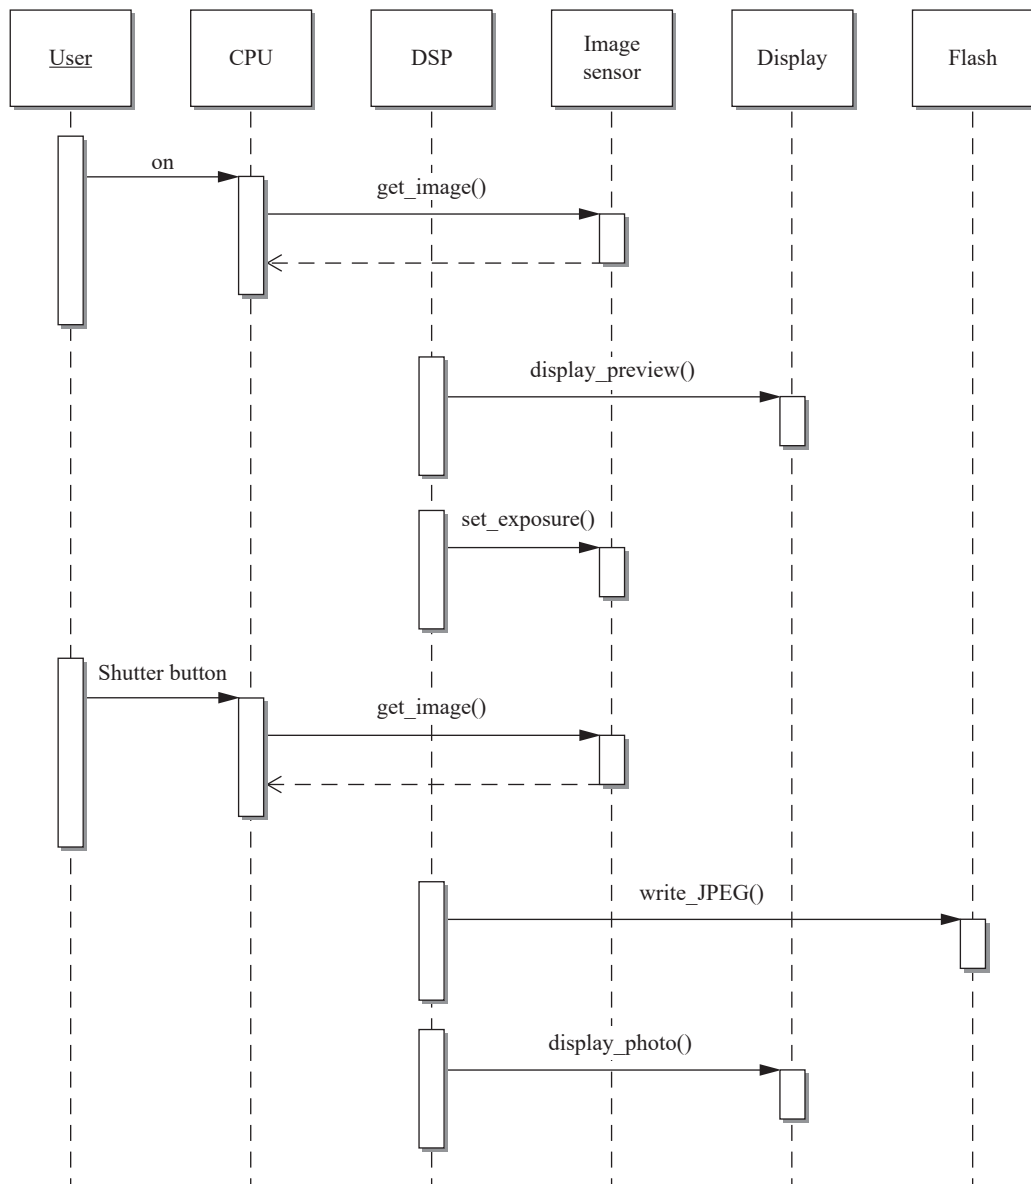
The multiple buffering points in the picture-taking process can help to simplify testing. Test file inputs can be introduced into the buffer by test scaffolding, and then, run with results in the output buffer checked against reference output.

#### 5.13.5 System integration and testing

Buffers help to simplify system integration, although care must be taken to ensure that the buffers do not overlap in main memory.

Some tests can be performed by substituting pixel value streams for the sensor data. Final tests should make use of a target image so that qualities such as sharpness and color fidelity can be judged.



**FIGURE 5.45**

Sequence diagram for taking a picture with a digital still camera.

---

### 5.14 Summary

The program is a very fundamental unit of embedded system design and it usually contains tightly interacting code. Because we care about more than just functionality, we need to understand how programs are created. Because today's compilers do not take directives such as "compile this to run in less than 1 microsecond," we must be able to optimize programs ourselves for speed, power, and space. Our earlier understanding of computer architecture is critical to our ability to perform these optimizations. We also need to test programs to make sure they do what we want. Some of our testing techniques can also be useful in exercising the programs for performance optimization.

---

### What we learned

- We can use data flow graphs to model straight-line code and CDFGs to model complete programs.
- Compilers perform numerous of tasks, such as generating control flow, assigning variables to registers, creating procedure linkages, and so on.
- Remember the performance optimization equation:  $execution\ time = program\ path + instruction\ timing$
- Memory and cache optimizations are very important to performance optimization.
- Optimizing for power consumption often goes hand in hand with performance optimization.
- Optimizing programs for size is possible, but don't expect miracles.
- Programs can be tested as black boxes (without knowing the code) or as clear boxes (by examining the code structure).
- Code signing can be used to authenticate software.

---

### Further reading

Aho, Sethi, and Ullman [Aho06] wrote a classic text on compilers, and Muchnick [Muc97] describes advanced compiler techniques in detail. A paper on the ATOM system [Sri94] provides a good description of instrumenting programs for gathering traces. Cramer et al. [Cra97] describe the Java just-in-time compiler. Li and Malik [Li97D] describe a method for statically analyzing program performance. Banerjee [Ban93, Ban94] describes loop transformations. Two books by Beizer, one on fundamental functional and structural testing techniques [Bei90] and the other on system-level testing [Bei84], provide comprehensive introductions to software testing, and as a bonus, are well written. Lyu [Lyu96] provides a good advanced survey of software reliability. Walsh [Wal97] describes a software modem implemented on an Arm processor.

## Questions

- Q5-1** Write C code for a state machine that implements a four-cycle handshake.
- Q5-2** Use the circular buffer functions to write a C function that accepts a new data value, puts it into the circular buffer, and then, returns the average value of all the data values in the buffer.
- Q5-3** Write C code for a producer/consumer program that takes one value from one input queue, another value from another input queue, and puts the sum of those two values into a separate queue.
- Q5-4** For each basic block given below, rewrite it in single-assignment form, and then, draw the data flow graph for that form.
- ```

x = a + b;
y = c + d;
z = x + e;

```
  - ```

r = a + b - c;
s = 2 * r;
t = b - d;
r = d + e;

```
  - ```

a = q - r;
b = a + t;
a = r + s;
c = t - u;

```
  - ```

w = a - b + c;
x = w - d;
y = x - z;
w = a + b - c;
z = y + d;
y = b * c;

```
- Q5-5** Draw the CDFG for the following code fragments:
- ```

if (y == 2) { r = a + b; s = c - d; }
else r = a - c

```
  - ```

x = 1;
if (y == 2) { r = a + b; s = c - d; }
else { r = a - c; }

```
  - ```

x = 2;
while (x < 40) {
    x = foo[x];
}

```
  - ```

for (i = 0; i < N; i++)
    x[i] = a[i] * b[i];

```

```

e. for (i = 0; i < N; i++) {
    if (a[i] == 0)
        x[i] = 5;
    else
        x[i] = a[i]*b[i];
}

```

**Q5-6** Show the contents of the assembler's symbol table at the end of code generation for each line of the following programs:

```

a.      ORG 200
p1:  ADR r4,a
      LDR r0,[r4]
      ADR r4,e
      LDR r1,[r4]
      ADD r0,r0,r1
      CMP r0,r1
      BNE q1
p2:  ADR r4,e

b.      ORG 100
p1:  CMP r0,r1
      BEQ x1
p2:  CMP r0,r2
      BEQ x2
p3:  CMP r0,r3
      BEQ x3

c.      ORG 200
S1:  ADR r2,a
      LDR r0,[r2]
S2:  ADR r2,b
      LDR r2,a
      ADD r1,r1,r2

d.      ORG 100
L1:  ADR r1,a
      LDR r0,[r1]
L2:  ADR r1,b
      LDR r1,a
L3:  CMP r0,r2
L4:  BEQ L19

```

**Q5-7** Your linker uses a single pass through the set of given object files to find and resolve external references. Each object file is processed in the order given, all external references are found, and then, the previously loaded files are searched for labels that resolve those references. Will this linker be able to successfully load a program with these external references and entry points?

Object file	Entry points	External references
o1	a, b, c, d	s, t
o2	r, s, t	w, y, d
o3	w, x, y, z	a, c, d


**Q5-8** Determine whether each of these programs is reentrant.

- ```
int p1(int a, int b) {
    return( a + b);
}
```
- ```
int x, y;
int p2(int a) {
    return a + x;
}
```
- ```
int x, y;
int p3(int a, int b) {
    if (a > 0)
        x = b;
    return a + b;
}
```


**Q5-9** Is the code for the FIR filter of Programming Example 5.3 reentrant? Explain.

**Q5-10** Provide the required order of execution of operations in these data flow graphs. If several operations can be performed in arbitrary order, show them as a set:  $\{a + b, c - d\}$ .


**a.**



b.



c.



**Q5-11** Draw the CDFG for the following C code before and after applying dead code elimination to the if statement:

```

#define DEBUG 0
proc1();
if (DEBUG) debug_stuff();
  
```

```

switch (foo) {
    case A: a_case();
    case B: b_case();
    default: default_case();
}

```

**Q5-12** Unroll the loop below:

**a.** two times

**b.** three times

```

for (i = 0; i < 32; i++)
    x[i] = a[i] * c[i];

```

**Q5-13** Apply loop fusion or loop distribution to these code fragments as appropriate. Identify the technique you use and write the modified code.

**a.** for (i=0; i<N; i++)

```
    z[i] = a[i] + b[i];
```

for (i=0; i<N; i++)

```
    w[i] = a[i] - b[i];
```

**b.** for (i=0; i<N; i++) {

```
    x[i] = c[i]*d[i];
```

```
    y[i] = x[i] * e[i];
```

```
}
```

**c.** for (i=0; i<N; i++) {

```
    for (j=0; j<M; j++) {
```

```
        c[i][j] = a[i][j] + b[i][j];
```

```
        x[j] = x[j] * c[i][j];
```

```
    }
```

```
    y[i] = a[i] + x[j];
```

```
}
```

**Q5-14** Can you apply code motion to the following example? Explain.

```

for (i = 0; i < N; i++)
    for (j = 0; j < M; j++)
        z[i][j] = a[i] * b[i][j];

```

**Q5-15** For each of the basic blocks of Q5-4, determine the minimum number of registers required to perform the operations when they are executed in the order shown in the code. You can assume that all computed values are used outside the basic blocks, so that no assignments can be eliminated.

**Q5-16** For each of the basic blocks of Q5-4, determine the order of execution of operations that gives the smallest number of required registers. Next, state the number of registers required in each case. You can assume that all computed values are used outside the basic blocks, so that no assignments can be eliminated.

**Q5-17** Draw a data flow graph for the code fragment of Example 5.6. Assign an order of execution to the nodes in the graph so that no more than four registers are required. Explain how you arrived at your solution using the structure of the data flow graph.

**Q5-18** Determine the longest path through each code fragment, assuming that all statements can be executed in equal time and that all branch directions are equally probable.

```
a.  if (i < CONST1) { x = a + b; }
      else { x = c - d; y = e + f; }

b.  for (i = 0; i < 32; i++)
      if (a[i] < CONST2)
          x[i] = a[i] * c[i];

c.  if (a < CONST3) {
      if (b < CONST4)
          w = r + s;
      else {
          w = r - s;
          x = s + t;
      }
  } else {
      if (c > CONST5) {
          w = r + t;
          x = r - s;
          y = s + u;
      }
  }
```

**Q5-19** For each code fragment, list the sets of variable values required to execute each assignment statement at least once. Reaching all assignments may require multiple independent executions of the code.

```
a.  if (a > 0)
      x = 5;
      else {
          if (b < 0)
              x = 7;
      }

b.  if (a == b) {
      if (c > d)
          x = 1;
      else
          x = 2;
      x = x + 1;
  }
```

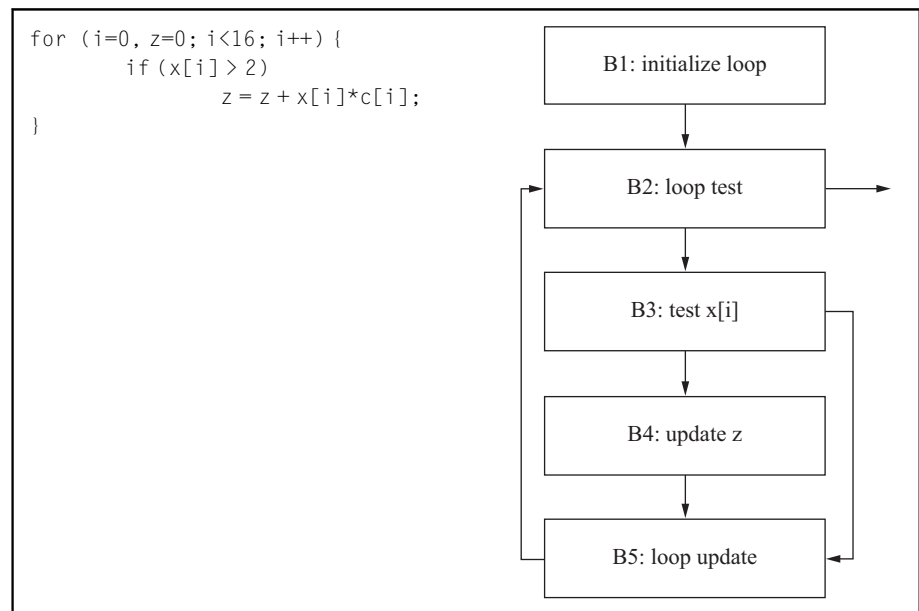


- c. `if (a + b > 0) {  
    for (i=0; i < a; i++)  
        x = x + 1;  
}`
- d. `if (a - b == 5)  
    while (a > 2)  
        a = a - 1;`

**Q5-20** Determine the shortest path through each code fragment, assuming that all statements can be executed in equal time and that all branch directions are equally probable. The first branch is always taken.

- a. `if (a > 0)  
    x = 5;  
else {  
    if (b < 0)  
        x = 7;  
}`
- b. `if (a == b) {  
    if (c > d)  
        x = 1;  
else  
    x = 2;  
    x = x + 1;  
}`

**Q5-21** You are given this program and its flowchart:



The execution time of the blocks is:  $B1 = 6$  cycles,  $B2 = 2$  if branch taken, 5 if not taken,  $B3 = 3$  if branch taken, 6 if not taken,  $B4 = 7$ ,  $B5 = 1$

- What is the maximum number of times that each block in your flowchart executed?
- What is the minimum number of times that each block in your flowchart executed?
- What is the maximum execution time of the program in clock cycles?  
What is the minimum execution time of the program in clock cycles?

**Q5-22** You are given this program:

```
for (i=0, z=0; i<16; i++) {
    z = z + x[i]*c[i];
}
```

A cache miss costs 6 clock cycles and a cache hit costs 2 clock cycles. Assume that  $x$  and  $c$  do not interfere in the cache and that  $z$  and  $i$  are held in registers. If the cache line can hold  $W$  words, plot  $T_a$ , the total number of cycles required for the array accesses ( $x$  and  $c$ ) during all 16 loop iterations for the values  $2 \leq W \leq 8$ .

**Q5-23** Write the branch tests for each conditional.

- `if ((a > 0) && (b < 0)) f1();`
- `if ((a == 5) && !c) f2();`
- `if ((b || c) && (a != d)) f3();`

**Q5-24** The loop appearing below is executed on a machine that has a 1-K-word data cache with four words per cache line.

- How must  $x$  and  $a$  be placed relative to each other in memory to produce a conflict miss every time the inner loop's body is executed?
- How must  $x$  and  $a$  be placed relative to each other in memory to produce a conflict miss one out of every four times the inner loop's body is executed?
- How must  $x$  and  $a$  be placed relative to each other in memory to produce no conflict misses?

```
For (i = 0; i < 50; i++)
    for (j = 0; j < 4; j++)
        x[i][j] = a[i][j] * c[i];
```

**Q5-25** Explain why the person generating clear-box program tests should not be the person who wrote the code being tested.

**Q5-26** Find the cyclomatic complexity of the CDFGs for each of the code fragments given below.

- a.**

```
if (a < b) {
    if (c < d)
        x = 1;
    else
        x = 2;
} else {
    if (e < f)
        x = 3;
    else
        x = 4;
}
```
- b.**

```
switch (state) {
    case A:
        if (x = 1) { r = a + b; state = B; }
        else { s = a - b; state = C; }
        break;
    case B:
        s = c + d;
        state = A;
        break;
    case C:
        if (x < 5) { r = a - f; state = D; }
        else if (x == 5) { r = b + d; state = A; }
        else { r = c + e; state = D; }
        break;
    case D:
        r = r + 1;
        state = D;
        break;
}
```
- c.**

```
for (i = 0; i < M; i++)
    for (j = 0; j < N; j++)
        x[i][j] = a[i][j] * c[i];
```

**Q5-27** Use the branch condition testing strategy to determine a set of tests for each of the following statements.

- a.**

```
if (a < b || ptr1 == NULL) proc1();
else proc2();
```
- b.**

```
switch (x) {
    case 0: proc1(); break;
    case 1: proc2(); break;
    case 2: proc3(); break;
```

```

    case 3: proc4(); break;
    default: dproc(); break;
}

```

- c. 

```
if (a < 5 && b > 7) proc1();
else if (a < 5) proc2();
else if (b > 7) proc3();
else proc4();
```

**Q5-28** Find all the def-use pairs for each code fragment given below.

- a. 

```
x = a + b;
if (x < 20) proc1();
else {
    y = c + d;
    while (y < 10)
        y = y + e;
}
```
- b. 

```
r = 10;
s = a - b;
for (i = 0; i < 10; i++)
    x[i] = a[i] * b[s];
```
- c. 

```
x = a - b;
y = c - d;
z = e - f;
if (x < 10) {
    q = y + e;
    z = e + f;
}
if (z < y) proc1();
```

**Q5-29** For each of the code fragments of Q5-28, determine values for the variables that will cause each def-use pair to be exercised at least once.

**Q5-30** Assume you want to use random tests on an FIR filter program. How would you know when the program under test is executing correctly?

**Q5-31** What role does a certificate play in signed code?

---

## Lab exercises

**L5-1** Compare the source code and assembly code for a moderate-size program. Most C compilers will provide an assembly language listing with the -S flag. Can you trace the high-level language statements in the assembly code? Can you see any optimizations that can be done on the assembly code?

- L5-2** Write C code for an FIR filter. Measure the execution time of the filter, either using a simulator or by measuring the time on a running microprocessor. Vary the number of taps in the FIR filter and measure execution time as a function of the filter size.
- L5-3** Write C++ code for an FIR filter using a class to implement the filter. Implement as many member functions as possible as inline functions. Measure the execution time of the filter and compare to the C implementation.
- L5-4** Generate a trace for a program using software techniques. Use the trace to analyze the program's cache behavior.
- L5-5** Use a cycle-accurate CPU simulator to determine the execution time of a program.
- L5-6** Measure the power consumption of your microprocessor on a simple block of code.
- L5-7** Use software testing techniques to determine how well your input sequences to the cycle-accurate simulator exercise your program.
- L5-8** Generate a set of functional tests for a moderate-size program. Evaluate your test coverage in one of two ways: Have someone else independently identify bugs and see how many of those bugs your tests catch (and how many tests they catch that were not found by the human inspector); or inject bugs into the code and see how many of those are caught by your tests.